

er account since we only use this for break glass security incidents or related administrative activity by the Infrastructure Realm Owners.

Accessing your GCP Project

1. Sign in to <https://gitlabsandbox.cloud> and navigate to Cloud Infrastructure.
1. If you have not created an account yet, click the Create Individual Account button and use the form to create a new account. If you already have an account, click on the title or the gear icon for the GCP project you want to access.
1. On the Cloud Account details page, click the View Credentials button in the top right corner to open up a popup modal window.
1. You will see the GCP Console URL and username. Since GCP uses Google authentication, you simply need to be signed in with your GitLab email address on Google. HackyStack has added the Owner GCP IAM role to your email address when the project was created. Your project ID is in the format of {emailHandle}-{cloudAccountShortId}. You can choose this project from the dropdown list when accessing a different project in the GCP console.
1. After accessing your project for the first time, you will be prompted to enable the service for each of the GCP services that you want to use in your GCP project. This is expected behavior and will take a few seconds and is a one time step during project initial configuration.

Domain Names

See the Domain Names and DNS Records IT guide internal handbook page for more details.

Terraform Environments

In the HackyStack v1.11 (November 2021) release, we introduced Terraform environment generation for GCP projects and a lot of underlying automation for GitOps with alpha stability. AWS will be supported in a future iteration.

- LucidChart Architecture Diagram Edit Source
- LucidChart Architecture Diagram View Source (password WAUBVZA3bneq4oWx)

How Terraform Environments Work

- New GitLab Omnibus instance for securely hosting GitLab Terraform (GitOps) projects for all team members when they create a Cloud Account with GitLab Sandbox Cloud.
- New GitLab Project templates with Terraform scaffolding and easy-to-use Terraform modules. We provide the foundation for you to use any of the Terraform.io Registry providers or modules with built-in support for the Google Cloud provider.
- Every GitLab Sandbox Cloud GCP project now has an automatically created GitLab group and a starter GitLab project with a GitOps Terraform configuration scaffolding with provisioning automation powered by GitLab CI. This allows team members to start deploying resources with Terraform in just a few minutes without dealing with Terraform set up, while complying with security best standards.
- We will have additional project templates released throughout the coming months that provide pre-configured environments that you can provision with just a few clicks. This includes

Omnibus/Runner/Cluster all-in-one environments, Kubernetes cluster environment, etc. We also have the foundation to be able to explore how to support GitLab Environment Toolkit.

- You can also easily create additional Terraform projects in the Sandbox Cloud UI for different environments or configurations in the same Cloud Account to allow you to isolate your module/resource configuration based on the use case that you're experimenting with.

How to Create a Terraform Environment

1. Sign into <https://gitlabsandbox.cloud>
1. Create a Cloud Account in GCP (GCP Project) or navigate to an existing project.
1. Click the Create Terraform Environment button and fill out the form:
 1. Choose your Cloud Account from the Cloud Account dropdown.
 1. Choose the template you wish to use from the Environment Template dropdown. If this is your first time, use the gcp-sandbox-environment-template-v2- template.
 1. Input a name for your environment in the Environment Name (Alphadash Slug) text field.
1. After the Environment is created, click the View Terraform Configuration button. This is hosted on a new GitLab instance at <https://gitops.gitlabsandbox.cloud>. Your GitLab instance credentials can be found in the View GitOps Credentials button modal.

How to Use Terraform Environments

1. Sign into <https://gitops.gitlabsandbox.cloud> using your generated credentials on <https://gitlabsandbox.cloud>. Keep in mind that this is {firstName}{lastName}-{hash} and not your normal GitLab username.
1. Navigate to the project for the Terraform environment that you just created. You can quickly access the project from the link on the Cloud Account page on <https://gitlabsandbox.cloud>.
1. On your local computer navigate to your ~/.ssh folder and generate an SSH key

```
shell
ssh-keygen -t rsa -b 4096 -C <nameofproject>
```

1. Navigate to terraform/main.tf on this project and copy and paste your public key. See the example below

```
shell
-----
Add your Terraform modules and/or resources below this line
-----

locals {
  sshkey          = "<RSA public key here>"
  normalizedenvprefix = "sr-${var.envprefix}"
  tags            = ["sr-firewall-rule", "${local.normalizedenvprefix}-firewall-rule"]
}
```

1. Run a new CI pipeline. After the Plan job completes, trigger the Deploy job. (Notice how you haven't had to do any configuration).
1. Watch the terraform apply outputs as your new environment is spun up with a sample Ubuntu

virtual machine for testing with. You can add additional Terraform resources as you see fit (see below).

1. Navigate to the GCP console using the link on <https://gitlabsandbox.cloud> to view the deployed VM. Feel free to connect to the VM via SSH using the `gcloud` command or Cloud Shell.

1. Run the GitLab CI job for Destroy to clean up your resources.

1. You can update the `terraform/main.tf` file in the Git repository to add more Terraform resources or modules.

1. Simply run the Deploy CI pipeline job to deploy your resources.

Delete an AWS account or GCP project

Currently, deleting an AWS account or GCP project must be performed manually by IT Ops. This is done during offboarding for operational consistency. However, outside of offboarding, you must make a best effort to delete all resources within the account yourself. There may be a nominal monthly cost of a few dollars a month to keep the account in existence. At this time, this cost is deemed acceptable.

Background Context and Problem Statement

The oversimplified user story is "I need to spin up VM(s) or cluster(s) in GCP or AWS to try something (anything, may not be GitLab product specific), what's the company infrastructure standards for doing that?"

The goal is to create a frictionless approach for technical team members that includes the tagging needed for cost allocation, best practice security configurations, and streamline the provisioning of infrastructure without needing to wait several days or weeks for an access request to be approved and provisioned.

This also reduces the burden on the accounting team that processes expense reports for team members each month. Each team member's account is now part of consolidated billing.

History

Over the years, our non-production infrastructure resources have grown organically without accountability, cost controls, naming conventions, provisioning automation, or security best practices. This includes resources created in the GCP `gitlab-internal` project, AWS `gitlab-np` account, and DigitalOcean using `dev-resources`.

Recent iterations

Epic 257 was created to iterate on our processes. In FY21-Q3, we created company-wide infrastructure standards which solved the "naming things is hard" problem with labels, tags, and naming conventions in our AWS and GCP organization accounts. The infrastructure standards define realms to create separate security boundary namespaces for different use cases. For our sandbox use cases, we've created a sandbox realm for individual users and department realms for shared collaboration projects, notably the Engineering Development realm which allows each of the department groups (functional teams) to have a shared AWS account or GCP project for creating infrastructure.

Current iteration

Jeff Martin developed the first release of HackyStack that powers the GitLab Sandbox Cloud. We developed the tooling in-house since the existing industry tools only solve 1-3 of the Technical Problems We're Solving and we wanted to automate the workflow end-to-end.

We are developing HackyStack as an open source project to allow other infrastructure or IT team to simplify their processes for provisioning sandbox accounts. HackyStack is not designed for individual use (yet). As we evolve, we'll be able to advocate HackyStack to partners and customers for deploying demo, testing, or training infrastructure without long manual provisioning documentation or burdening internal infrastructure team members.

Business and financial impact

- All infrastructure resources are associated with a user and department for cost allocation
- Reduced or eliminated expense reports with AWS invoices for individual usage (consolidated billing)
- Budgets and cost controls with Slack notifications to reduce abandoned test environment costs
- Automated access request and provisioning process for IT Ops
- Standardized organizational hierarchy and naming schema for AWS accounts and GCP projects
- Automated security best practice controls and least privilege rights

Technical problems that HackyStack is solving

1. Self-Service Provisioning: Creating an "easy button" for technical users at a company to get access to an AWS account or GCP project with zero manual provisioning by the IT team.
1. Cloud Agnostic: Providing a universal interface that is cloud provider agnostic so you don't need to create different architecture and provisioning processes for AWS, GCP, etc.
1. Hierarchy: Defining a standard reference architecture for organizational unit hierarchy.
1. Auto Labeling/Tagging: Apply labels and tags to resource for cost management, infrastructure-as-code, and security policy compliance without users needing to remember to add tags.
1. Billing Costs per User: Unified billing metrics across all cloud providers on a per-user, per-account/project, and per-group/team level.
1. Automated Access Requests: Supplementing single sign-on (SSO) providers with pre-auth automated group membership provisioning with seamless manager approval(s)
1. Automated Access Approval Provisioning: Supplementing single sign-on (SSO) providers with post-auth provisioning of infrastructure resources in one or more provider APIs.
1. GitOps Infrastructure-as-Code Provisioning: Automatically creating Git projects with Terraform infrastructure-as-code scaffolding with security best practices that use CI/CD automation.
1. Standardized Infrastructure-as-Code Library: Linking to a curated library of Terraform modules for easily deploying common infrastructure elements that follow company security best practices.
1. Daily Workflow Cost Controls: Slack bots and notifications for users to easily provision or destroy infrastructure and threshold cost/usage report notifications.

Tech Stack

- Laravel - web portal, CLI application, API provisioning handler
- MySQL v5.7 - database
- Terraform v0.13+ - Infrastructure as Code configuration

- AWS API
- Google Cloud
- GitLab API - For Git SCM of Terraform configurations
- GitLab CI - For automated Terraform deployments

This project was built using Laravel instead of other viable languages due to Jeff's prior experience and proficiency with Laravel to achieve the most efficient time to business value. This builds on the success of the GitLab Demo Systems that is powered by the demosys-portal.

For those who are not familiar with Laravel, it is the PHP equivalent of Ruby on Rails and Django and has seen tremendous community popularity in recent years since PHP has made revolutionary improvements in recent years with PHP 5.x and PHP 7.x. This project also allows us to dogfood GitLab CI/CD capabilities for PHP projects.

Roadmap

See the issue trackers for the latest up-to-date information.

- HackyStack issue tracking (open source code feature development)
- CorpSec] Infrastructure issue tracking (strategic or GitLab-specific issues)

Current Projects

1. Released in v0.3 Add GCP project provisioning
1. Released in v1.11 Create GitOps project per Cloud Account
1. Released in v1.11 Add Terraform module library for users
1. it-infra86 Project Playground: Deprecate shared AWS accounts and GCP projects with legacy configurations including dev-resources and support-resources.
1. Released in v0.2 Create new group accounts with Admin CLI provisioning
1. Add self-service provisioning and member management for group accounts
1. Add automated access request audit reporting with IT ops issue tracker
1. Add cost and usage reports for AWS and GCP
1. Add Slack bots for easy infrastructure cost reporting and destroy button
1. Implement and automate the GitLab Environment Toolkit (GET)

Future Planning Themes

Phase 4 - Automated provisioning of AWS accounts and GCP projects for each user and team with streamlined/automated access requests (aka "Automate the manufacturing of everyone's green LEGO board"). This is being achieved with the HackyStack open source project that Jeff is building.

Phase 4.5 - Migrate everyone's resources in shared accounts into respective isolated accounts and apply labels/tags for cost management and reporting. See it-infra86 Project Playground for details.

Phase 5 - Curate centralized library of Terraform modules, Ansible roles, Packer images, Docker images, and other scripts that have best practice security standards are used for deploying common infrastructure (aka "Provide everyone a box of LEGO bricks and the tools to deploy them"). Integrate GitLab Environment Toolkit for deploying GitLab in decentralized test environments (user sandboxes, community member environments, etc). This will be open source

with the community so partners and customer POCs can take advantage of what we have. This will solve Sid's request to ensuring we're all on the same page and using the same library for the millions of GitLab users.

Phase 6 - Create "easy button" for deploying the library of infrastructure (aka the LEGO kits) into a topology builder.

Product and Revenue Enablement - Since a lot of provisioning functionality uses GitLab CI/CD and GitOps, we are dogfooding the GitLab product and allows users to manage their infrastructure-as-code in a GitLab repository and extend capabilities with other GitLab features as they see fit.

HackyStack and GitLab Premium Features - We will eventually add premium features to HackyStack that would require features included with a GitLab paid subscription.

How to Contribute

Please post your ideas on Slack or in the issue tracker so we can discuss the best ways to implement them.

The GitLab Sandbox Cloud is a part-time side project for Jeff Martin, so we are not able to commit to firm timelines for feature development. We will try to resolve bugs promptly, however please comment in an issue and tag @jeffersonmartin if there is a feature or capability that needs to be prioritized.

Please direct any questions about the Sandbox Cloud or HackyStack to Jeff Martin or Dave Smith.

SECTION: company/infrastructure-standards/realms/sandbox/labels-tags.md

Quick Links

- Infrastructure Standards - Sandbox Realm Documentation
- Infrastructure Standards - Global Labels and Tags Standards
- Infrastructure Standards - Policies
- Infrastructure Standards - Tutorials
- Infrastructure Standards - Helpdesk

Quick Reference

We use the gl prefix for all labels and tags. All keys use underscores (snakecase). All values should use hyphens (alpha-dash for slug'd values), however underscores are allowed.

Global Labels/Tags

See the list of global labels and tags that should be applied to each resource.

Sandbox Realm Labels/Tags

Slug/Label/Tag/Key	Usage	Human Name
Documentation		
----- ----- -----		
----- ----- -----		
glsandboxshutdownafterdays	Required	Shutdown After Days
Usage Documentation		
glsandboxshutdownworkinghours	Optional	Shutdown Outside Working Hours
Usage Documentation		
glsandboxshutdowndate	Auto Calculated	Shutdown Date
Usage Documentation		

Shutdown after Days (glsandboxshutdownafterdays)

{{% alert color="success" %}}

This label/tag is required. Some Terraform modules will automatically calculate this for you.

{{% /alert %}}

terraform

glsandboxshutdownafterdays: '14'

The value set is used to calculate the glsandboxshutdowndate. Our sandbox automation scripts will handle the shutdown of compute and other resources for you.

Our scripts will rescan shutdown resources and permanently destroy them after 90 days if they have not been powered up since the scripts powered them down.

Expected Values

Value	Recommended Usage
----- -----	
-	
14	(default) The default for all environments unless overridden.
30 to 45	This is the recommended range for longer-term test environments.
2 to 90	The range of accepted integer values that do not require infra/security review unless a vulnerability is detected.
91 to 365	These values can be used but an infra/security review will be performed.
null	This is the recommended value for perpetual environments that do not power off. All environments with a null value are subject to an infra/security review.

Shutdown resources outside of working hours (glsandboxshutdownworkinghours)

{{% alert color="info" %}}

This label/tag is recommended. Some Terraform modules will automatically calculate this for you. This is for future proofing and may not be fully implemented in this iteration.

{{% /alert %}}

terraform

glsandboxshutdownworkinghours: 'true'

Expected Values

Value	Description
true	(default) Resources can be systematically powered down overnight between 6pm (18:00) and 6am (06:00) and on weekends during the environment owner's timezone as set in envtimezone. This reduces compute hour spending from 168 hours per week to 60 hours per week, a 65% savings.
false	All resources will run 24/7 without any disruption. This consumes 720 compute hours per month.

Expiration Date for Resources to Shutdown (glsandboxshutdowndate)

{{% alert color="info" %}}

This label/tag is optional. Some Terraform modules will automatically calculate this for you.

{{% /alert %}}

terraform

glsandboxshutdowndate: '2020-06-30'

The expiration time is calculated as 23:59:59 UTC on the date calculated from the glsandboxshutdownafterdays value and the timestamp the resources was initially provisioned in the cloud provider. Due to limitations with the allowed characters in labels, we do not include the time since it requires colons and plus symbols for timezones.

This value can be updated/extended systematically as part of our expiration scripts (and future Slack notifications) if a user needs more time. After the expiration date has passed, an automated terraform destroy CI pipeline will run to remove all resources in the environment.

SECTION: company/infrastructure-standards/realms/security/_index.md

Quick links

- Global infrastructure standards
- Global labels and tags
- Infrastructure policies
- Infrastructure helpdesk

Overview

This realm is for the engineering security team to deploy shared and team-specific infrastructure resources.

{{% panel header="Future Iteration with Engineering Infrastructure Handbook Pages" header-bg="info" %}}

The Engineering Infrastructure Environments handbook page is the current SSOT for environments.

As the WIP initiative to iterate on our company-wide infrastructure standards evolves, the Engineering Infrastructure pages will be refactored incrementally as the standards are documented, implemented, and changes to environments take place.

{{% /panel %}}

Access requests

To request access to a group, please see [group access request tutorial](#).

> For email authenticity security reasons, only GitLab issues or Slack messages to owners or counterparts are allowed for infrastructure requests.

Realm Owners

Name	GitLab.com Handle	Group Role	Job Title
-----	-----	-----	-----
-----	-----	-----	-----
Marco Lancini	mlancini	Owner	Staff Security Engineer - Infrastructure
Paulo Martins	pmartinsgl	Counterpart	Senior Security Engineer - Infrastructure
James Ritchey	jritchey	Counterpart	Sr. Manager, Product Security

Realm labels and tags

The global labels/tags and realm labels/tags should be applied to each resource.

Realm Groups

Each gldeptgroup has a shared GCP project and/or AWS account for group members.

If a group has not been implemented yet, please contact the realm owner for assistance. After a group is implemented, a separate handbook page is created with usage documentation.

Group Name (AWS Account/GCP Project Name)	Usage Documentation (Empty cells are not implemented yet)
-----	-----
-----	-----
eng-security-shared-infra	<!--Group Docs-->
eng-security-shared-services	<!--Group Docs-->
eng-security-ops-red	<!--Group Docs-->
eng-security-ops-incident-response	<!--Group Docs-->
eng-security-ops-trust-safety	<!--Group Docs-->
eng-security-risk-compliance	<!--Group Docs-->
eng-security-eng-app-sec	<!--Group Docs-->
eng-security-eng-automation	<!--Group Docs-->
eng-security-eng-research	<!--Group Docs-->

Usage guidelines

This is a placeholder for the realm owner to provide instructions on best practices and usage guidelines for this infrastructure.

SECTION: company/infrastructure-standards/realms/security/labels-tags.md

Quick Links

- Infrastructure Standards - Sandbox Realm Documentation
- Infrastructure Standards - Global Labels and Tags Standards
- Infrastructure Standards - Policies
- Infrastructure Standards - Tutorials
- Infrastructure Standards - Helpdesk

Quick Reference

We use the gl prefix for all labels and tags. All keys use underscores (snakecase). All values should use hyphens (alpha-dash for slug'd values), however underscores are allowed.

Global Labels/Tags

See the list of global labels and tags that should be applied to each resource.

Realm Labels/Tags

There are no labels and tags specific to this realm yet.

SECTION: company/infrastructure-standards/tutorials/_index.md

Overview

This is a list of tutorials that are available.

If a tutorial is not available, please contact the realm owner on Slack or in a GitLab issue for assistance.

Realms

- Create Realm

Groups

- Access Request

SECTION: company/infrastructure-standards/tutorials/groups/_index.md

Overview

This is a list of tutorials that are available.

If a tutorial is not available, please contact the realm owner on Slack or in a GitLab issue for assistance.

Groups

- Access Request

SECTION: company/infrastructure-standards/tutorials/groups/access-request.md

Overview

This is a placeholder for a tutorial on requesting access to a group.

In the interim, please contact the realm owner on Slack or in a GitLab issue for assistance.

SECTION: company/infrastructure-standards/tutorials/realms/_index.md

Overview

This is a list of tutorials that are available.

If a tutorial is not available, please contact the realm owner on Slack or in a GitLab issue for assistance.

Realms

- Create Realm

SECTION: company/infrastructure-standards/tutorials/realms/create-realm.md

Overview

This is a placeholder for a tutorial on creating a new realm.

In the interim, please contact Dave Smith or Jeff Martin on Slack or in a GitLab issue for assistance.

SECTION: company/offsite/_index.md

The E-Group offsite happens every quarter. It is usually four days after the Board of Directors meeting..

Goal

The goal is to have 25 or 50 minute discussions around topics that benefit from in-person dialogue, require more context and where the E-Group is likely to disagree.

The agenda should include discussions that are:

1. Top of mind
1. Actionable
1. Impactful to the trajectory of the company
1. Cross-functional

Attendees

1. Executives
1. Chief of Staff to the CEO; when not possible, the Director, Strategy and Operations
1. CEO Shadows
1. Executive Business Admin to the CEO (optional)
1. Invited participants: folks invited to participate in one or more specific session

While attendees will include this list of folks, there will be conversations that are most relevant to a sub-set of people. In these cases, calendars will specify who is "recommended" and who is "optional." In some instances, E-Group members may be among the list of optional attendees.

Roles

Executives are committed to working through difficult discussions and problems during the event.

They commit to being active participants by proposing topics to the agenda before hand and being executive sponsors for those discussion topics.

They establish criteria for success for the discussion topics by, for example, identifying decisions to be made or establishing clarity on what the next steps need to be.

If decisions need to be made to stop, continue, or change discussion when there is disagreement on path forward, the decision is responsibility of the topic executive sponsor.

The CEO Shadows are responsible for taking thorough notes throughout the event, so that the E-Group can be focused on the discussion.

CEO Shadows will also tasked with making Merge Requests on behalf of an Executive.

Please follow the below outlined process for announcing and merging the changes.

The Chief of Staff to the CEO or other team member is responsible for facilitating the event.

They will work with the EBA closely to ensure the event runs smoothly.

The CoS to the CEO is the on-the-ground person ensuring that the event is kept on-schedule,

discussions are kept on-subject, helping steer the conversation when necessary, guiding conversations towards action items, and ensuring that implementation is about 50% of time.

The Staff Executive Business Administrator to the CEO is responsible for organizing and coordinating the Offsite, including travel, lodging, agendas, and meals.

Logistics

Since most of the E-Group is in the San Francisco Bay Area, we most often go to a location close to a member of the E-Group that is drivable or a short direct flight from San Francisco. The E-Group offsite can also be hosted at the house of the CEO.

The EBA to the CEO is responsible for booking the hotel and meeting rooms for the people attending in person.

The EBA to the CEO is responsible for sending out the calendar invites associated to each meeting in the agenda, one calendar entry per meeting, with unique Zoom links for the meeting when the attendees change.

If the CEO is traveling, the CoS to the CEO (preferred) or the EBA to the CEO should work with the CEO to bring the Owl.

Hybrid calls are hard, but occasionally the Offsite will need to take a hybrid meeting form.

When this is the case, the EBA to the CEO will ensure that the calendar invites for the offsite sessions include Zoom links.

The Zoom links should have a waiting room attached to it since the Zoom URL is on calendars and is discoverable internally.

This also allows the E-Group to pull in folks as-needed into the room without switching Zoom rooms, as people won't just jump in and out without being noticed or before the E-Group is ready to move onto that subject.

Location: House of the CEO

We hosted the June 2022 E-Group Offsite meetings at the house of the CEO in San Francisco. When hosting at this location, the following guidelines apply.

The EBA to the CEO will coordinate with on-site contact of the house and collaborate on all details taking place at the house. This includes:

1. Ensuring the AV setup is available for use
1. Meals and catering
1. Schedule including arrival time and departure times
1. Attendee count and details
1. Share a summarized agenda detailing when meetings are taking place, which room will be used for each meeting, and the count of people in each room. If a room will be used by the Directs-Group while they are not in the E-Group Offsite meeting, note this as well.

The maximum attendees that can fit in the main room is 16, however, if additional attendees are needed in the main room, coordinate with the on-site contact to secure additional accommodations such as an extra table, chairs, etc... and share which meetings this would be needed for.

AV Equipment is stored at the house. The following items are on-site:

- Logitech GROUP Video Conferencing System
 - This is powered by a Mac Mini using Zoom Rooms and an iPad as the controller
 - Logitech GROUP Expansion Mics are being ordered for future offsites, and will need to be connected at the October, 2022 E-Group Offsite
- Anker 655 USB-C Hub
- Additional 20' HDMI Cable
- Power Strips
- 2x MacBook Pro Chargers

It is recommended that someone arrive an hour before the first meeting on the first day to set up the AV Equipment, as well as ~30 minutes before the first meeting on subsequent days to ensure everything is logged in appropriately.

The EBA to the CEO should coordinate sharing the needed passwords for the AV Equipment from the CEO to the person who is setting up the equipment in case the accounts have been logged out.

These passwords are for the Zoom Rooms account, the iCloud account, and the iPad.

The AV Equipment is on a cart with wheels and can be moved into the room where needed.

It is recommended to set up the cart at the head of the table, placing the main speaker/microphone in the center of the table while placing the expansion mics towards the two heads of the table.

Schedule

The offsite is a quarterly meeting scheduled over 3 or 4 days.

In-person offsites happen three times a year. The January, end of Q4, offsite is remote. Q1 and Q3 offsites happen during the same weeks and in the same locations as in-person Board.

This section captures all three meeting variants.

In-person offsite with Board Meeting

Scheduling for the in-person event around a Board Meeting generally follows:

1. First Day (Monday): Half day meetings. Start around 1pm PST. A social dinner in the evening.
1. Second Day (Tuesday): 6-7 hours of meetings starting with breakfast at 8am. Includes some unstructured time with the team and team activities.
1. Third Day (Wednesday): 6-7 hours of meetings starting with breakfast at 8am. Includes some unstructured time with the team. Concludes with a dinner with the Board.
1. Fourth Day (Thursday): No scheduled meeting. Time reserved for ad hoc meetings and time with the Board.

In-person offsite without Board Meeting

Scheduling for the in-person event that is not around a Board Meeting generally follows:

1. First Day (Monday): Half day meetings. Start around 1pm PST.
1. Second Day (Tuesday): 6-7 hours of meetings starting with breakfast at 8am. Includes some unstructured time with the team and team activities.
1. Third Day (Wednesday): 6-7 hours of meetings starting with breakfast at 8am. Includes some unstructured time with the team and a team dinner.

1. Fourth Day (Thursday): Meetings start at 8am and run through 12pm PST.

Fully remote offsite

Scheduling for a fully remote offsite generally follows:

1. Monday through Thursday from 8:30am-12:20pm Pacific time each day.

Functional leaders meetings

The E-Group hosts a Functional Leaders Zoom call in the week after the E-Group Offsite

Offsite dates and times

The E-Group offsite generally occurs in:

- March or early April
- June or early July
- September or October
- January

The following date(s) have been confirmed for future E-Group offsite(s) (starting dates):

- 2024-12-09
- 2025-03-31
- 2025-06-23
- 2025-09-22
- 2025-12-08

The following dates were previous E-Group offsites (starting dates):

- 2021-06-28
- 2021-11-01
- 2022-01-10
- 2022-04-04
- 2022-06-13
- 2022-10-04
- 2023-01-09
- 2023-05-03
- 2023-06-26
- 2023-09-27
- 2024-01-08
- 2024-03-25
- 2024-06-24
- 2024-09-23

Offsite topic calendar

While E-Group Offsites often happen about two months in advance of Board Meetings, we reserve time to discuss upcoming Board Meeting topics during each offsite.

Q1 Offsite

- 1. E-Group Offsite and Board Meeting Month: March
- 1. Start review of Three Year Strategy (0.5 day)
- 1. Non-product Engineering review, e.g. architecture, infrastructure

Q2 Offsite

- 1. E-Group Offsite and Board Meeting Month: June
- 1. Review 3 to 5 year Long Range Outlook (LRO) including scenarios, capabilities and dependencies
- 1. Finish the review of Three Year Strategy (1 day)
- 1. Talent assessment and succession planning (2x session per year, 1.5 hrs per session, ~10 mins per function). CEO Shadows and EBAs do not attend this session
- 1. Director + promotion nominations

Q3 Offsite

- 1. E-Group Offsite and Board Meeting Month: September
- 1. Start Annual Plan that includes the budgets and kick off the work streams to get to a final annual plan a quarter later.
- 1. Go-to-market (GTM) learnings
- 1. Org design

Q4 Offsite

- 1. E-Group Offsite and Board Meeting Month: January
- 1. Final Annual Plan that includes the budgets
- 1. Product Roadmap including competitive differentiation
- 1. Talent assessment and succession planning (2x session per year, 1.5 hrs per session, ~10 mins per function). CEO Shadows and EBAs do not attend this session
- 1. Director + promotion nominations

Recurring discussion topics

In addition to the topics from the Topic Calendar above, these topics are discussed at every E-Group offsite:

- 1. Offsite peer feedback session (10 minute prep, 3 minute sessions, 10 minute closing conversation) (50 minutes). This always happens on the first day of the offsite before lunch unless not everyone is present. If not everyone is present, this will happen within half a day of everyone being together
- 1. Director+ Promotions (as-needed based on number of people, estimate ~10 mins per person)
 - 1. To have an efficient meeting and focus on areas that require discussion, we will not spend time discussing a candidate unless questions or concerns are raised. At that point, the E-Group sponsor for the promotion can refer to the candidate promotion case and answer questions. If an E-Group Member wants to document support for a candidate or highlight something positive about a candidate who has not been discussed, they can write a comment in the agenda and note that it is a "[won't verbalize]" item.
 - 1. When a new person is added to the promotion or future promotion list, the sponsorship

exec should provide a summary on the person, their contributions, and any areas to be addressed in advanced of the intended promotion window

1. Content Discussion (30 minutes - 1 hour)
1. Aligning on quarterly priorities (OKRs) (25 minutes)

Pre-offsite discussion topics

These topics are discussed by E-Group before the offsite:

1. Planning discussions. These will be facilitated by the Finance Team. Many of them will occur during E-Group Weeklies.

Collaborating with EBAs on the Offsite

EBAs to E-Group members should have access to the E-Group Offsite outline. This does not mean all individual agenda docs. The EBAs to the CEO should ensure that they have an understanding of the schedule and are aware of which sessions their E-Group leader is leading. They should help to ensure that their E-Group member is inviting folks who should be in sessions, outside of E-Group, to agendas.

Timeline

1. The CoS to the CEO will ask for agenda suggestions at least 3 weeks in advance of the offsite
1. 3 weeks in advance of the offsite, the CEO will share a suggested agenda and solicit additional feedback. If there is a content selection, the content should be shared
1. 1 week, on the Tuesday before the offsite, pre-work should be available

Prep Work

Because the Offsite is a very expensive meeting, we want to be sure we are as efficient as possible.

One of the ways we do this is by asking participants to do prep work ahead of time.

Examples of prep work include:

- Reading the content that was selected
- Executives working with their People Business Partners to prep the Performance/Potential Matrix of their direct reports
- Executives working with their People Business Partners to prep any Director + promotion nominations
- Reviewing any Section Direction Materials for reviews, such as the Growth Section's Direction
- Executives propose draft Key Results for their functions
- Watch this short clip on the danger of sticky bears

Anyone who is presenting for a topic should share materials with the prep work.

Guidance for invited participants

There are a number of reasons why participants are invited to join specific Offsite meetings.

These can range from giving team members a chance for greater context to expecting an invited participant to lead a discussion. Before you attend the Offsite, please:

1. Understand expectations for your participation

1. Ensure that you have context about what is being discussed. If you are invited to a meeting and don't have context, proactively start an async conversation or schedule a 15 minute meeting with the E-Group sponsor and other participants. You should not need to attend a meeting without knowing what will be discussed or do any pre-work without clear direction

1. Don't spend hours preparing unless there is a direct ask to do so

Agenda and Documenting

We take notes in agenda docs that come from individual meeting invites. These are the SSoT for offsite meetings.

At least 1 month prior, an outline document is created. Proposed topics for discussion should be added to the bottom of the doc. Each proposed topic should include a clear definition of what we are trying to achieve in the discussion (a decision made, a chance for E-Group to give feedback, etc.). E-Group sponsors are responsible for providing context to team members who they plan to include in discussion preparation or participation. This could take the form of a 15 minute meeting to align on goals for the discussion and what material should be prepared in advance.

(The CoS to the CEO will pull from here when prioritizing and planning.)

Please add links to relevant materials, issues, or proposals up front.

When there is an issue or doc linked, we take notes there, instead of in the overall doc.

There is a doc template that can be used as a starting point.

If we can conclude a topic early we move on to one from a parking lot.

The CoS to the CEO is responsible for maintaining the schedule, optimizing discussion schedules for energy levels, and having topics prepared.

Meeting agendas will guide conversations. If an E-Group member wants to speak, the team member should put thoughts in the agenda. In instances where the comment cannot be typed out, for reasons of confidentiality or given time constraints in typing out the comment, the member can type out: [INSERT NAME:!). This serves as a time placeholder for the team member, so the team member can contribute at the appropriate point in the agenda.

We will document agreed changes directly to the handbook and any other relevant SSoT during the meeting.

Every item will get a MR maker and MR reviewer assigned.

Most of the time the MR maker will be the CoS to the CEO, one of the CEO shadows, or the EBA to the CEO.

When the MR is ready, the reviewer is at-mentioned in the public E-Group channel in Slack.

The reviewer communicates with the maker via that Slack thread.

The goal is to merge it the same day, preferably within 15 minutes.

Communication guidelines

We have communication guidelines to enable more efficient and productive conversations:

1. Encourage participation from others. We should actively encourage folks who aren't speaking much to speak--especially if we are discussing a topic that should have their direct input. This includes directly asking other folks for their opinions and allowing space for folks to join the conversation.

1. Proactively stop conversations only involving two people if they go on for more than a few minutes and do not require engagement from the broader group. These conversations can be taken offline.

Topic owner responsibilities

If you are the person who owns a topic, you are responsible for:

1. Ensuring that any pre-work is shared at least 5 days before the offsite.

1. The conversation is being managed toward your desired outcome.

1. Encouraging folks to stick to the agenda rather than speaking out of order.

1. The conversation includes all folks who want to contribute. This may mean actively encouraging folks to participate if they have not spoken and may have an item to add.

Team norms

In our discussions, we'll consider the following:

1. Is this topic worth debating?

1. What should we do to evaluate the priority of this idea?

1. The why behind the what. How will we communicate the rationale behind change/priority/direction via multi-modal communication and repetition?

We agreed that we'd use some chosen words to signal that we should pause and reorient. Specifically:

1. Delegate: something is being discussed that should be solved below the E-Group level

1. Telephone: this conversation is not happening directly with folks who should be involved and should be revisited with the appropriate people

Break guidelines

Break periods are scheduled into the agenda by the CoS to the CEO and should follow these guidelines:

1. Breaks should be scheduled for 1/6 of the time (same as speedy meeting)

1. Breaks should be a minimum of 5 minutes

1. Going over time on any topic should result in moving the break. It does not reduce it.

Follow up

Every discussion should start by clarifying "What decision needs to be made from this?"

All follow up actions need to be captured as to-dos, noted with TODO Person in the doc.

If there is not an exec's name tied to the to-do, it belongs to the CEO Shadows/and or the CoS to the CEO.

Before emphasizing follow ups, many conclusions never landed and/or resulted into action.

Follow ups can take the form of an:

1. Merge request to the handbook
1. Issue created
1. Meeting scheduled
1. Notes shared with the rest of the company in Slack
1. Etc.

Functional Leaders Meeting

In the week following the Offsite, there is a 25-minute Zoom call before 12pm Pacific time for the Functional Leaders Meeting. This is called the Functional Leaders Meeting.

The goal of this call is to communicate:

1. Strategic or visionary updates that are crucial to priorities
1. Key decisions made
1. Key messaging that leaders are enlisted to help distribute
1. Action items which may need cross-functional collaboration

In advance of each meeting, we will:

1. Provide a summary of key topics
1. Pick a topic with multiple view points and recap the conversation in greater detail

The Agenda will be organized by meeting topics. Designated Functional Leader team members who led or participated in specific sessions may be asked to provide written summaries and flag highlights. Participants should review the meeting material in advance of the meeting. They can put their comments and questions below each summary. We will run through the agenda during the meeting.

This time will not work for everyone. The meeting will be privately recorded via Zoom and linked in the meeting agenda.

Functional Leaders Responsibilities after Meeting

Functional Leaders play a key role in cascading what is happening in the business and communicating key implications for teams and team members. After each Functional Leaders Meeting, participating team members should:

1. Identify key changes or information of greatest relevance to their teams
1. Reach out to their manager if they have unresolved questions
1. Recognize any sensitivities and think about how best to frame messaging to the team
1. Communicate key insights and asks for team members
1. Help to operationalize any asks or decisions that have been made

Logistics for Functional Leaders Meeting

This Zoom call is a separate invite.

The EBA to the CEO is responsible for setting up the invite, the doc for the call, and inviting

Functional Leaders.

The CoS to the CEO will moderate the call.

The CoS to the CEO is responsible for meeting material. The CoS to the CEO will share the meeting recap material no later than 72 business hours after the offsite concludes and at least 48 hours in advance of the meeting. Material will include an event summary and may include a highlight video message from the CEO or others. All materials will be signed off on by the Legal Team in advance of sharing.

A version of the notes shared with Functional Leaders may be shared with all team members within 4 business days of a Functional Leader meeting. The Office of the CEO prepares materials excluding Functional Leaders Meeting Q&A or any material non-public information and shares them in the company-fyi-private Slack Channel. Functional Leaders are encouraged to discuss relevant information within their teams. You can think of the materials shared with all team members as a "Meeting in a Box."

Content Choice

The offsite includes a 30 minute to 1 hour discussion on material chosen by the CEO. This could be a book, article, recording, or other artifact. The selection will be finalized no less than two weeks prior to the offsite. At the time that the selection is shared, the CEO will also share three conversation topics. E-Group team members are expected to come to the discussion prepared to discuss these prompts.

We will share discussion highlights and takeaways in E-Group Offsite meeting notes that are shared with all team members. This also allows team members to engage in the conversation.

Team Members may expense E-Group offsite material in the quarter that it is discussed.

Material that the E-Group finds noteworthy should be added to the Leadership Books.

SECTION: company/offsite/beware-sticky-bears.md

Be clear

As a leader it's important to remember that what you say leads people to take action. If you voice a vague frustration it might send people on your team looking for a solution you don't need.

Here's an extreme example of what not to do from Gavin Belson, CEO of Hooli.

{{< youtube "i92Ws7qPTRg" >}}

SECTION: company/okrs/2017-q3.md

Summary

With a great team make a popular next generation product that grows revenue.

Objective 1: Grow incremental ACV according to plan

- CEO: Build sales qualified leads (SQL) pipeline for new business. 300% of incremental annual contract value (IACV) plan
- CMO: Identify CE installations within strategic organizations with installations with over 500 users and launch demand generation campaigns. => Achieved. 50 SQOs generated.
- CMO: More demand coming from inbound marketing activities. 50% of sales pipeline quota. => Achieved on SQL \$, miss on
 - CMO: Increase EE Trials. 21% Quarter over Quarter (QoQ) => Miss
 - CMO: Increase EE Trial SQL \$ pipeline contribution 43% QoQ => Achieved over 2x goal
- CFO: Forecast accuracy and delivery of relevant information
 - Controller: Zuora to Salesforce reconciliation complete and tested. DONE
 - CFO: Forecasting model projects forecast +/- 10% normalized for big deals. DONE
 - Controller: Zuora to Salesforce to Zendesk data synch completed. DONE
- CFO: Deliver 10 SQLs > \$10k during quarter through investor outreach 50% ACHIEVED
- VP Prod: Salesforce shows usage. Version check, usage ping, and ConvDev index are available in SFDC. => All shipped
- VP Prod: Import documentation for all major platforms. Five new import docs shipped. => 5/5 shipped. 10 in total now.
- CRO: 100% of quota-carrying reps at 3x pipeline of quarterly IACV goal MISSED
 - Alliances: generates 10% of sales pipeline
 - Federal: generates 10% of sales pipeline
 - AE/Customer Success: generate 15% of sales pipeline quota
 - SDR lead: Generate 25% of sales pipeline quota => 36% of SCLAU was SDR sourced in Q3. Up from 12% in Q2.
 - SDR lead: Get an SQL from Fortune 1000 companies identified using CE. 50% of identified companies using CE. => 50 SQOs generated.
 - SDR lead: Double Win Rate. 10% => 5% of Q3 SDR sourced SCLAU have been won so far. Given sales cycle expect increase over time in the Q3 opportunities.
- CRO: Increase new business sales velocity. Grow 200% YoY (Q3 2016 to Q3 2017) in new business deals closed won. ACHIEVED
- CRO: Decrease age of non-web direct deals by 10% ACHIEVED
- Sales Ops: New sales stages and exit criteria. Launch and Train
- Customer Success: Solution Architects managing PoC's for new business opps within strategic/large accounts. 50%
 - PMM: 2 ROI calculators => New competitive ROI calculator published.
 - PMM: 3 case studies => 3 case studies started. New /customers page launched.
- CRO: Keep sales efficient. Sales efficiency ratio > 1.8 (IACV / sales spend) ACHIEVED
- CMO: Keep marketing efficient. Marketing efficiency ratio > 3.4 (IACV / marketing spend) => 3.43
- CEO: Understand and improve existing account growth. 200% YoY
- VP Eng: Geo DR successful deployments. Test fail-over with GitLab.com => 20% done. Some customers running Geo in alpha now. Failover targeted for Q1 2018.
 - Platform: Get Geo reliable with a testing framework. => 10% done, moved to Q4 with GitLab QA
 - Distribution: Deliver PostgreSQL High Availability (HA) in omnibus and Terraform to enable HA
- VP Prod: Make it easier to discover and use EE features. Ship EE by default. Trials

increase 100%. => Shipped. Trials % TODO

- Technical Writing: CE docs link to EE wherever relevant. Increased visits to EE docs. => Done. Defaults to EE docs now (i.e. over 100% increase)
- VP Prod: Ship paid subscriptions for GitLab.com. All EE features are behind subscription. Ability to buy higher storage limits. => Done. Storage limits deprioritized.
- Head Prod: Improved support for Java development lifecycle. 2 projects done.
- CFO: Make our sales process data driven
- CFO: Lead hired DONE
- Data and Analytics Architect: Data and Analytics vision and plan signed off by executive team NOT STARTED
- Data and Analytics Architect: Create a user journey funnel NOT STARTED
- CMO: Educational Email campaign series to educate users on full solution capabilities and how to get started. TODO of campaign conversions => Issue board anniversary awareness campaign on blog and social. Initiated CI webcast series.
- Customer Success: License utilization for large and strategic accounts at a 90% average
- Customer Success: Double spend for scheduled renewals valued at \$20,000+
- Customer Success: Double spend within large/strategic account segmentation
- Customer Success: Identity the trigger(s) to purchase for large/strategic accounts.
- Support: Provide faster, more knowledgeable support. Reduce Average Time to Solve by 5%. -> Not Hit. SLAs were breached which became the new focus.
- CEO: Increase average sales price (ASP) (IACV per won) by 25% QoQ
- CRO: Triple sales assisted average new business deal YoY ACHIEVED
- CRO: 50% of IACV from EEP ACHIEVED
- CMO: More sales assets. 3 new for each stage (TOFU, MOFU, BOFU) of the sales process => Miss. PMM hire/onboard.
- CMO: Presentation generation based on conversational development index. => Initial investigation of technology options.
- CMO: 3 minute video product demo targeted at buyer audience. Published => Not completed
- PMM: Sales Demo easy to give. At least one Account Executive can do it in 10 minutes.=> Not completed
- CRO: Sales Pitch Deck and Messaging trained. 75% of customer facing team passes test. ACHIEVED
- VP Prod: Conversational Development Index in product and SFDC. => Done
- VP Prod: Improve JIRA support. Better than Bitbucket. Support transitions, references and development panel. => Done: dev panel, references. Not possible yet: transitions.
- Technical writing: Create a page on our JIRA integration and how to move to Issues. 100% increased usage of JIRA integration. => Done. % TODO

Objective 2: Popular next generation product

- CEO: GitLab.com ready for mission-critical tasks. 99% of user requests have first byte < 1s
- VP Scaling: Highly Available. 99.9%
 - Production: Scalable and HA setup for cache and background jobs storage. Persistence of cache (Redis) and background job storage (Sidekiq) split and HA is set up for persistent storage. => Done.
 - Production: Robust backups with automatic periodic restores. Counter on dashboards.gitlab.com shows days since last automatic restore of database and file system. => Backups are tested but restores are not automatic, and counter is missing.
 - Production: Multi-canary deployments enabled. PoC of two canaries via kubernetes with Distribution team's Helm charts. => One canary available, not via kubernetes.

- Production: Enable GitLab.com search using Elasticsearch. Provide a cluster that is 99% available for > 2 months in Q3. => Infrastructure available but application requires further work.
- Database: Run pgbouncer from Omnibus. Shipped. => Done
- Database: Use Omnibus provided software for HA / failover instead of using a GitLab.com specific setup. Shipped and In Use. => Not yet in use.
- VP Eng: Enable High Availability
- Platform: Enable graceful degradation when file servers are down. => Circuitbreakers built but not yet fully deployed.
- Platform: Support multiple Redis clusters for persistent and cache stores. => Done.
- Platform: Use Geo DR to move between clouds and Area Zones => 10% done, begin testing with Geo testbed, pushing to Q4.
- Support: Improve GitLab provided debugging tools. We'll log all times we need terminal access and create an issue for each to improve GitLab tools. => 50% done, see <https://gitlab.com/gitlab-com/support/issues/754>
- Prometheus: Reach parity with Prometheus metrics for Unicorn, Sidekiq, and gitlab-shell. Deprecate InfluxDB. => 30% done with Prometheus metrics for Unicorn. Sidekiq and gitlab-shell remain for Q4.
- Edge: Make GitLab QA test backup/restore, LDAP, Container Registry, and Mattermost => 25% done with Mattermost. Backup/restore + LDAP targeted for Q4.
- CI/CD: Decrease monthly per minute cost of shared runners. Cost down by 30%. => 25% done, Google Compute Engine support implemented but not deployed.
- CI/CD: Improve responsiveness for pipelines without any existing builds. 99% start within 1s.
- VP Scaling: Lower latency. 99% of user requests have first byte < 1s
- Gitaly: Gitaly service active on file-servers. => Done.
- Gitaly: Roll out Gitaly migrations. 24 additional endpoints migrated to Gitaly and in acceptance testing. => Done.
- Gitaly: Reduce "idea to production" time of migrations. 80% of all migrations started in Q3 are enabled on GitLab.com within two GitLab releases. => Done by changing process.
- Database: Reduce the p99 of SQL timings across the board to 200 ms (100-200 ms less than what we have now). => Significant progress, but not all p99 < 200 ms. See [gitlab-ce34535](#) for more context.
- Production: Provide reliable internal & external baseline monitoring of overall service health. Implement baseline end to end monitoring for GitLab.com (api/web/git(ssh/https)) and define SLO's based on this baseline. => Baseline not defined.
- Production: Solve performance issues. Implement CDN for GitLab.com. => CDN was implemented but had to be rolled back.
- VP Eng: Lower latency in application
- Discussion: Solve performance issues. Reduce p95 of discussion-related actions with over 10 hits/day to < 1 s. Reduce p99 to < 3 s. => 25% done. 10% (~150 of 1476 endpoints) still exceed thresholds. <https://stats.gitlab.com/1902794> is at 1.2 s.
- Platform: Solve performance issues. Reduce p95 of platform-related actions with over 10 hits/day to < 1 s. Reduce p99 to < 3 s. 25% done. 10% (~150 of 1476 endpoints) still exceed thresholds.
- Frontend: Do a Manual performance audit and deploy improvements. Implement Top 3 actions => Done
- Frontend: Package optimization and CDN Hosting for .com => Done

- Edge: Ship large database seeder for developers. => 25% done
- Edge: Enable Bullet by default on the CI. => May not be practical given <https://gitlab.com/gitlab-org/gitlab-ce/issues/30129note35800389>
- VP Eng: Eliminate critical stability issues
 - Discussion: Merge requests get merged 100% without ever getting into stuck locked state => 20% done. Improvements made in <https://gitlab.com/gitlab-org/gitlab-ce/mergerequests/13207>, but still occurring. We need to measure rate.
- Platform: Project authorizations fast (< 1 s) and consistent (requires no manual refreshes). DONE
- Platform: Project imports and forks complete 100% without ever getting stuck => 10% done. Import issues need to be broken out separately.
- Platform: Namespace and project renames work 100% of the time => 30% done. Requires hashed storage support
- Platform: Repository cache state eventually consistent within minutes (no manual expiration needed) => In progress with <https://gitlab.com/gitlab-org/gitlab-ce/mergerequests/14980>
- Frontend: Implement cross-browser automated testing. Catch at least one browser regression before release date. => 95% Done Not automated. Has to be turned on.
- Frontend: Accomplish at least 3 frontend improvement issues per release cycle. => 60% Done
- CI/CD: Make runners work on Google Compute Engine without dying halfway. Done
- CI/CD: Track and ensure the number of job failures due to system failure. Number < 0.01%
- VP Scaling: Secure platform
 - Security: Improve defenses. Implement top 10 actions from Risk Assessment. => Implemented top 12.
 - Security: Vulnerability testing. Conduct external testing. => Done.
 - Security: Improve security practices through people and processes. Backlog of security issues reduced by 50% (currently 148 issues in gitlab-ce, 9 in gitlab-ee, 69 on infrastructure). => Not done, backlog is larger.
- CEO: Increased usage of idea to production. 100% growth
- CMO: Generate more company and product awareness including overtake Bitbucket in Google Trends. => Achieved.
- VP Prod: Issue boards usage increase. 100% (.com and usage ping) => ~100% increase
- VP Prod: Service Desk usage increase. 100% (.com and usage ping) => ~1000-2000% increase
- VP Prod: Subgroups usage increase. 100% (.com and usage ping)
- Head Prod: Monitoring usage increase. 100% (.com and usage ping)
 - Prometheus: Increase adoption. 3 issues shipped.
- Head Prod: Pipelines usage increase. 100% (.com and usage ping)
 - CI/CD: Improve onboarding. 10 issues shipped.
- Head Prod: Environments usage increase. 100% (.com and usage ping)
- Head Prod: Review apps usage increase. 100% (.com and usage ping)
- CTO: Make autodeploy more practical. Undefined yet.
- CTO: Zero-configuration CI. Make it work for 5 popular frameworks.
- CTO: Lead update to Rails 5.1. Shipped.
- CMO: Better explain our solution, features and value proposition. Flow page that links to feature pages. => New home page with new content for 7 steps launched.
- UX: Measure usability of critical user flows to identify areas needing improvement. => 75% Research done, issues opened. Implementation and re-measure not done.
 - Propose optimizations of critical user flows based on results found in research.
 - Implement optimization and re-measure to ensure the user experience has improved.

- Frontend: Optimize Frontend Code to support speed. For larger projects and ones with thousands of comments. => Done a lot of Performance related optimisations, ongoing task
- Distribution: Simplify HTTPS configuration. In Omnibus and Helm. For Rails app, container registry, and pages. Consider Let's Encrypt.
- Dir Partnerships: First team of OSS projects starts using GitLab for proof-of-concept/testing at Drupal, Gnome and Kubernetes.
- Dir Partnerships: Get major partner to use it for CI.
- Dir Partnerships: AWS QuickStart guide published
- CMO: More contributors from the wider community each month. Unique contributors grow 10% QoQ => Missed.
- CMO: Initiate AR engagement with key analysts and achieve Leader in the Forrester Wave CI research report. => Leader in Forrester CI Wave. Account engagement with Forrester/Gartner. Analyst day and multiple briefings.
- Head Prod: Make sure the installation process is great. Complete all items.
- VP Eng: Make sure existing features are used by GitLab.com. Cycle time analytics working.
- Support: GitLab.com uses Service Desk. In use for one process. => Set up for Security, but not in use. Legal is using Service Desk.
- Distribution: Deliver service specific Docker images and Helm charts usable for GitLab.com and our users. Rails, Workhorse, Pages, Registry, Gitaly, Shell.
- CI/CD: Help GitLab.com use our deployment features. CD/Kubernetes/Helm/Canary deploys/Review Apps/Service Desk
- Edge: GDK based on Kubernetes (e.g. minikube).
- Edge: CE Pipelines run in 30 minutes.
- Edge: Flaky tests don't break master.
- Edge: CE is automatically merged to EE daily.
- Edge: Triage policies are automatically enforced.
- UX: Make it easier to find and use advanced GitLab features => 25% Implemented <https://gitlab.com/gitlab-org/gitlab-ee/issues/3054> which was then reverted. Successfully implemented <https://gitlab.com/gitlab-org/gitlab-ee/issues/3049>
- Frontend: Support > 2000 comments/code changes on issues, diffs, and merge requests. With speed and without jank. => 60% done. Needs work on MR discussion.
- Frontend: Reduce frontend render times of comments and diffs with a goal of %70 faster. => <https://gitlab.com/gitlab-org/gitlab-ce/issues/4058> Speed Index Before 8081, Now 2757 (65% faster)
- Head Prod: Ensure all idea to production features are usable by GitLab.com
- CEO: Next generation product. Move vision forward every release.
- Head Prod: Create and radiate vision for GitLab DevOps. Publish vision video.
- Head Prod: Auto DevOps. Ship kernel of Auto DevOps.
- UX: Improved navigation. Iterate on it every month. => 100% All work completed, work continues
- UX: Improve perceived performance => 0% Issues continually re-scheduled (UX Ready)
- Frontend: Instant user feedback => Done
- Frontend: Repo as editor => PoC Done, move to Beta as a stable feature in progress
- Frontend: Get a moonshot (e.g. VSCode/Monaco editor, Team dashboards) from PoC to Feature - Image Comments released + Multi file editor Beta in 10.3

Objective 3: Great team

- CEO: Effective leadership. NPS of leaders as rated by their reports 10% improvement. Update: This will be pushed back until 360 reviews are established in 2018.

- VP PO: Management completed basic management (MGR) classes. 75% completed
- VP PO: 5 GitLab specific courses for the management team. Published and delivered.
- CFO: Real-time analytics platform. 40% of metrics live 50% ACHIEVED - 20% OF METRICS
- CEO: Attract great people. New hire score 2% improvement
- VP PO: More sourced recruiting. 20% of total hires
- VP PO: Hires from lower cost locations. 20% increase
- VP PO: Interviewer rating (ELO). Implemented
- VP PO: Quicker hiring cycle (plan interviews upfront). Hiring time -25%
- VP PO: Improve our global compensation framework. Structured on better data.
- CEO: Retain great people. eNPS 0% change
- Customer Success: Hiring talent to plan and ramping up SA's so that on day 31 they are certified and ready to be assigned to accounts.
- VP PO: Improve diversity. Two initiatives deployed.
- VP PO: Train all Senior level and above team members on how to interview.
- VP PO: Develop and document benefits plans (current and considering) for each entity.

SECTION: company/okrs/2017-q4.md

Objective 1: Grow incremental ACV according to plan

- CRO: Sales efficiency ratio > 1.2 (IACV / sales+marketing)
- CRO: Field Sales efficiency ratio > 1.8 (IACV / sales spend)
- CMO: Marketing efficien