1. HPT in Product Division
1. Keep around references
1. Product Accessibility
1. Secure Govern Database Decomposition
1. User journey map

Past Working Groups (alphabetic order)

1. Account Escalation Process
1. AI Integration
1. API Vision
1. Architecture Kickoff
1. Bounded Contexts
1. Category Leadership
1. China Service
1. CI Queue Time Stabilization
1. CI/CD Build Speed time-to-result
1. ClickHouse Datastore
1. Cloud Licensing
1. Commercial & Licensing
1. Consumption Add-Ons
1. Continuous Scanning
1. Contributor Growth
1. Cross Functional Prioritization
1. Dashboards
1. Database Scalability
1. Demo & Test Data
1. Development Metrics
1. DevSecOps Adoption
1. Digital SMB + Solutions Architects
1. Dogfood Plan
1. Ecommerce Motion
1. Emerging Talent
1. Engineering Career Matrices
1. Engineering Internship
1. Enterprise Market Leadership
1. Event Stream
1. Expense Management
1. Experimentation
1. Feature Testing
1. First Order
1. Frontend Observability
1. Frontend Vision
1. Githost Migration
1. GitLab Administration
1. GitLab Dedicated
1. gitlab-ui (CSS and Components)
1. GitLab.com Cost
1. GitLab.com Disaster Recovery

1. GitLab.com Revenue
1. GitLab.com SAAS Data Pipeline
1. Government Support Offerings (Internal only)
1. GTM Product Analytics
1. IACV - Delta ARR
1. IC Gearing
1. Improve Ops Quality
1. Internal Feature Flag usage
1. Internship Pilot
1. Isolation
1. Issue Prioritization Framework
1. Leading Organizations
1. Learning Experience
1. Learning Restructure
1. Licensing and Transactions Improvements
1. Lighthouse Metric Definitions
1. Log Aggregation
1. Logging
1. Maintainership
1. Major Releases
1. Merge Request Report Widgets
1. Minorities in Tech (MIT) Mentoring Program
1. MLOps
1. Modern Applications Go-To-Market
1. Multi-Large
1. Next Architecture Workflow
1. Object Storage
1. Performance Indicators
1. Pipeline Validation Service Operations
1. Product Analytics
1. Product Career Development Framework
1. Product Development Flow
1. Product Engagement Actions (FY21)
1. Project Matterhorn: Premium Price Tier Increase. Limited access
1. Purchasing Reliability
1. Python Stewardship
1. Rate Limit Architecture
1. Real-Time
1. Revenue Globalization
1. Runtime Update Process
1. SaaS Free User Efficiency
1. SaaS Reliability
1. Secure Offline Environment Working Group
1. Self-Managed Scalability
1. Sharding
1. Shift LAM to Primary Metric
1. Simplify Groups & Projects
1. Single Codebase
1. Software Supply Chain Security
1. SOX PMO

1. Talent Acquisition SSOT
1. TeamOps Sales and Marketing Group
1. Tiering
1. Token Management
1. Transient Bugs
1. Upgrade Improvements
1. Upstream Diversity
1. Usage Reporting
1. User Engagement
1. Vue.js 3 Upgrade
1. Webpack (Frontend build tooling)

SECTION: company/working-groups/account-escalation-process.md

Attributes

Property	Value
Date Created	2022-11-15
Target End Date	2023-03-31
Slack	wgaccountescalation (only accessible from within the company)
Google Doc	Working Group Agenda (only accessible from within the company)

Goal

The Account Escalation Process Working Group aims to improve the cross-functional collaboration when a customer is in an escalated state.

Overview

The account escalation process can involve stakeholders from CS, Support, Engineering as well as
Sales or other areas of the company. This is a keystone process that requires cross-functional attention to improve.

Goals

This is a non-comprehensive list of topics to be discussed.

- What is the single source of truth for knowing which customers are escalated?
- How can we make account escalations easier to declare? (see: Declaring an Incident)
- How can we better prepare CS and Support ICs to work within escalations
- Reduce the ambiguity in roles between CS and Support

Exit Criteria

1. Create an SSOT for knowing which customers are escalated.
1. Create training materials suitable for a future OKR for CS and Support
1. Refine process of declaring an account escalation so that anyone in the company can do it with minimal effort (MR: www-gitlab-com!114418
1. Refine roles and responsibilities of CS and Support within account escalations to reduce ambiguity
1. Develop a method of communication appropriate for e-group stakeholders, for example a "hot customer list" published in the e-group channel at some cadence. (gitlab-com/chief-of-staff-cto67)

Roles and Responsibilities

Working Group Role	Username	Person	Title
Executive Stakeholder	@spatching	Sherrod Patching	VP Customer Success Management
Facilitator	@lyle	Lyle Kozloff	Director of Global Support Readiness
Member	@christiaanconover	Christiaan Conover	Director of Customer Success Management (AMER)
Member	@manuel.kraft	Manuel Kraft	Customer Success Manager (EMEA)
Member	@ricardoamarilla	Ricardo Amarilla	Customer Success Manager (LATAM)
Member	@afappiano	Anthony Fappiano	Engineering Manager, Reliability

---
SECTION: company/working-groups/administration.md

Attributes

Property	Value
Date Created	2023-03-01
End Date	2023-07-27
Slack	wgadministration (only accessible from within the company)
Google Doc	Working Group Agenda
Issue Board	Administration Working Group
Epic	GitLab Administration Working Group
Overview & Status	Closed

Charter

The GitLab Administration Working Group will determine the work needed to maintain GitLab Administrator functions healthy, identify owning groups for requests, and redirect issues to corresponding groups for resolution. This group is a temporary owner for all GitLab Administrator functions until a team is funded to be a permanent owner.

Context

The maintenance and enhancement of GitLab Administration functions has been a shared responsibility among Development groups. This often triggers the same questions about shared responsibilities - who owns this and who should work on this. Before a development owning group can be identified, a working group will help us to maintain the related functions to run GitLab (the software) smoothly.

Exit Criteria

1. Owners of the items in the Admin Area of GitLab are identified, acknowledged, and documented in the handbook. => DONE The owner list is in the issue description, we (GitLab as a whole) will continue seeking clarification for the remaining unowned areas. Meanwhile, continue following Shared responsibility functionality for unowned areas.
1. Owners of things related to instance management (upgrade, backup and restore, metrics, performance) and configuration are identified, acknowledged, and documented in the handbook. => SKIPPED Didn't investigate due to other priorities.
1. A proposal for a new team in case there is a need for certain items of the GitLab Administrator features. There is a chance this is unnecessary if the above two items cover all features. => SKIPPED Didn't investigate due to other priorities.
1. A process of identifying Owners for future administration features is documented in the handbook. => DONE Updated handbook page with process of adding new admin section.

Roles and Responsibilities

Working Group Role	Person	Title
Executive Stakeholder	Christopher Lefelhocz (@clefelhocz1)	VP of Development
Facilitator/DRI	Peter Lu (@plu8)	Engineering Manager, Distribution:Deploy
Product Management DRI	Dilan Orrino (@dorrino)	Product Manager Distribution
Member	Gosia Ksionek (@mksionek)	Backend Engineer, Manage::Authentication & Authorization
Member	Christina Lohr (@lohrc)	Senior Product Manager, Data Stores::Tenant Scale
Member	Joseph Longo (@jlongogitlab)	Manager, Governance and Field Security
Member	Oriol Lluch Parellada (@o-lluch)	Engineering Manager, Infrastructure
Member	Hannah Sutor (@hsutor)	Senior Product Manager - Auth
Member	Falko Sieverding (@fsieverding)	Customer Success Manager, EMEA
Member	Shawn Sichak (@ssichak)	Senior Security Engineer, Trust and Safety
Member	Manuel Kraft (@manuel.kraft)	Customer Success

Manager, EMEA |
| Member | Sean Carroll (@seancarroll) | Engineering Manager, Create::Source Code |

SECTION: company/working-groups/ai-integration.md

Attributes

Property	Value
Date Created	2022-03-23
End Date	2022-11-13
Slack	CLOSED wgaiintegration - Slack channel for the working group and the high level alignment on getting AI ready for Production
Slack	aiintegrationdevlobby - Channel for all implementation related topics and discussions of actual AI features
Slack	gaiframework - Channel for the AI Framework Team which is building the base for all features (experimentation API, Abstraction Layer, Embeddings, etc.)
Slack	aistrategy - Discussion on strategic and business initiatives surrounding AI/ML at GitLab.
Slack	ai-infrastructure - Infrastructure/Platform support for AI integration. See also &969.
AI Architecture Documentation	Doc
Google Doc	Working Group Agenda
Feature Tracking	Sheet
YouTube playlist	Playlist on GitLab Unfiltered
Parent Epic	Parent epic
Epic/Issue Working Group label	wg-ai-integration issue board and wg-ai-integration epic search
Epic label for prioritized prototypes	wg-ai-integration-prioritized-prototype
Issue Board for AI Framework group	Issue board link using label group::ai framework
Overview & Status	See Exit Criteria below
Meeting schedule	Monday, Tuesday, and Wednesday at 8am Pacific and Thursdays at 1pm Pacific

Goal

The GitLab AI Integration Working Group aims to define, coordinate and ramp up integration of AI capabilities into all product areas

Overview

Corresponding Epic

We want to enable all product teams to be able to use advanced AI capabilities for improving and adding functionality to the product so users can be faster and more productive in their

DevSecOps lifecycle. For product teams we want to establish a clear and fast way for going from idea, experiment to production when using AI functionality. Incorporating the services, models and knowhow that the MLOps group has built over time and can provide to the wider team.

The working group will facilitate fast experimentation and prototyping of AI capabilities. We will also advise on what must be considered (and in some cases, get explicit approval on) before moving to production, including legal approval, ethical use of AI, potential necessary changes to terms of service, performance implications, hosting cost implications, infrastructure readiness, security readiness, licensing of 3rd party software/services, appropriate GitLab licensing levels for features, value add in helping users achieve their goals and needs, etc.

More information on the effort and plans can be found in the internal handbook.

Goals

This is a list of topics that we want to discuss:

- Classification of AI possibilities and complexities
- Ideation and understanding how experimentation can be done
- Rapid ML Prototyping
  - API's and Framework for experimentation
  - Feature Flagged Prototypes
  - Determine how teams without Product Design representation can progress with validation
- Groundwork
  - Base API's for all teams
  - Infrastructure
  - Documentation
  - Understanding of Jobs and user needs to be affected
- Existing ML features
- Road to production for features
  - Different gates

Exit Criteria

The following criteria should be met for the group to disband:

- Product teams have a clear method to build and integrate AI into GitLab product areas.
- The integration platform should have a product group handling maintanence and feature development
- We have a structured methodology for evaluating new AI models, adding them to the integration platform to allow them to be consumed by product teams.
- We have a roadmap plan to achieve GA for our initial AI experiments.
- Documented process for handling AI feature proposals as part of the prioritization framework
- Move SAFE content from the internal handbook to the public handbook where appropriate and SAFE.
- Develop an evaluation process of user experience between options to make more intelligent decisions on which engineering solution we recommend.

Q2 OKRs

Deliver X experimental, Y beta, and Z GA AI features

Roles and Responsibilities

Working Group Role	Username	Person	Title
Executive Stakeholder	@hbenson	Hillary Benson	Senior Director, Product Management - Sec, Data Science & Monitor
Executive Stakeholder	@timzallmann	Tim Zallmann	Senior Director of Engineering, Dev
Facilitator	@tmccaslin	Taylor McCaslin	Group Manager, Product - Data Science
Functional Lead - AI Model Validation	@mray	Monmayuri Ray	Engineering Manager AI Model Validation
Functional Lead - UX	@jmandell	Justin Mandell	Product Design Manager: Analytics, Govern, ModelOps, and Secure
Functional Lead - UX	@pedroms	Pedro Moreira da Silva	Staff Product Designer
Functional Lead - Legal	@mtaylor	Matthew Taylor	Sr. Director of Legal
Pricing representative	TBH	TBH	Principal Pricing Manager, Product
Product representative	@mushakov	Melissa Ushakov	Group Manager, Product - Plan
Product representative	@sarahwaldner	Sarah Waldner	Group Manager, Product - Create
Product representative	@abellucci	Alana Bellucci	Senior Product Manager, Govern:Threat Insights
Product representative	@joshlambert	Joshua Lambert	Director of Product, Enablement
Product representative	@tlinz	Torsten Linz	PM, Source Code
Development representative	@johnhope	John Hope	SEM, Plan
Development representative	@andr3	André Luís	FEM: Source Code
Development representative	@cdu1	Chun Du	Director of Engineering, Enablement
Development representative	@igor.drozdov	Igor Drozdov	Staff Backend Engineer, Source Code
Development representative	@jeromezng	Jerome Ng	Director of Engineering, Fulfillment
Development representative	@pcalder	Phil Calder	Senior Engineering Manager, Anti-abuse, Govern, and Growth
Development representative	@nmccorrison	Neil McCorrison	Engineering Manager, Govern: Threat Insights
Development representative	@carlad-gl	Carla Drago	Senior Backend Engineer, Manage: Import & Integrate
Development representative	@donaldcook	Donald Cook	EM, Project Management
Legal representative	@jbackerman	Jesse Backerman	Managing Legal Counsel
Vulnerability Research Representative	@idawson	Isaac Dawson	Staff Vulnerability

Researcher |
Vulnerability Research Representative	@dbolkensteyn	Dinesh Bolkensteyn	Sr. Vulnerability Researcher
Third Party Security Risk Representative	@tdilbeck	Ty Dilbeck	Security Risk Manager
Governance and Field Security Representative	@jlongogitlab	Joseph Longo	Governance and Field Security Manager
Security Compliance Representative	@kbray	Ken Bray	Sr. Security Compliance Engineer (Dedicated Markets)
Security Compliance Representative	@lcoleman	Liz Coleman	Security Compliance Manager (Commercial)
Security Automation Representative	@agroleau	Alexander Groleau	Senior Security Engineering Manager (Automation)
Security Automation Representative	@imand3r	Ian Anderson	Staff Security Engineer (Automation)
Application Security Representative	@greg	Greg Myers	Security Engineer (Application Security)
Solutions Architecture Representative / Rapid Prototyping Team Member	@bartzhang	Bart Zhang	Channel Solutions Architect
Product Marketing Representative	@laurenaalves	Laurena Alves	Senior Product Marketing Manager
Developer Relations Representative	@johncoghlan	John Coghlan	Director, Developer Advocacy
Privacy Representative	@emccrann	Eugene McCrann	Lead Legal Counsel, Privacy
Quality Engineering Representative	@at.ramya	Ramya Authappan	Engineering Manager, Quality, Dev & Analytics Section
Infrastructure	@lmcandrew	Liam McAndrew	Engineering Manager, Scalability Frameworks
Infrastructure	@igorwwwwwwwwwwwwwwwwwwww	Igor Wiedler	Staff SRE, Scalability Frameworks
Infrastructure	@mbursi	Michele Bursi	Engineering Manager, Delivery System
Support	@ralfaro	Ronnie Alphero	Support Engineering Manager
Enablement	@cs.wang	Christopher Wang	Sr. Manager, Enablement (Sales Development)

Engineering Groups

We currently have two core AI Development groups at GitLab: AI Framework group and AI Model Validation group.

AI Model Validation group

The AI Model Validation group helps all product groups to match the right model(s) and AI/ML-based techniques to the user problem they must solve. They do that by evaluating, building, training, and tuning many of the models GitLab uses as well as by proactively sharing AI resources and experience. Today, they also directly build and maintain some user-facing AI features.

- AI Model Validation group develops in-house AI features that run natively within GitLab infrastructure. The group develops these models to meet capability, quality, customizability, privacy, and cost requirements. These components include an inference engine, abstraction layer, and models.
- The custom-built models help in use cases when there is a need to train models on customer-

proprietary data (like all merge requests and commits for a customer) and when 3rd party models do meet our needs.
- Their currently released features are code suggestions which is currently in customer beta and suggested reviewer which is GA.
- This group can help evaluate models for functional correctness and model perplexity based on metrics and (often large) benchmark datasets, which is a more statistical evaluation than manual testing. This can help determine the most quality model for a feature's use-case.
- They work in many languages, including Ruby on Rails, Golang, Python (for machine learning and data science), and Typescript (for the VS Code Plugin). The group comprises of ML Scientistists, MLOps Engineers , ML Infrastructure Engineers, and Fullstack engineers.

You can contact this group via Slack in gaimodelvalidation. View their issue board here. To see who is engaged on this effort please see here.

AI Framework

AI Framework exposes AI services and the underlying models (third party or native GitLab models) to all product groups.

- The AI Framework group enables the rest of the development department to build AI features through the abstraction layer.
- The abstraction layer supports OpenAI and is being extended to support equivalent Google AI functionality. Other commercial, open-source, and GitLab custom-built models are also being considered.
- This group empowers other groups to evaluate models via manual human testing, through the Experimentation API.
- This group works with Ruby on Rails as they make it easy for the GitLab product to add AI functionality through the gitlab/gitlab-org repo.

You can contact this group via Slack in gaiframework. View their issue board here.

AI Engineering Allocation

Because of the dynamic nature of the AI work and folks to be engaged, we are putting the AI work under an Engineering Allocation.   This means that assignments may change rapidly as focus and priorities shift. Current focus is on Code Suggestions adoption.

Name	Role	Area of Work
Alexandru Croitor	Senior Backend Engineer	AI Enablement
Eulyeon Ko	Backend Engineer	AI Enablement
Nicolas Dular	Senior Backend Engineer	AI Enablement
Denys Mishunov	Staff Frontend Engineer	AI Enablement
Jan Provaznik	Staff Backend Engineer	AI Enablement
Miko·aj Wawrzyniak	Staff Backend Engineer	AI Enablement
Pavel Shutsin	Senior Backend Engineer	AI Enablement
Max Woolf	Staff Backend Engineer	AI Enablement
Tan Le	Senior Fullstack Engineer	AI Enablement
Andras Herczeg	Backend Engineer	AI Enablement
Sebastian Rehm	Engineering Manager	AI Enablement

Daniel Tian	Senior Frontend Engineer	Threat Insights
Gregory Havenga	Backend Engineer	Threat Insights
Kerri Miller	Staff Backend Engineer	Code Review
Stanislav Lashmanov	Senior Frontend Engineer	Code Review
Simon Knox	Senior Frontend Engineer	Plan:Project Management
Nikola Milojevic	Senior Backend Engineer	Application Performance
Aleksei Lipniagov	Senior Backend Engineer	Application Performance
Matthias Käppler	Staff Backend Engineer	Application Performance
Roy Zwambag	Backend Engineer	Application Performance
Paul Phillips	Engineering Manager	Application Performance
Igor Drozdov	Staff Backend Engineer	Source Code
Patrick Cyiza	Backend Engineer	Source Code
Natalia Radina	Frontend Engineer	Source Code
Alper Akgun	Staff Fullstack Engineer	VS Code Extension
Tomas Vik	Staff Fullstack Engineer	VS Code Extension
Enrique	Staff Frontend Engineer	VS Code Extension
André Luís	Engineering Manager	Editor Extensions
Mike Eddington	Staff Backend Engineer	Editor Extensions (Visual Studio)
Ross Fuhrman	Senior Backend Engineer	Editor Extensions (Visual Studio)
Gabriel Mazetto	Senior Backend Engineer	Editor Extensions (JetBrains)
Naman Gala	Senior Backend Engineer	Editor Extensions (JetBrains)
Marshall Cottrell	Principal	Editor Extensions (Code Suggestions/VS Code)
Illya Klymov	Senior Frontend Engineer	Editor Extensions (Code Suggestions/VS Code)
Lena Horal-Koretska	Senior Frontend Engineer	Editor Extensions (Language Server)
Erran Carey	Staff Incubation Engineer	Editor Extensions (Neovim)
Ash McKenzie	Staff Backend Engineer	Editor Extensions (Neovim)
Jay Swain	Engineering Manager	Model Evaluation
Andrei Zubov	Senior Frontend Engineer	Deploy:Environments, Model Evaluation
Allison Browne	Senior Backend Engineer	Model Evaluation, Verify:Pipeline Execution
Dylan Bernardi	Associate Backend Engineer	Model Evaluation
Stephan Rayner	Senior Backend Engineer	Model Evaluation
Igor Wiedler	Staff Site Reliability Engineer	Model Evaluation
Alejandro Pineda	Staff Site Reliability Engineer	Model Evaluation

---
SECTION: company/working-groups/ai-security.md

Attributes

Property	Value
Date Created	June 12, 2025
End Date	TBD
Slack	wgaisecurity (internal)
Google Doc	Working Group Agenda (internal)
Epic	Main Project Epic (internal)
Handbook Page	GitLab's FedRAMP Authorization Program

Exit Criteria

1. Best practices for implementing AI prompts are documented in our Contributor documentation.
1. Proof-of-concepts are executed and recorded to understand what AI security tooling could offer SaaS, Dedicated, and Self-Managed customers.
1. Our CI/CD pipeline will trigger code review of merge requests and provide actionable advice for contributors.
1. Our CI/CD pipeline will block merge requests that do not meet secure development standards we establish for our AI offerings.

Roles and Responsibilities

Working Group Role	Team Member Name	Role
Executive Sponsor	Jamie Dicken	Director, Security Platforms and Architecture
Executive Sponsor	Julie Davila	VP, Product Security
Executive Sponsor	Tim Zallmann	VP, AI Engineering
Functional Lead	Erran Carey	Staff Fullstack Engineer
Functional Lead	Jessie Young	Principal Engineer
Functional Lead	Joern Schneeweisz	Principal Security Engineer
Member	Daniel Hauenstein	Application Security Engineer, Product Security
Member	Vitor Meireles De Sousa	Senior Manager, AppSec, Product Security

---

SECTION: company/working-groups/api-vision.md

Attributes

Property	Value
Date Created	2022-02-07
End Date	2023-11-17
Slack	wgapivision (only accessible from within the company)
Google Doc	Working Group Agenda (only accessible from within the company)
Issue Board	Issue Board
Overview & Status	See Exit Criteria below

Goal

The GitLab API Vision Working Group aims to improve the current APIs and define their future evolution.

Overview

We don't have a cohesive view between the REST and GraphQL APIs. We specify that the GraphQL API is the primary means of interacting programmatically with GitLab, but we often don't follow this criteria. Both APIs cover a different set of features, but none is feature-complete.

Goals

This is a list of topics that we want to discuss:

- Responsibilities, Directly Responsible Individual, and technical experts. At the moment, the Manage:Integrations group is the DRI of the APIs but there is also the @graphql-experts group.
- General vision of the GitLab API:
  - REST / GraphQL API consistency.
  - REST first vs. GraphQL first vs. another approach.
- Review APIs:
  - General architecture.
  - Permissions and scopes.
  - Feature coverage.
  - Performance.
- Testing:
  - Coverage.
  - Automated testing.
  - Tools (eg. Postman collections).
- API deprecations lifecycle and strategy:
  - REST v5 API or further iterations.
  - GraphQL deprecation process.
- API standards, including OpenAPI specification.
- Documentation:
  - Improve current documentation.
  - Review the first-time API user experience.
  - Automation of the documentation.
  - Full catalog of all public and internal APIs.
- Learning and contributions:
  - Review contributors' documentation.
  - Create learning paths for team members, especially about GraphQL.

Exit Criteria

The table below lists all exit criteria for the working group. This is the top-level epic.

	Completed Date	Progress	DRI	Criteria

1	TBD	10%	@g.hickman	Define the vision of the GitLab API for the future years
2	TBD	0%	@mgill	Set the foundation of a cohesive development strategy going forward
3	TBD	0%		Capture work needed for next generation API
4	TBD	15%	@.luke	API Deprecation and Lifecycle policies
5	TBD	15%	TBD	Create a concept and roadmap to automatically generate API documentation
6	TBD	5%	TBD	Define minimum levels of performance and stability, with appropriate checks and monitoring
7	TBD			Robust Open API

Roles and Responsibilities

Working Group Role	Username	Person	Title
Executive Stakeholder	@timzallmann	Tim Zallmann	Senior Director of Engineering, Dev
Facilitator	@arturoherrero	Arturo Herrero	Engineering Manager, Manage:Integrations
Facilitator	@g.hickman	Grant Hickman	Senior Product Manager, Manage:Integrations
Functional Lead	@.luke	Luke Duncalfe	Senior Backend Engineer, Manage:Integrations
Functional Lead	@axil	Achilleas Pipinellis	Senior Technical Writer, Enablement
Functional Lead	@Andysoiron	Andy Soiron	Senior Backend Engineer, Manage:Integrations
Member	@grzesiek	Grzegorz Bizon	Principal Engineer, Verify
Member	@fcaplette	Frédéric Caplette	Senior Frontend Engineer, Verify:Pipeline Authoring
Member	@bmarjanovic	Bojan Marjanovic	Senior Backend Engineer, Manage:Integrations
Member	@kerrizor	Kerri Miller	Senior Backend Engineer, Create:Code Review
Member	@lauraX	Laura Montemayor	Backend Engineer, Verify:Pipeline Authoring
Member	@nagyv-gitlab	Viktor Nagy	Senior Product Manager, Configure
Member	@kpaizee	Kati Paizee	Senior Technical Writer, Growth and Ecosystem
Member	@fabiopitino	Fabio Pitino	Staff Backend Engineer, Verify:Pipeline Execution

Member	@dstull	Doug Stull	Staff Fullstack Engineer, Growth:Expansion
Member	@ntepluhina	Natalia Tepluhina	Staff Frontend Engineer, Plan:Project Management
Member	@avielle	Avielle Wolfe	Backend Engineer, Verify:Pipeline Authoring

---

SECTION: company/working-groups/architecture-kickoff.md

Attributes

Property	Value
Date Created	August 10, 2020
End Date	October 27, 2020
Slack	wgarchitecture-kickoff (only accessible from within the company)
Google Doc	Working Group Agenda (only accessible from within the company)
Issue Board	TBD

Exit Criteria

The charter of this working group is to produce the first iteration of a 12 month technical roadmap, due on September 1, 2020.

- 12 month artifact in handbook
- 6 month artifact(s) in handbook
- 3 month artifact(s) in handbook
- FY21-Q4 OKR proposals

Roles and Responsibilities

Working Group Role	Person	Title
Executive Stakeholder	Eric Johnson	Chief Technology Officer
Facilitator	Gerardo "Gerir" Lopez-Fernandez	Engineering Fellow, Infrastructure
Functional Lead	Stan Hu	Engineering Fellow, Development
Functional Lead	Andrew Thomas	Principal Product Manager, Enablement
Functional Lead	Andrew Newdigate	Distinguished Engineer, Infrastructure
Functional Lead	Kamil Trzci·ski	Distinguished Engineer, Ops and Enablement
Functional Lead	Philippe Lafoucrière	Distinguished Engineer, Threat Management, Secure, Protect

Member	Antony Saba	Senior Security Engineer
Member	John Jarvis	Staff Site Reliability Engineer
Member	Grzegorz Bizon	Staff Backend Engineer
Member	Craig Gomes	Backend Engineering Manager

SECTION: company/working-groups/automotive-development.md

Attributes

Property	Value
Date Created	2023-04-29
Target End Date	TBD
Slack	wgautomotivedevelopment (only accessible from within the company)
Google Doc	Working Group Agenda (only accessible from within the company)

Goal

The Automotive Development Working Group is a cross-functional group focused on gathering requirements and providing product support for automotive development use cases.

Overview

With the increasing focus on software-defined cars, traditional automotive OEMs are shifting away from contracted software and towards in-house development. Iteration, efficiency and compliance have to be squared with complex development workflows using a wide array of proprietary development tools. GitLab with its built in versioning, CI/CD, security and compliance features is poised to become the leading platform for automotive development.

Automotive software development is characterized by a complex combination of:

1. embedded development close to the kernel
1. using a wide array of proprietary development tools (Autosar, Matlab/ model driven development) whose outputs need to be orchestrated, integrated and automated
1. based on huge repositories with a large number of artifact-like binary files
1. thus very long pipeline runtimes (due to repo checkout and artifact handover between stages) measured in terms of hours
1. complex QA requirements with deployments into custom hardware test clusters
1. compliance requirements regulated by ASPICE including long-term versioning and pipeline reproducibility requirements measured in years to decades ("Baselining") as well as traceability of changes and dependency / license management (SBOM)
1. Strong focus on application lifecycle management

Goals

- Establish GitLab as the leading platform for automotive software development
  - Capture requirements, use cases and product gaps in collaboration with major automotive customers
  - Close product gaps and develop blueprints for automotive development workflows
  - Obtain TISAX certification
  - Develop solutions for ASPICE compliance
- Enable Sales and CS to drive customer value and growth based on the developed solutions
  - Provide Sales/SA/CSM enablements on automotive development use cases
  - Develop a Sales strategy for automotive customers

Exit Criteria

1. Top 10 product gaps for automotive customers identified and go-forward plans for each defined
1. Established Automotive Customer Advisory Board / Automotive Round Table
1. Solution blueprints successfully deployed in 5 major automotive customers

Roles and Responsibilities

Working Group Role	Username	Person	Title
Executive Stakeholder	@david	David DeSanto	Chief Product Officer
Functional Lead Compliance	@hbenson	Hillary Benson	Senior Director of Product Management, Sec, Monitor, & Data Science
Functional Lead Verify	@jreporter	Jackie Porter	Director of Product Management, Verify & Package
Functional Lead Alliances	@DarwinJS	Darwin Sanoy	Principal Solutions Architect, GitLab Alliances
Working Group DRI	@mbruemmer	Martin Brümmer	Customer Success Engineer
Facilitator	@bmarnane	Bartek Marnane	VP of Incubation Engineering at GitLab
Facilitator	@fsieverding	Falko Sieverding	Customer Success Manager
Facilitator	@jesswanggitlab	Jess Wang	Customer Success Manager
Facilitator	@vchacon	Vianney Chacon Madriz	Customer Success Manager
Facilitator	@jgaetjens	Julia Gätjens	Solutions Architect
Facilitator	@svij	Sujeevan Vijayakumaran	Solutions Architect
Facilitator	@Heddy2022	Hedayat Paktyan	Strategic Account Leader

---
SECTION: company/working-groups/bounded-contexts.md

Attributes

Property	Value
Date Created	2023-11-28

Target End Date	2024-03-31
Actual End Date	2024-05-20
Slack	wgboundedcontexts (only accessible from within the company)
Google Doc	Agenda doc (only accessible from within the company)

Summary

Today code is namespaced using Ruby namespaces but we don't have explicit rules or guidelines on how to organize code.
New namespaces are constantly being created and often there is no need to since new concepts could be nested inside
existing namespaces to better represent bounded contexts.
Having consistently namespaced code is the pre-requisite for a modularized codebase.

With this working group we want to:

- Identify the bounded contexts that compose the GitLab Rails codebase
- Stop the proliferation of new top-level namespaces without explicit reason.
- Have modules with rich domain logic that are aligned to the engineering organization.

Context

The proposal to split GitLab Rails monolith into a modular monolith
describes in details why modularization is important for the current stage of the GitLab codebase.
In order to make progress towards a modularized monolith we first need to understand what are the modules and bounded-context
that are present today in GitLab Rails monolith.

Activity

Members of the working group will:

- Classify the files related to their bounded context.
- Participate to sync meetings, bring questions and communicate progress.
- Provide feedback on the process and the guidelines to identify bounded contexts.
- Act as modularization evangelist for their respective stage/domain.
- Help review the final list of bounded contexts to ensure it's comprehensive.

Communication

- The working group would meet weekly initially alternating between a APAC and AMER friendly timezone.
  Then we can make it bi-weekly.
- A wgboundedcontexts slack channel would be created for anyone to ask questions and communicate async.
- After the initial meeting we expect the work to be carried out asynchronously.
  Each team member can work off their own list of files to categorize.
- Working on a spreadsheet (or alternatives) will make it easy to track progress on categorized files and remaining work to do.

- DRI to facilitate the meetings and communicate weekly status updates will be @fabiopitino and anyone else who wants to contribute.
- Eventual issues and follow-up ideas would be created using ~modularmonolith label.

Exit Criteria

The working group can be disbanded when:

Exit criteria	Resulting artifact
We have a published list of identified bounded contexts that acts as a documentation for developers.	config/boundedcontexts.yml
We have a simple process for creating, deleting and renaming bounded contexts, to allow the codebase evolve over time.	Docs updated
Bonus: Create a Rubocop Cop to enforce top-level namespaces using the list of identified bounded contexts.	Rubocop static analyzer added

Details

The working group would map the Ruby codebase, in particular the domain code into app/ and lib/ folder and come up
with a list of bounded contexts. This may exclude app/controllers and app/views initially since we want to focus primarily
on the domain layer (the core) of the Hexagonal Architecture.

During the process of mapping the codebase we would need to do:

- List all the Ruby files in a spreadsheet and categorize them into components following the defined guidelines.
  Files under the same directory can generally be categorized in bulk, so this should not take a lot of time.
- Inside the lib directory, distinguish between generic code that should be extracted as gem and domain code that should be namespaced
  the same as code in app. All code related to the same domain should have the exact same namespace. For example, we should not have
  Ci:: (from app) and Gitlab::Ci:: (from lib).
  - We will have a special platform category for code that should be
    extracted as gems and cross-cutting concerns in app and lib (generic base classes, utility, loggers, framework code, etc.).
- It's important to focus only on the top-level bounded contexts without getting carried away with renaming nested classes.
  This can be a refactoring activity at the discretion of the team owning this feature, but not in scope of the working group.
- Identify similar namespaces and the merging strategy. It could be any of the below:
  - Create a new namespace that represents a larger domain and move the classes in there.
  - Identify a representative namespace for the bounded context and move all other classes in there.
- Distinguish ambiguous namespaces by putting them into separate bounded context.
  When 2 namespaces seem similar but they operate in completely different contexts, move them under the respective bounded context.

- Identify namespaces that have multiple responsibilities and should be split. In this case classes could go into other namespaces
  that better represent their domain.
- Try to reuse existing top-level namespaces whenever possible.
- Ignore everything else (non Ruby code, config, specs, controllers, views, rake tasks, Grape API, GraphQL).

The list of bounded-context doesn't have to be perfect. It should be a good enough representation of the domains at play in the GitLab Rails
codebase where each class has only 1 namespaces where it can exist and there is no ambiguity.
We will iterate on this list and will need to have a process in place for creating, deleting, renaming bounded contexts.

Roles and Responsibilities

The Working Group should be composed of representative of all the various GitLab stages.
Each representative would be in a functional lead role representing their domain as a backend expert for their given GitLab DevOps stage.

Participants would ideally be interested in driving results related to modularization and Domain-Driven Design.
We will communicate the creation of the working group in development, backend and engineering-fyi slack channels.
For stages still missing participants we will seek support from Engineering leaders who work in the given DevOps stage.

Working Group Role	Person	Title
Executive Sponsor	Sam Goldstein	Director of Engineering, Ops
Facilitator	Fabio Pitino	Principal Engineer, Verify
Member	Chad Woolley	Senior Backend Engineer, Create::Remote Development (Workspaces)
Member	Thong Kuah	Principal Engineer, Data Stores
Member	Lucas Charles	Principal Engineer, Secure & Govern
Member	Furkan Ayhan	Senior Backend Engineer, Verify:Pipeline Authoring
Member	Smriti Garg	Senior Backend Engineer, Govern:Authentication
Member	Aboobacker MK	Senior Backend Engineer, Govern:Authentication
Member	Sean Carroll	Engineering Manager, Create:Source Code
Member	Ahmed Hemdan	Senior Backend Engineer, Secure:Static Analysis
Member	Ethan Urie	Staff Backend Engineer, Secure:Secret Detection
Member	Gabriel Mazetto	Senior Backend Engineer, Systems:Geo
Member	Vitali Tatarintev	Senior Backend Engineer, Create:Code Creation
Member	Sashi Kumar Kumaresan	Staff Backend Engineer, Govern:Security Policies

Member	Vasilii Iakliushin	Staff Backend Engineer, Create::Source Code
Member	Remy Coutable	Principal Engineer, Quality
Member	Eduardo Bonet	Staff Incubation Engineer, MLOps
Member	Vij Hawoldar	Senior Backend Engineer, Fulfillment::Utilization
Member	Suraj Tripathi	Senior Backend Engineer, Fulfillment::Utilization
Member	Kassio Borges	Staff Backend Engineer, Plan::Knowledge
Member	Aakriti Gupta	Senior Backend Engineer, Systems:Geo
Member	Peter Leitzen	Staff Backend Engineer, Engineering Productivity
Member	Abhinaba Ghosh	Engineering Manager, Test & Tools Infrastructure
Member	Mohamed Hamda	Backend Engineer, Fulfillment::Provision
Member	Andrei Zubov	Senior Frontend Engineer, Environments
Member	Fabien Catteau	Staff Backend Engineer, Secure::Composition Analysis

---
SECTION: company/working-groups/category-leadership.md

Attributes

Property	Value
Date Created	2021-07-20
Estimated Date Ended	2021-10-31
Slack	wg-category-leadership (only accessible from within the company)
Google Doc	Category Leadership Working Group - Agenda

Business Goal

Solidify GitLab's 'The DevOps Platform' messaging, including messaging for:

1. One sentence
1. One pager
1. Golden pitch
1. Website
1. 'Airport banner' / Brand agency
1. Category leadership
1. Commit keynote
1. Key external audience messaging

Exit Criteria

Achieve the business goals. Being a category leader is a multi-year effort and the focus of marketing, outside the working group.

Roles and Responsibilities

Working Group Role	Person
Facilitator	Christine Lee
Functional Lead(s)	Craig Mestel
Member	Harsh Jawharkar, Melissa Smolensky, Stella Treas, Mike Pyle, Monica Galleto
Executive Stakeholder	Sid Sijbrandij

SECTION: company/working-groups/china-service.md

Attributes

Property	Value
Date Created	December 17, 2020
End Date	April 30, 2021
Slack	jihu-working-group (only accessible from within the company)
Google Doc	Working Group Agenda (only accessible from within the company)
Issue Board	https://gitlab.com/gitlab-jh/gitlab-jh-enablement/-/issues

Exit Criteria

The charter of this working group is to operationalize a new plan for servicing China, due on April 30, 2021. Ongoing plans for engagement should be clearly identified, agreed upon, and documented.

- JiHu FAQ: https://about.gitlab.com/pricing/faq-jihu/
- GitLab Blog on JiHu: https://about.gitlab.com/blog/2021/03/18/gitlab-licensed-technology-to-new-independent-chinese-company/

Roles and Responsibilities

GitLab Inc JiHu Enablement Team

The GitLab Inc JiHu Enablement Team includes team members who are directly enabling and supporting JiHu.

Function	Primary Contact	Primary Contact GitLab Handle	Supporting Contact Name	Supporting Contact GitLab Handle
Exec Stakeholder	Sid Sijbrandij	@sytses		
General	Stella Treas, Jerome Ng	@streas, @jeromezng		

| Global Communications | Lisa Boughner | @lboughner |
| |
| Strategic Marketing | Melissa Smolensky | @melsmo
| Nout Boctor-Smith | @nbsmith |
| Brand | Emily Kyle, Luke Babb | @emily
@luke | |
|
| Sales - APAC | Anthony McMahon |
@amcmahon88 | |
|
| Sales - Solutions Architecture | Cherry Han |
@cherryhan | |
|
| Sales - Enablement / Education | Monica Jacobs | @monicaj
| David Somers | @dcsomers |
| Resellers | Michelle Hodges | @mwhodges
| Ed Cepulis, Amelia Seow, Kim Jaeger, Colleen Farris | @ecepulis, @aseow, @kjaeger, @cfarris |
| Field Operations | James Harrison, Tanuja Paruchuri | @dhong,
@jamesharrison, @tparuchuri | |
|
| Professional Services | David Sakamoto |
@dsakamoto | Michael Lutz, Ryan O'Nell, Mike Pyle
| @michaellutz, @ronell, @mpyle |
| Product | Kenny Johnston |
@kencjohnston | Justin Farris |
@justinfarris |
| Engineering - Development / Quality | Mek Stittri | @meks
| Lin Jen-Shin, Tanya Pazitny, Grant Young | @godfat-gitlab, @tpazitny, @grantyoung |
| Engineering - Infrastructure / Delivery | Steve Loyd | @sloyd
| Marin Jankovski, Amy Phillips | @marin, @amyphillips |
| Engineering - Distribution | Chun Du | @cdu1
| Steven Wilson, DJ Mountney, Balasankar 'Balu' C | @mendeni @twk3, @balasankarc
|
| Engineering - Fulfillment | Chase Southard |
@csouthard | Krasimir Angelov | @krasio
|
| Engineering - Security | Johnathan Hunt |
@JohnathanHunt | Steve Manzuik, Dominic Couture | @smanzuik,
@dcouture |
| Customer Support | Shaun McCann |
@shaunmccann | Jason Colyer, Anton Smith, Mike Lockhart, Alvin
Gounder, Harsh Chouraria, Emily Chang, Arihant Godha, Alexander Tanayno
| @jcolyer, @anton, @mlockhart, @alvin, @hchouraria, @emchang, @arihantar, @atanayno
|
| Finance | Craig Mestel | @cmestel
| | |
| Accounting | Dale Brown | @daleb04
| Sarah McCauley | @smccauley |
| Legal | Rashmi Chachra | @rchachra
| Rob Nalen, Emily Plotkin, Lynsey Sayers, Cathy Hilling, Julie Braughler, Miguel Silva, Robyn

Hartough | @rnalen, @emilyplotkin, @lasayers, @chilling32, @jbraughler, @mdsilva, @rhartough |

GitLab Inc Other Informed Team Members

The Other Informed Team Members lists people who have access to confidential information about JiHu that are not directly involved in enabling the JiHu team.

Function	Primary Contact	Primary Contact GitLab Handle	Supporting Contact Name	Supporting Contact GitLab Handle
Executive Group	Sid Sijbrandij Eric Johnson, Michael McBride, Todd Barr, Robin Schulman, Brian Robins, Wendy Barnes	@sytses, @edjdev, @mmcb, @tbarr, @rschulman, @brobins, @wendybarnes	Cheri Holmes, Robyn Hartough, Katie Gammon	@cheriholmes, @rhartough, @kgammon
Internal Comms	Simon Liang	@sliang2		
IT	Bryan Wise	@bryanwise	Marc DiSabatino, Peter Kaldis	@marcdisabatino-admin, @pkaldis
Engineering - Development	Christopher Lefelhocz	@clefelhocz1		
Product - Distribution	Josh Lambert	@joshlambert		
Investor Relations Messaging	Tony Righetti	@trighetti		
Community Messaging	John Coghlan	@johncoghlan	Michael Friedrich	@dnsmichi
Sales - Channel Sales	Alan Geary	@ageary		
Market Research and Customer Insights	Traci Robinson-Williams	@traci		

SECTION: company/working-groups/ci-build-speed.md

Attributes

Property	Value
Date Created	2023-12-01
Estimated Date Ended	2024-03-31
Slack	(only accessible from within the company)
Google Doc	CI Build Speed Working Group Agenda (only accessible from within the company)

Overview

We aim to create a working group to establish a repeatable process and framework to measure CI build speed and performance (time-to-result). The objective is to communicate to the market and customers a GitLab point of view, position GitLab CI build performance as the market leader, and provide guidance to customers on optimization considerations so that they can maximize developer efficiency while balancing compute costs for CI builds.

This group is primarily focused on measuring CI build speed, comparing GitLab-hosted runners to 3rd party solutions and other platform-hosted runners such as GitHub-hosted runners, Circle CI.

Context:

CI build speed and performance (time-to-result) and CI build cost efficiency are essential competitive vectors especially given the improved maturity of the CI/CD solutions in the market. The brand Q4 FY23 qualitative research study data indicates that "GitLab leads the pack in associations with speed." However, our internal benchmark testing, (slides, report, of CI build performance on GitLab SaaS had mixed results. Therefore GitLab SaaS customers' perception of CI build performance could be different than self-managed customers or even GitLab SaaS customers that choose to manage their own CI build environment.

Competitors like Github and Harness.io have generally focused on build time duration when discussing CI build speed and performance. While build time duration is the default performance measure, the working group must evaluate if there are other discrete measurements to include in the framework.

Additionally, new players are entering the market advocating for up to 5x GitLab SaaS runner build speed such as puzl.cloud and actuated.dev, shifting the perception of GitLabs positioning as a leader in CI build speed.

Out-of-scope

The visibility & observability of build speed for customer jobs will not be scope of this working group, as it's part ot CI Insights from the Fleet Visibility team.

Business Outcomes

1. Establish a repeatable process and framework to measure CI build speed and performance. The process must include the frequency of measurements (quarterly, semi-annual, annual).
1. Conduct an in-depth analysis of competitors product features that are focused on improving CI build speed and efficiency.
1. Conduct a technical analysis on CI build speed and performance improvements.
1. Recommendation to product leadership regarding new investment or prioritization of features to improve GitLab CI competitiveness specific to build speed and performance.
1. Establish comprehensive material for customers on how to improve CI build speed and cost efficiency on GitLab.
1. Review and agreement by legal on how we can communicate CI build speed data externally in blog posts or other marketing materials.

Topic	DRI

CI benchmarking framework - Design doc	@grzesiek
CI benchmarking framework - Implementation	tbd
Competitor analysis	tbd
Technical improvements	tbd
Product recommendation	@gabrielengelgl
Customer guide	tbd
Communication review	@gabrielengelgl

Exit Criteria

1. CI build speed benchmark process codified in the handbook
1. Internal charts for continuous monitoring of CI build speed
1. Initial blog post published on how to improve GitLab CI build speed

Roles and Responsibilities

Working Group Role	Person
Executive sponsor	Mike Flouton @mflouton
Facilitator & Member	Gabriel Engel @gabrielengelgl
Member	Allison Browne @allison.browne
Member	Grzegorz Bizon @grzesiek
Member	Arran Walker @ajwalker
Member	Oliver Falk @ofalk
Member	Marius Bobin @mbobin
Member	Cheryl Li @cheryl.li

SECTION: company/working-groups/ci-queue-stability.md

Attributes

Property	Value
Date Created	November 1, 2019
Date Ended	January 22, 2019
Slack	wgciqueuestability (only accessible from within the company)
Google Doc	CI Queue Stability Working Group (only accessible from within the company)
Issue Label	wgCIQueueStability (gitlab-com/-org)

Business Goal

Increase the stability and predictability of the CI job queue times on GitLab.com.

Intent is to:

1. Analyze and remediate situations where our CI job queue times for shared runners exceed reasonable expectations

1. Define metrics and tune alerting that more precisely correspond to the expectations of the CI job queues
1. Develop troubleshooting and investigation guides to use in cases of excessive CI job queue times
1. Perform predictive analysis on system health and growth and create issues to remediate anticipated future bottlenecks

Exit Criteria

1. Creation and tuning of metrics and alerts that trigger when system behaviour no longer matches expectations -> Done
1. 1-week of running with above mentioned tuned alerts without them going off -> Done
1. Published or updated documentation of runbook information on how to diagnose, respond, and restore abnormal behavior into being normal -> Done

Artifacts

1. Updated runbook for ci-runners service has a apdex score (latency) below SLO alert that could be triggered

Roles and Responsibilities

Working Group Role	Person	Title