iews, indexes, data types, functions, stored procedures and operators, see docs | | |
| GitLab DB Schema | An application-level table classification schema that abstracts away the underlying database connection, see docs | | |
| Server | A database server is a physical or virtual system running an operating system that is running one or more database instances. | Physical Database | |
| Table | A database table is a collection of tuples having a common data structure (the same number of attributes, in the same order, having the same name and type per position) (source) | | |
| Table Partitioning | A table that contains a part of the data of a partitioned table (horizontal slice). (source)| Partition | |
| Dataset | A set of tables and their contained data that is contained within a logical database. | | The Sec Dataset includes all tables related to GitLab's security features, including but not limited to vulnerability and dependency tracking. |
| Featureset | A set of features associated with some kind of concept within GitLab for ease of reference. | | Core, Sec |
| Core | Referred to in terms of Dataset or Featureset, this is information of functionality related to standard GitLab operations, such as Projects, Namespaces, Users and others.  | | |
| Sec | Referred to in terms of Dataset or Featureset, this is information of functionality related to standard GitLab operations, such as Vulnerabilities, Dependencies (SBOM), Security Findings, Policies and more. | | |

Overview

There is high impetus within GitLab to reduce pressure on the primary GitLab database server. The Database and Scalability teams have been taking a variety of steps to mitigate the ongoing pressure on the database server to maintain the growth and stability of GitLab in the long term. One such endeavour is Cells, however, there is desire to provide further mitigation in the short to medium term. Decomposition of the Sec dataset from the primary database was identified as a strong possible solution, similar to how the CI decomposition aided in this regard in the past.

Decomposition of the Sec dataset is a significant engineering effort due to the magnitude of the data interactions related to these features. The domain accounts for 25% of all database write traffic, and is only set to grow as we expand our feature set and grow our customer base. Further statistics and technical details can be found on the associated epic.

As this has become a scalability and stability concern for all of GitLab.com, as well as significantly constraining the ability of the Stages in the Sec section to implement new features due to continuously growing performance concerns, it is necessary to form an organised effort to effectively achieve this project.

We have the benefit of being able to lean heavily on the prior art and experience of the database-scalability working group who decomposed the CI database to achieve this goal. However, some key challenges we may face is the scale of the existing Sec codebase, and the need to maintain ongoing operations with (no/minimal) disruption to our customer base. A full GitLab.com downtime is heavily disfavoured due to our uptime SLA agreements with customers, but the scale of our operations may mean that some processes for this kind of decomposition may not be feasible.

Benefits

1. Reduce write pressure on the GitLab.com primary Write database in advance of Cells 1.5
2. Improve stability of GitLab operations, by isolating the primary database from Sec feature pressure
3. General performance improvement for both the Core and Sec feature sets due to seperation of concerns.
4. Improve iteration speed of Sec feature development without significant concern for compromising stability of the platform.

## Risks

1. Significant developer commitment for a currently unknown duration.
2. Increased database maintenance requirement for the new decomposed database and it's associated replicas.
3. Possibly unable to be delivered prior to the delivery of Cells 2.0
4. May require a full downtime of GitLab.com, which may be difficult to arrange with our customers.

## Interdependencies

Sec Data has a high degree of integration with CI and standard GitLab data, such as Users, Projects and Namespaces. The past CI decomposition has succesfully delinked query interdependency of the associated CI dataset, however, significant effort will be necessary to do the same between the core GitLab dataset and Sec functionality.

## Timeline

The group has determined that a gradual rollout is not advisable since it doesn't relieve pressure
on the main database cluster, and increases the amount of work to be done. As such, all slices will
be rolled-out simultaneously to gitlab.com.

## Progress

```
mermaid
gantt
dateFormat YYYY-MM-DD
title Estimated Decomposition Timeline

section Timeline
Gitlab Decomposition Ready :active , decompose, 2024-07-01, 2025-02-14
Non-Slice Work :active, nonslicework, 2024-07-15, 2025-02-14
Slice 1 :active, slice1, 2024-07-23, 2025-01-13
Slice 2 :active, slice2, 2024-08-06, 2024-12-30
Slice 3 :active, slice3, 2024-07-15, 2025-02-10
Gitlab Application Ready for Decomposition :milestone, allslices, after slice1 slice2 slice3
nonslicework, 0d
Phase 1 & 2 : phase12, 2024-09-11, 16w
Phase 4 : phase4, 2025-03-14, 3w
Phase 5 : phase5, after phase4, 1w
```

Phase 6 : phase6, 2025-03-21, 3w
Phase 7 : phase7, after phase6, 1w
Rollout complete :milestone, rollout, 2025-04-19, 1d
axisFormat  %Y-%m-%d

Source.

## Decomposition

| Slice | % Done | Estimated completion |
| --- | --- | --- |
| Slice 1 | 100% | Complete |
| Slice 2 | 100% | Complete |
| Slice 3 | 100% | Complete |
| Non-slice work | 97% | 2025-02 |

Last update: 2025-02-18.

## Plan

1. Introduce separate gitlabsec schema
1. Introduce gitlabsec database connection (defaulting to fallback to using gitlabmain database)
1. In parallel, begin decomposition of foreign keys and cross-database transactions following the loose order of SBOM, Security, and Vulnerability code boundaries. For each slice perform the following breakdown:
    1. Migrate tables with low referentiality (few foreign keys)
    1. Migrate tables with higher referentiality (many foreign keys)
    1. Identify and allowlist cross-joins to be addressed
    1. Identify and allowlist cross-database transactions to be addressed
    1. Remove previously identified cross-joins and cross-database transactions allowances
1. Formulate a logical replication path for the safe migration of the Sec dataset to a new physical database.
1. Open Change Request to migrate tables using a single replication event for all tables in scope of decomposition

## Data Migration Proposal

See rollout for full details

1. With physical-to-logical replication we replicate the full DB before converting to logical replication for the relevant sec tables
    1. Deploy the decomposed database instance as a streaming replica of main
    1. Begin replicating the full Sec data to the new database instance
    1. Establish a separate sec DB connection pointed at the same main DB
    1. Write the necessary code to enable GitLab.com to begin utilising the new DB connection generically and for all Sec features.
    1. As this is a potentially risky operation, ensure production snapshots are ready and that customers are sufficiently informed of potential problems or dataloss in the event of failure.

1. Begin testing transition of the Sec featureset to using the new database instance as it's new primary (gstg -> canary -> grpd)
   1. If successful, globally rollout usage of the decomposed database for the full featureset.
2. Truncate legacy sec tables from the GitLab main database.

## Roles and Responsibilities

| Working Group Role | Name | Title |
| ----------- | ----------- | ----------- |
| Executive Stakeholder | Jerome Ng | Engineering Director, Expansion |
| Functional Lead | Gregory Havenga | Senior Backend Engineer, Govern: Threat Insights |
| Functional Lead | Lucas Charles | Principal Software Engineer, Sec |
| Facilitator AMER | Neil McCorrison | Manager, Software Engineering |
| Facilitator APAC | Thiago Figueiró | Manager, Software Engineering |
| Member | Fabien Catteau | Staff Engineer, SSCS: Pipeline Security |
| Member | Arpit Gogia | Backend Engineer, AST: Dynamic Analysis |
| Member | Schmil Monderer | Staff Backend Engineer, APM: Threat Insights |
| Member | Ethan Urie | Staff Backend Engineer, AST: Secret Detection |
| Member | Jon Jenkins | Senior Backend Engineer, Database |
| Member | Ved Prakash | Staff Data Engineer, Data Science |
| Member | Dylan Griffith | Principal Engineer, Create |
| Member | Thong Kuah | Principal Engineer, Data Stores |
| Member | Rick Mar | Manager, Core Infrastructure |

## Related Performance Projects

1. Tuple Reduction
   - Brian Williams (DRI)
   - Fabien Catteau
   - Michael Becker
1. Vulnerability Management Application Limits and Vulnerability Management Retention Policy
   - Mehmet Emin Inaç (DRI)
   - Joey Khabie
1. Cells 1.0
   - Subashis Chakraborty (DRI)

## Useful References

| Reference | Description |
| --- | --- |
| Link | Proposal for support levels for multiple databases in GitLab deployment architecture. |
| Link | Epic dashboard for tracking outstanding work towards completion of decomposition |

Thanks

Much information and inspiration and experience is being enjoyed from the database-scalability working group who accomplished the succesful decomposition of the CI database.

---
## SECTION: company/working-groups/task-groups/_index.md

What's a Task Group?

A Task Group is a streamlined, goal-oriented team formed to accomplish a specific, well-defined objective. These groups are characterized by their lean structure and focused approach, minimizing overhead and maximizing efficiency. Task Groups are composed of subject matter experts who collaborate to rapidly and effectively achieve their singular goal. This format is ideal for targeted initiatives requiring specialized knowledge and swift execution.

Differences between a Task Group and a Working Group

While both Task Groups and Working Groups are integral to achieving organizational goals, they differ in structure, scope, and function.

Structure and Scope

- Task Group:
  - Focuses on a single, well-defined objective.
  - Composed of subject matter experts collaborating to achieve a specific goal.
  - Lean and agile, with minimal hierarchy and bureaucracy.
  - Ideal for tasks that require rapid, focused execution.

- Working Group:
  - Address broader, high-impact business goals.
  - Typically involves more complex and varied functions across the organization.
  - Has defined roles and responsibilities with a broader scope than Task Groups.
  - Requires an executive sponsor and often involves collaboration across multiple functions.

When to Use Each

- Use a Task Group when:
  - The objective is clear-cut and specific.
  - Speed and efficiency are priorities.
  - The task requires specialized knowledge from experts.
  - The organizational overhead of a Working Group can be avoided.

- Use a Working Group when:
  - The goal is complex and multifaceted, impacting various parts of the

organization.
- The project requires collaboration across different functions and teams.
- There's a need for a structured approach with clearly defined roles.
- The project has strategic importance warranting executive oversight.

## Roles and Responsibilities

### Task Group Directly Responsible Individual (DRI)

This role is pivotal, with the individual taking full accountability for the group's success. They coordinate activites, lead meetings, and ensure the group's objectives are met efficiently. The DRI is the main point of contact and decision-maker for the group, bridging communication and action.

### Member

Members are subject matter experts who actively participate in meetings and contribute to task completion. They are responsible for staying informed via the Task Group's communication channels, sharing insights with peers, and gathering feedback. Members may work independently on tasks within the Task Group or collaborate with others, playing a critical role in driving the group's task forward.

## Guidelines

1. A Task Group is created, in part, to prevent proliferation of working groups.
1. Participation in some Task Groups requires a significant time commitment. Participants should have a clear understanding of their roles and the expectations for their deliverables within the Task Group. They should manage their time and capacity and quickly escalate if they feel unable to serve in or deliver in their role.
1. It is highly recommended that anyone in the Task Group with OKRs aligns them to the effort.

## Process

1. Create an MR with an overview page in the task-groups/ directory. Ensure a brief description with an actionable objective is included. It is generally recommended that your manager approves the MR prior to merging.
1. Optionally, create an associated channel in Slack.
1. Announce the creation of the Task Group to the subject matter experts.
1. Interested members can create MRs adding themselves to the Task Group, ensuring that their manager is informed.
1. The DRI can use the Slack channel to set up initial communication guidelines.
1. The group should determine a roadmap for completion of the task, with an expected end date, as soon as they can.
1. Once the task is complete, the Task Group should be disbanded.

Active Task Groups (alphabetic order)

1. Accessibility Audit Triage
1. BootstrapVue Removal

Past Task Groups (alphabetic order)

1. CSS Utilities
1. GitLab UI @vue/compat Compatibility
1. Vue 3 Router 4 Test Compatibility

---
## SECTION: company/working-groups/task-groups/bootstrap-vue-removal.md

Attributes

| Property | Value |
| --- | --- |
| Date Created | 2024-09-10 |
| Target End Date | 2025-08-31 |
| Slack | tgbootstrapvueremoval (only accessible from within the company) |
| Google Doc | Agenda doc (only accessible from within the company) |
| Zoom Recordings | Recordings (only accessible from within the company) |

Context

Previously, we chose to use BootstrapVue because we were already employing Bootstrap, and creating a custom UI library would be a significant undertaking. At that time, BootstrapVue was an appropriate choice. Initially, BootstrapVue was part of the monolithic codebase, but later it was moved to the GitLab UI repository.

While BootstrapVue was suitable in the past, we are now considering migrating away from it. There are several disadvantages that hinder maintenance and usage, as detailed below. The primary reason for migration is that BootstrapVue is now unsupported and impedes the migration of gitlab-ui (and ultimately, gitlab) to Vue 3.

Goals

1. Primary Goal: Enable Vue 3 Migration: BootstrapVue does not support the Compat Build mode for Vue 3. Previous attempts to build a compat mode have proven as time-consuming as migrating to our own components. While enabling the Vue 3 migration is our immediate priority, removing BootstrapVue would prevent potential future blocks, such as issues with upcoming Vue versions like Vue 3.5, which had been announced September 1st 2024, or even Vue 4 in the future. By maintaining our own components, we ensure continued flexibility and freedom from external library constraints.

2. Secondary Goal: Improve Developer Experience (DX)
   - Single Source of Truth: The GitLab UI repository currently 'mirrors' BootstrapVue components. We maintain two documentation sources: the Pajamas design system and the Storybook for Vue Components. The Storybook advises referring to BootstrapVue documentation for complete information. Consolidating these resources into a single source of truth would enhance DX and facilitate a clearer understanding of our design system.
   - Design Token Migration: Using BootstrapVue components complicates the addition of design tokens, requiring workarounds like overriding nested CSS selectors and contending with Bootstrap styles and HTML structure.
   - Design Systems Team Efficiency: Expanding components for GitLab developers often requires substantial effort, leading to manual implementations in the Monorepo instead of using BootstrapVue components.
   - Component Usability: The component <table> proofed to be not useable for most of our usecases, and custom implementations had to be done.
   - Ownership and Responsibility: Developing our own components ensures clear responsibilities within the Component System and a dedicated team for building and maintaining them.
   - Enabling other repositiories: E.g. transitioning to a Compat Mode or Vue 3 mode could be considered if we no longer rely on BootstrapVue.

## General Advantages of Maintaining Our Own System

- Reduced Dependencies: Removing BootstrapVue would eliminate dependencies from our codebase. See here for a list of dependencies BootstrapVue installs.
- Accessibility Improvements: Some complex components in BootstrapVue (like Popover) have accessibility issues. Creating and maintaining our own UI library would allow us to improve accessibility and reduce time spent on workarounds.
- Enhanced Flexibility: We would have the opportunity to integrate various features and tools as needed without relying on an external library.
- Maintenance Simplification: Although maintaining a UI Component System involves effort, avoiding reliance on external libraries like BootstrapVue would eliminate the need for ongoing updates and related maintenance challenges.
- Repository Impact: The removal of BootstrapVue will take place in the gitlab-ui repository. This will result in reduced bundle sizes for the gitlab repository by shipping less code and provide increased flexibility to address GitLab's specific needs moving forward.

## Challenges

- Compatibility: Ensuring that our custom components maintain the same APIs as before is crucial. As with most migrations, the focus will be on replicating the functionality of BootstrapVue components without unnecessary changes.
- Global Events: BootstrapVue supported global events (bv::). Rebuilding this functionality might be complex.

These are the known challenges, but there may be additional issues that have yet to be identified.

## Exit Criteria

1. Unused files from /src/vendor/bootstrapvue are removed: Epic &13075.
2. Documentation is updated and made accessible: Issue 2754.

3. MIT License usage is clarified: Issue 2318.
4. Files are migrated: Epic &15178
   1. Directives are moved to the /src folder.
   2. Components, plugins, mixins, and constants are rewritten using the 'usual' Vue 2 Options API syntax and moved to the /src folder.
5. The directory /src/vendor/bootstrapvue is removed.
6. Bootstrap CSS Utilities in GitLab UI are replaced with (TailwindCSS) GitLab CSS Utilities: Epic &15765.

Roles and Responsibilities

| Task Group Role | Person | Title |
| --- | --- | --- |
| DRI | Peter Hegman | Senior Frontend Engineer, Tenant Scale::Organizations |
| Member | Paul Gascou-Vaillancourt | Senior Frontend Engineer, Foundations::Personal Productivity |
| Member | Lukas Eipert | Staff Frontend Engineer, Foundations::Personal Productivity |
| Member | Lorenz van Herwaarden | Senior Frontend Engineer, Govern::Threat Insights |
| Member | Chaoyue Zhao | Frontend Engineer, Create::Source Code |
| Member | Thomas Hutterer | Senior Fullstack Engineer, Foundations::Personal Productivity |

---
## SECTION: company/working-groups/task-groups/css-utils.md

Attributes

| Property | Value |
| --- | --- |
| Date Created | 2024-02-20 |
| Target End Date | 2025-01-31 |
| End Date | 2024-09-24 |
| Slack | tgcssutils (only accessible from within the company) |
| Google Doc | Agenda doc (only accessible from within the company) |

Context

The frontend department faces recurring challenges around writing CSS utilities, using them in

the
product, controlling CSS bundles' size and ensuring pages only import the style definitions
they
actually need. This task group aims at alleviating those pain points through the following
driving
factors:

- Reduce CSS bundle sizes.
- Compartmentalize style definitions.
- Keep enforcing Pajamas-compliance through consistent and design token ready CSS utilities.
- Increase our dark theme readiness.
- Improve developer experience (DX).

## Exit Criteria

### September 2024 exit criteria update

In September 2024, we have reassessed the task group's exit criteria. At this point, we had
mostly
completed the migration to Tailwind CSS and were only left with a few cleanup tasks. When
reviewing
our goals, we noticed that the remaining scope might be too large for a structure such as a
task
group. Task groups are meant to be short-lived, and to focus on completing well-defined tasks
in a
timely fashion. For posterity, here are our original exit criteria:

> 1. CSS utilities are generated with Tailwind:
<https://docs.gitlab.com/ee/architecture/blueprints/tailwindcss/>.
>     1. They strictly enforce Pajamas specifications through the use of design tokens.
>     1. Color utilities support dark theme overrides.
>     1. Any necessary migrations have been completed: https://gitlab.com/groups/gitlab-org/-/epics/12108.
>     1. Development guidelines have been updated to reflect any direction change in how we write CSS.
> 1. Page-specific style definitions have been moved to individual bundles:
https://gitlab.com/groups/gitlab-org/-/epics/3694.
> 1. CSS rules have been migrated to CSS utilities: https://gitlab.com/groups/gitlab-org/-/epics/8326.
> 1. Color declarations have been consolidated to only use Pajamas-compliant ones:
https://gitlab.com/groups/gitlab-org/-/epics/8599.

As a result of the re-assessment, we are acknowledging that:

- The CSS utilities to Tailwind CSS migration is done.
- The "page-specific CSS bundles" and "CSS rules to CSS utilities" initiatives are conflicting
with
  each other. One of these should be abandoned in favor of the other.
- We, as a task group, don't have the means to tackle the aforementioned initiatives. These
should

be owned by the frontend department, or we should spawn a new task group dedicated to these.
- What we can do however is provide guidance on what we think would be the best approach going forward. This has been done through documentation updates.
- Regarding the color declarations consolidation, the fact that Tailwind CSS is now available in our main projects puts us in a better spot to achieve this goal. However, this is again outside of the scope of this task group. As we are still in the process of migrating our styles to design tokens, this initiative can be owned by the Design System group going forward.

Given all of the above, here are our updated exit criteria:

1. CSS utilities are generated with Tailwind in all active GitLab projects using GitLab UI.
1. Guidance as to how styles should be written going forward has been documented.

As of September 24, 2024, these goals have been attained and we are therefore disbanding the task group.

Roles and Responsibilities

| Task Group Role | Person | Title |
| --- | --- | --- |
| DRI | Paul Gascou-Vaillancourt | Senior Frontend Engineer, Foundations::Personal Productivity |
| Member | Florie Guibert | Senior Frontend Engineer, Plan::Product Planning |
| Member | José Iván Vargas López | Senior Frontend Engineer, Verify::Pipeline Execution |
| Member | Peter Hegman | Senior Frontend Engineer, Data Stores::Tenant Scale |
| Member | Savas Vedova | Staff Frontend Engineer, Govern::Threat Insights |
| Member | Vanessa Otto | Senior Frontend Engineer, Foundations::Design System |
| Member | Lukas Eipert | Staff Frontend Engineer, Foundations::Personal Productivity |

---
## SECTION: company/working-groups/task-groups/gitlab-ui-vue-compat.md

Attributes

| Property | Value |
| --- | --- |

| Date Created | 2024-10-14 |
| Target End Date | 2024-12-24 |
| End Date | 2025-02-03 |
| Slack | tggitlabuivuecompat (only accessible from within the company) |

Context

The scope of work to migrate everything in GitLab to Vue 3 is huge. As a result, the Vue 3 Working Group has disbanded, in favour of more focussed task groups, like this one.

Since the GitLab UI library is used by various projects (GitLab, Switchboard, Editor Extensions), it is a major blocker for Vue 3 adoption. While a lot of work has already been done to make GitLab UI work with @vue/compat, much remains.

Goals

- All components and directives of GitLab UI should be usable under @vue/compat.

Non-goals

- We are not going to bring any components or directives closer to MODE: 3 compatibility, since it doesn't block adoption of @vue/compat.

Exit Criteria

- All known issues with @vue/compat fixed, or work-arounds documented.
- All screenshot tests passing under @vue/compat

As of 2025-02-03, the above criteria have been met, so this task group is disbanded.

While some issues remain,
they are considered low priority, and/or cannot be fixed until we drop support
for Vue 2.

Roles and Responsibilities

| Task Group Role | Person | Title |
| --------------- | ------------- | --------------------------------------------------|
| DRI | Mark Florian | Staff Frontend Engineer, Foundations::Design System |
| Member | Marina Mosti | Sr. Frontend Engineer, Switchboard |
| Member | Miguel Rincon | Staff Frontend Engineer, Verify:Runner |

---
## SECTION: company/working-groups/task-groups/vue3-quarantined-tests.md

Attributes

| Property | Value |
| -------- | ----- |
| Date Created | 2025-02-28 |
| Target End Date | 2025-08-31 |
| Slack | tgvue3quarantinedtests |

Context

A major effort is underway to ensure that the GitLab unit tests are passing and compatible with Vue compat.
A major cause of the current failures is a compatibility issue in how tests are written.

There is no single reason for tests to fail but most of the known reasons and how to handle it were gathered under epic Fix failing jest unit tests for Vue.js 3 in gitlab-org/gitlab |

Goal

Fix all tests that are currently failing due to incompatibility.

The list of tests suites that need be fixed will be maintained in this quarantined list. More user friendly representation can be found on this page to keep this up to date.

The strategy for fixing these tests has been documented in the epic.

Non-goals

There is separate working group Vue 3 Router 4 Test Compatibility Task Group, dedicated specifically for fixing tests with Vue router compatibility errors

Exit Criteria

Tests issuing any compatibility errors fixed (where applicable, see non-goals above).

Roles and Responsibilities

| Task Group Role | Person | Title |
| --------------- | ------ | ----- |
| DRI | Artur Fedorov | Sr. Frontend Engineer, Security Policies |
| Member | Alexander Turinske | Staff Frontend Engineer, Security Policies |

---
## SECTION: company/working-groups/task-groups/vue3-router4-tests.md

Attributes

| Property | Value |
| -------- | ----- |
| Date Created | 2025-01-02 |
| Target End Date | 2025-03-31 |

| Slack | tgvue3router4tests |
| --- | --- |

## Context

A major effort is underway to ensure that the GitLab unit tests are passing and compatible with Vue compat.
A major cause of the current failures is a compatibility issue in how tests are written when incorporating Vue Router.

Vue 2's router version is Vue router 3, whereas Vue 3's router version is Vue router 4. There are several differences between the two, explained in the migration guide.

These differences between Vue Router 3 and 4 require that unit tests are re-written based on the findings of the initial investigation.

## Goal

Fix all tests that are currently failing due to Vue Router incompatibility.

The list of tests suites that need be fixed will be maintained in this section of the original issue. A script that fetches the router errors is available to keep this up to date.

Once these tests are fully addressed, it will be necessary to take another look at the jest speed reporter to make sure that no other tests are failing due to router compatibility. The above investigation has assumed that tests failing due to compatibility issues are exposed by the string "Vue router" in the failure message.

The strategy for fixing these tests has been documented in the Testing Vue router section of the Vue 3 testing handbook.

## Non-goals

Tests that are not directly related to the compatibility of Vue Router 3 and 4 won't be considered, as they are not covered by the migration strategy.

## Exit Criteria

Tests issuing any router compatibility errors fixed (where applicable, see non-goals above).

As of 2025-03-07 the above criteria has been met.

No tests related to router compatibility are pending correction, and Vue router 4 testing documentation has been created.

## Roles and Responsibilities

| Task Group Role | Person | Title |
| --------------- | ------ | ----- |
| DRI | Marina Mosti | Sr. Frontend Engineer, Switchboard |

---
## SECTION: customer-experience/_index.md

Customer Experience (CX)

Welcome to the Customer Experience team page.

About Customer Experience (CX)

Customer Experience (CX) is a field that encompasses the entire end-to-end customer journey, from initial awareness, evaluation and purchase to post-sales adoption and expansion, focusing on creating positive, seamless, and meaningful interactions to build loyalty and customer satisfaction.

Goals of Customer Experience (CX)

Increase customer satisfaction, loyalty, and retention through delivering exceptional experiences and building strong customer relationships to drive repeat business
Reduce customer effort and friction across touchpoints by identifying top areas of friction to remediate cross-functionally.
Build a positive brand. Drive positive word-of-mouth and advocacy, and differentiate from competitors in crowded markets.
Increase customer lifetime value and reduce acquisition costs by retaining customers and encouraging repeat purchases to increase the long-term value for each customer.
Align organizational processes and culture around customer needs, creating customer-first organization and product.

Read more about Customer Experience (CX) as a discipline.

Our Mission

To create exceptional value through deep customer understanding, enabling GitLab and our customers to achieve mutual success through meaningful engagement, actionable insights, and continuous innovation.

Our Vision

We create meaningful connections with customers, transforming every interaction into an opportunity to build trust. By anticipating needs and delivering personalized experiences, we make customers feel valued and supported, turning them into passionate advocates who drive our shared success. (Powered by Claude)

Key Objectives

| Drive Customer-Centric Decision Making|Enable Customer Success via Strategic Engagement| Champion Customer Innovation and Growth|
| :---- | :---- | :---- |
| Establish comprehensive Voice of Customer (VoC) programs for feedback loops, customer intelligence frameworks, and biannual customer satisfaction  insights mechanism to provide

actionable experience insights for the organization. | Drive customer engagement through strategic programs, executive relationships, peer learning platforms, and advisory boards to accelerate value realization. | Foster customer growth by identifying innovative use cases, targeted education opportunities, enabling customer influence on strategy, and demonstrating value via success stories and foundational metrics. |

Our Work

CSAT surveys, VoC programs
Customer Journey Research Projects
Foundational (·Lighthouse·) metrics for CX
·and more to come as our team grows\!

Current Priority Programs (FY26)

| Program | Objective | Why | Cadence |
| :---- | :---- | :---- | :---- |
| Customer Experience Journey Research | Map our customer·s journey in purchasing and adopting GitLab, and experience with our teams, to identify opportunities to accelerate value. | Improving customer experience, results in increased customer satisfaction, and driving revenue through stronger lands, renewals, and expansions. | Quarterly |
| All Customer Satisfaction Survey | Customer Satisfaction benchmark to provide insights and measures to the business to improve the value of GitLab for our customers. | By keeping a pulse on customer satisfaction and feedback themes, we can identify areas of improvement to improve customer satisfaction. | Biannually (May, Nov) |
| Customer Health Framework (page coming soon) | Instrument and baseline customer health metrics | Identify areas of opportunity to focus on improving customer health | TBD |

Meet Our Team

Meet our small but mighty CX team\!

Brittney Sinq \- Head of Customer Experience (CX)
Sarah Schuster \- CX Strategy
Brandon Butterfield \- CX Strategy Analytics

Learn more about how we collaborate with our cross-functional partners.

Get in touch\!

We are always seeking to talk to customers like you. While this team can·t help with technical support, billing or licensing issues, we can influence how our field teams can improve upon partnering with you to better understand your needs and help you realize value faster through GitLab. Reach out at csfeedback@gitlab.com.

---
## SECTION: customer-experience/about-cx.md

Deep dive into Customer Experience (CX)

Understanding how this critical practice drives customer-centricity in GitLab

## What is Customer Experience (CX)?

Customer Experience (CX) encompasses the entire customer journey, from initial awareness, evaluation and purchase to post-sales adoption and expansion, focusing on creating positive, seamless, and meaningful interactions to build loyalty and satisfaction. CX focuses on understanding customer needs, expectations, and emotions to create meaningful, seamless experiences that foster loyalty and advocacy.

## What Customer Experience (CX) is not

- Just customer service: Customer service is one component of the overall customer experience
- Limited to UX: UX is an important element of the customer experience, but it is not the totality of the experience.
- A one-time fix: CX is not a one-time project but rather an ongoing strategic approach
- A department-specific initiative: CX requires a cross-functional approach, with all departments working together to create a consistent and positive experience.

## Goals of Customer Experience (CX)

- Increase customer satisfaction, loyalty, and retention through delivering exceptional experiences and building strong customer relationships to drive repeat business
- Reduce customer effort and friction across touchpoints by identifying top areas of friction to remediate cross-functionally.
- Build a positive brand. Drive positive word-of-mouth and advocacy, and differentiate from competitors in crowded markets.
- Increase customer lifetime value and reduce acquisition costs by retaining customers and encouraging repeat purchases to increase the long-term value for each customer.
- Align organizational processes and culture around customer needs, creating customer-first organization and product.

## Common Customer Experience (CX) Deliverables

Customer Journey Maps. Visual representations of a customer's end-to-end experience with a company, highlighting touchpoints, emotions, pain points, and opportunities.

Voice of Customer (VoC) Programs. Systematic collection and analysis of customer feedback through surveys, interviews, social media monitoring, and other channels.

Customer Personas. Detailed, research-based profiles of typical customer segments that help teams understand and empathize with different user needs and behaviors to drive more personalized experiences.

Experience Design Frameworks. Structured approaches for designing new experiences or improving existing ones based on customer insights.

CX Data Analytics & Reporting. Define and track CX metrics and KPIs like Customer Satisfaction (CSAT), Customer Lifetime Value (CLV), and churn rate to track CX performance and more.

Service Blueprints. Operational diagrams that visualize service processes, including front-stage customer interactions and backstage support activities.

Real-World Examples

Zappos: Built their entire business model around exceptional customer experience, with policies like 365-day returns and 24/7 customer service.

Apple: Created seamless integration across products, services, and physical retail spaces with the Genius Bar to provide consistent, high-quality experiences.

Amazon: Pioneered one-click ordering, personalized recommendations, and efficient delivery options to make shopping effortless.

The most successful CX initiatives align customer needs with business objectives and transform insights into meaningful improvements across all touchpoints in the customer journey.

---
## SECTION: customer-experience/raci.md

Our Team

Who we are, what we do, and who we work with across GitLab to drive customer excellence.

Meet The Team

We are a small, but mighty, team for now. More teammates will be added in future quarters\!

| Role | Responsibilities | Outputs |
| :---- | :---- | :---- |
| Head of Customer Experience | Champion for Customer Experience team across GitLab | Unblocks CX team, advocates for cross-functional resources across the org. |
| Customer Experience Strategist | Leads the overall customer experience strategy and programs, ensuring the voice of the customer drives business decisions and customer-centricity across GitLab. Leads customer research initiatives and manages feedback programs to provide actionable insights across the organization. | Customer journey maps, customer expeirence research insights,straetgic recommendations to improve high-friction areas |
| Customer Experience Data Analyst | Drives data-driven decision making across CX programs by managing customer analytics, measuring program effectiveness, and providing actionable insights on customer adoption patterns and program success. | End-to-end dashboards and reporting, partnering with Aalytics & Insights teams to instrument new metrics |

Cross-Functional Partners

| Department | How We Partner | Knowledge Share |
| :---- | :---- | :---- |
| Product (PM, PMM, UX, Data) | Research best practices Product usage data and trends | USAT survey insights Feedback on product roadmaps, releases |
| Sales | Aligning customers for interviews Insights into internal team pain points, needs |

Patterns/trends from the field Deal insights (wins, losses, trends) |
| Customer Success CSM/A/E Renewals Success Services | Aligning customers for interviews Insights into internal team pain points, needs Actioning improvements from research work (·the how·) | Patterns/trends from the field Wins/loss stories |
| Professional Services |  |  |
| Support | Insights into internal team pain points, needs Actioning improvements from research work (·the how·) | Insights/trends from support ticket data  |

Roles & Responsibilities Matrices

CX Customer Journey Research: RACI Matrix

While our team is in its start-up phase, our team must be flexible in how we contribute and lead the projects in our portfolio, often expanding into roles and tasks outside of our commonly defined scope. Below is the R\&R matrix specific to customer journey research. (Powered by Claude, refined by me)

| Role | Responsibilities |
| :---- | :---- |
| CX Strategist | Acts as primary project manager and owner Leads qualitative research design and execution Serves as primary interviewer and insights synthesizer Owns project timeline and deliverables Manages stakeholder communication and expectations Presents findings to leadership |
| CX Analyst | Leads quantitative research components Designs and analyzes metrics for measuring customer experience Collaborates on research methodology and sample design Provides data-driven insights to complement qualitative findings Develops data visualizations and analytical frameworks Ensures statistical validity of findings |
| Coross-functional Collaborators | Facilitate access to customers within their domains Provide contextual knowledge from their areas of expertise Co-facilitate interviews when domain knowledge is required Share relevant insights from their customer interactions Help implement recommendations within their areas Bridge organizational silos to ensure comprehensive insights |
| Steering Committee | Provides final approval on project scope and direction Removes organizational barriers Allocates resources and resolves budget issues Makes final decisions on strategic recommendations Champions findings to broader organization Holds team accountable for outcomes and impact |

Key

R \= Responsible (Does the work)
A \= Accountable (Ultimate approver/owner)
C \= Consulted (Provides input before decisions)
I \= Informed (Kept updated on progress/decisions)

Note: This RACI matrix should be reviewed and adjusted at each project kickoff to ensure alignment with organizational structure and project requirements.

Research to Strategy RACI

| Activity | CX Strategist | CX Analyst | XFN Collaborators | Steering Committee |

| | | | | |
| ----- | :---: | :---: | :---: | :---: |
| Planning & Scoping |  |  |  |  |
| Define research objectives | R/A | C | C | A |
| Develop research timeline | R/A | C | I | A |
| Stakeholder identification | R | C | C | A |
| Research Design |  |  |  |  |
| Methodology selection | R/A | C | C | I |
| Survey/interview design | R | C | C | I |
| Sampling strategy | R | R/C | C | I |
| Research tools selection | R | R | I | I |
| Participant Recruitment |  |  |  |  |
| Participant criteria definition | R | C | C | I |
| Recruitment outreach | R | I | R | I |
| Scheduling interviews/sessions | R | I | R | I |
| Participant management | R | I | C | I |
| Data Collection |  |  |  |  |
| Interview facilitation | R | C | C | I |
| Data gathering logistics | R | R | C | I |
| Recording/documentation | R | C | C | I |
| Analysis & Insights |  |  |  |  |
| Qualitative data analysis | R/A | C | C | I |
| Quantitative data analysis | C | R/A | C | I |
| Pattern identification | R | R | C | I |
| Insights generation | R | R | C | C |
| Reporting & Recommendations |  |  |  |  |
| Draft findings report | R | R | C | I |
| Develop recommendations | R | R | C | C |
| Results presentation | R | R | I | A/C |

Strategy to Execution RACI

| Activity | CX Strategist | CX Analyst | XFN Collaborators | Steering Committee |
| ----- | :---: | :---: | :---: | :---: |
| Action Planning |  |  |  |  |
| Prioritize recommendations | R | C | C | A |
| Resource allocation | I | I | C | A |
| Implementation planning | C/I | C/I | R | A |
| Success metrics definition | C | C | C | A |
| Project Management |  |  |  |  |
| Team coordination | I | I | C | A |
| Timeline management | I | I | C | A |
| Risk management | I | I | C | A |
| Status reporting | I | I | C | A |
| Decision Points |  |  |  |  |
| Approve research plan | R | C | I | A |
| Methodology changes | R | C | I | A |
| Timeline extensions | R | C | I | A |
| Final recommendations approval | R | C | I | A |

Decision Flow and Escalation Path

1. Day-to-day decisions · CX Strategist
2. Analytical methodology decisions · CX Analyst (with Strategist input)
3. Cross-functional coordination · Cross-Company Collaborators
4. Strategic direction, scope changes, resource allocation · Steering Committee

---
## SECTION: customer-experience/cx-journey/_index.md

Customer Experience Journey Research

Understanding the current customer experience, identifying opportunities where we can improve, and leveraging metrics to advocate key experience improvement initiatives.

About Customer Experience (CX) Journey Research

Customer Experience (CX) research takes three perspectives into account:

1. Customer Lens: Understanding the customer·s key activities along our core customer journey, their needs, who on their side is engaged in purchase decisions, through to onboarding and adoption, and the needs and expectations of GitLab for success.
2. GitLab Teams: Understanding the experience of the GitLab team collaborating as a team in service of helping customers move through their purchasing and adoption journey with us, and where there are areas of friction inhibiting us from providing our best service to customers.
3. Customer \+ Gitlab Intersection: This intersection captures the customer experience through the engagement model between their teams and ours to realize value in their investment with us.

Where·s Product/UX  in this? Our in-product experience absolutely impacts our customer experience in a very real way. The focus of this team is centered on the engagement

Objective

Map our customer·s journey and experience evaluating, purchasing, and adopting GitLab in partnership with our customer-facing teams to identify opportunities to accelerate value. We seek to understand this across our core customer journey phases: Awareness, Evaluation, Purchase, Onboarding, Adoption, and Expansion across all customer segments and regions.

Key Results

Improved customer experience
Increased customer satisfaction
Faster value realization for our customers

The Research

This year, the research conducted by this team will focus on building Foundational Maps: from aligning the many journey stages across the business into a unified journey view, mapping the current state end-to-end journey, a visionary future state map for what good looks like, and a single pane dashboard connecting the pre- to \-post sales metrics for these journeys to establish the baseline for future experience measures.

The phases below outline how we will get to these artifacts over the next year:

(Visual coming soon)

As Foundations maps are established, the team will shift to Iterative maps which may include deep dives into specific phases or pain points in the journey, specific verticals, and/or regions. We will continue to iterate on our foundational maps annually to demonstrate improvements made to the overall customer experience.

Scheduled Research Projects

| Project | Quarter | Objective | Deliverables |
| :---- | :---- | :---- | :---- |
| (In Progress) Understanding Premium to Ultimate Uptiers for Large, Self-Managed Customers | Q1-Q2 | What is the experience for our largest customers to consider Ultimate, what factors get in the way of purchase and adoption, and what are the core elements of successful adoption? This journey also factors in the highest touch engagement model GitLab currently has to offer | Completion of 10 contextual inquiry interviews with customers Current & future state journey maps List of recc·s to improve this experience |
| (Up Next) Mapping Mid-Market First Order Lands with Premium and Duo | Q2 | What is the experience like for medium-sized customers to evaluate and purchase GitLab with our scaled services and teams? This journey may or may not include high-touch elements, depending on customer needs. | Completion of 20 contextual inquiry interviews with customers Current & future state journey maps List of recc·s to improve this experience |
| (Backlog) Product-Led Growth | Q3 | What is the experience like for customers who never talk to a human at GitLab from evaluating and purchasing GitLab, to onboarding and ultimately adopting GitLab for their organization? | Completion of 20 contextual inquiry interviews with customers Current & future state journey maps List of recc·s to improve this experience |
| (Backlog) Multi-year Lands (Migration, Onboarding, Adoption) | Q4 | What is the experience like for customers who purchase our Ultimate tier as a first-time customer, and also must migrate data, manage two tools concurrently, fully onboard their team onto GitLab and adopt for business success? | Completion of 10 contextual inquiry interviews with customers Current & future state journey maps List of recc·s to improve this experience |

Get in touch\!

We are always seeking to talk to customers like you. While this team can·t help with technical support, billing or licensing issues, we can influence how our field teams can improve upon partnering with you to better understand your needs and help you realize value faster through GitLab. Reach out at csfeedback@gitlab.com.

---
## SECTION: customer-experience/cx-journey/prem-ult-epic.md

Expansion Journey: Premium to Ultimate Up-tier

Purpose

Create the foundational customer journey to describe our customer experience with the GitLab teams expanding into Ultimate from awareness, through purchase and onboarding, to adoption and renewal for this tier of customer.

While there are many micro journeys that were created by teams that describe a moment in the journey or a product interface user experience, the holistic pre- through post-sales journey currently does not exist.

Essential to this research is not only understanding what the journey looks like for customers where Ultimate is the right fit, but also understanding where Ultimate tier is not a fit and why.

### Core Team

| Person | Role | What they own |
| :---- | :---- | :---- |
| Sarah Schuster | Customer Experience Strategy | Owns overall CX strategy, insights Lead research efforts, mapping journeys, communications |
| Brandon Butterfield | CX Strategy Analytics | Owns overall CX metrics, dashboards and reporting |
| Jong Lee | Data Analytics | Consultation on data analytics, insights |
| Ben Leduc-Mills | UX Research | Consultation on research methodology, best practices |

### Cross-functional Counterparts

| Product | Go-to-Market | Operations |
| :---- | :---- | :---- |
| Product Marketing, Product Management, User Experience (UX), Data & Analytics | Sales, Solution Architects, Customer Success, Services, Support | Strategy & Operations |

### Scope of Work

Customers who have considered, bought, and/or down-tiered from Ultimate tier
Mapping the customer journey (current state, future state)
Mapping of current metrics and data systems, dashboards
Contextual interviews of 15-20 customers
Interview internal teams ICs, leaders across field teams for input
Determine strategic recommendations to improve the customer experience

### Deliverables

Customer journey framework
Completed current and future state maps for this cohort
Customer Journey dashboard with baseline metrics
Three strategic recommendations to improve customer experience

### Key Resources

For more details on this research \- see program poster (add link)
GitLab Epic (add link)
Slack (internal only): \cxs-ultimate-tier-customer-journey

---
## SECTION: customer-success/_index.md

The Customer Success department is part of the GitLab Sales function who partners with our customers to deliver value and positive business outcomes throughout their journey with GitLab.

The team can be reached in Slack channel (internal only).

Mission Statement

To deliver value to all customers by engaging in a consistent, repeatable, and scalable way across defined segments so that customers see the value in their investment with GitLab, and we retain and drive growth within our enterprise customers.

The mission of the Customer Success Department is to provide these customers with experience in order to:

- Accelerate initial customer value
- Maximize long-term, sustainable customer value
- Improve overall customer satisfaction & referenceability
- Maximize the total value of the customer to GitLab

North Star Metrics

Our top-level metrics are:

1. Renewal Rate of ATR (available to renew)
1. Growth ARR
1. Customer Outcomes Realized

Customer Success Teams

- Customer Success Manager handbook
- Customer Success Engineer handbook
- Customer Success Architect handbook
- Renewals Managers handbook

Digital Strategy

- Digital Strategy handbook

Demo Systems

- Demo Systems documentation

Customer Success Decision Tree

CSMAE Decision Tree - internal only

## Account Team

The account team is comprised of the Strategic Account Executive/Account Executive, Solutions Architect (Enterprise), and Customer Success Manager.

More information about the account team

## Overlap Between Solution Architects and Customer Success Managers or Architects

SA owns 1) pre-sales technical evaluation and relationships prior to the initial sale and 2) tier upgrades and new business units (i.e., connected new) within an existing customer. CSM owns 1) post-sales customer relationship and 2) license upgrades within an existing customer.

More information on the transition and ownership between Pre-Sales and Post-Sales

## Other Resources

## Education and Enablement

As a Customer Success team member, it is important to be continuously learning more about our product and related industry topics. The education and enablement handbook page provides a dashboard of aggregated resources that we encourage you to use to get up to speed.

## Customer Success Playbooks

See the Playbooks Page

## Customer Workshops

CSM-Created, Enablement Focus:

- All CSM-created workshops

CSE-Created, Enablement Focus:

- GitLab User Webinars and Labs

## Using Salesforce within Customer Success

Visit this page for more info on using Salesforce within Customer Success.

## Using Gainsight with Customer Success

Visit this page for more information on using Gainsight within Customer Success.

## Dogfooding

Outside of Engineering the Customer Success team has the largest concentration of tooling development capability. The team has unique needs that can't always be solved by GitLab's single DevOps platform. However, it is important to dogfood and avoid dogfooding anti-patterns.

As a result the Product organization heavily weights internal customers when considering prioritization. If you are considering building tooling in support of Customer Success priorities outside of GitLab, please follow the dogfooding process.

## Customer Success AWS Test Account

In an effort to keep AWS spend down, initiatives are being taken to automatically clean up our AWS account. This account is primarily used as a proof of concept for IaC and creating demos for GitLab customers. An automated cleanup script is currently being tested that will tag, shutdown and delete old resources as they are no longer needed. The automation will:

- Turn off and Tag Un-named resources. When resources are created a "Name" tag should be created with a value that's meaningful and indicates who deployed the resource. Example: {initials}-GitLabRunner
- New Resources will be automatically tagged with a Discovered and Expiration tag
- The Expiration tag is 14 days after the discovery. The script will only a tag an instance once. If you need additional time, please change the date to a reasonable date for cleanup (Add a month or two for prospective customers)
- If a resource needs to be permanent please set termination protection on the instance. This should also include tagging the instance with an explanation on why it's permanent and what its for
- On expiration the resource will be shut off and left for 7 days
- In 7 days if the instance is still off a snapshot will be taken and it will be terminated
- If the instance is still on but the expiration has not been changed it will be terminated

## Customer Success Tools and Scripts

By customer or internal request, we sometimes develop tools to automate certain GitLab tasks using the API. The resulting tools and scripts are publicly available for everyone to use and contribute to in the GitLab CS Tools group.
Note: Those tools are not supported by GitLab Support.

## Communities of Practice

Community of Practice are cross-functional groups of SME's (or aspiring to be!) within the CS organization dedicated to a topic within GitLab or the broader DevOps space. The goal is to build assets, best practices, demonstrations, and share experiences we learn from prospects and customers. In turn, CoP will build broader technical depth within our CS organization to better advise our customers and influence our product roadmap.

## Customer Terrain Mapping Engagements

Terrain Mapping discovery engagements provide customers with the benefit of GitLab's experience with DevOps methodologies, Git, GitLab, CI, CD, and monitoring by brainstorming a high level, first draft discovery of the elements of a success plan to address various challenges. They are also mapped to professional services that can help with some of the elements identified in the engagement.

See the Terrain Mapping Engagements Page

Frequently Asked Questions

Customer Success team members maintain a FAQ to keep questions customers ask documented in a place where everyone can view and contribute to.

Customer Success resource links outside handbook

- Customer Reference Sheet
- Sales Collateral
- GitLab University
- Our Support Handbook
- Customer Collaboration Project template
- GitLab Demo Portal
- Workflow SA Demo Scenarios (Internal Only)
- SA-Created - Hands-On Workshops

Other Sales Topics

- Sales Handbook
- Sales Operations
- Sales Skills Best Practices
- Sales Discovery Questions
- EE Product Qualification Questions
- GitLab Positioning
- FAQ from prospects
- Client Use Cases
- Proof of Value Guidelines
- Account Planning Template for Large/Strategic Accounts
- Sales Demo
- Sales Development Group Handbook
- With Whom to Talk to Ask Questions or Give Feedback on a GitLab feature

Customer Success Meetings

Customer Success has a few standing meetings:

- CS Team Monthly All-Hands - Monthly on the second Wednesday
- CS Skills Exchange - once or twice Monthly

The different groups within CS also have standing meetings, including meetings for the Customer Success teams and Renewal Managers, regional groups, and social calls.

---
## SECTION: customer-success/account-team.md

---

The account team works together to drive value, success, and growth for a customer.

## Sales Segments

### Enterprise

An Enterprise account team is comprised of the following people:

- Strategic Account Executive (SAE)
- Solutions Architect (SA)
- Customer Success Manager (CSM) for qualifying accounts

### Commercial/Mid-Market

A Mid-Market account team is comprised of the following people:

- Account Executive (AE)
- Customer Success Manager (CSM) for qualifying accounts

In Mid-Market, Solutions Architects are pooled so they are not aligned with specific account teams.

## Roles and Responsibilities

### Strategic Account Executive (SAE) / Account Executive (AE)

- Negotiates contracts and price on initial sale and renewals
- Should have a strategy for prospects and how to close
- Works with the Solutions Architect (SA) on pre-sales prospects and post-sale upgrades
- Works with the Customer Success Manager (CSM) on account planning and renewals

### Solutions Architect (SA)

- Should lead discovery calls to find out where and why the prospect is having issues in the SDLC
- Runs POVs and scopes out SOWs
- The technical advisor to the SAE for prospects and upgrades
- Handle all pre-sale technical activities to be handed off to the CSM after the technical win stage
- Owns the technical side for responding to RFP/RFIs

### Customer Success Manager (CSM)

- Trusted strategic advisor to the customer
- Responsible for the post-sales customer journey
  - Onboarding
  - Success Planning
  - Stage adoption & expansion
  - Organizing Workshops and Enablement Sessions for Use-case Enablement & Expansion
  - Executive Business Reviews
  - Risk Mitigation
  - Renewal Discussions

- Leads regular cadence calls and keeps the customer collaboration projects up-to-date
- Maintains Gainsight records of customer health and triaging as necessary
- Account escalation point of contact for High or Critical Escalations

The Account Team Roles & Responsibilities Framework provided below is designed to serve as a guiding framework to improve collaboration across teams. Please note that this is applicable only to accounts managed by a CSM or CSA. For accounts managed by a CSE, refer to the CSE's Handbook on Rules of Engagement. Typically, an account will be overseen by either a CSM or a CSA, and the responsibilities outlined in the Account Team Roles & Responsibilities Framework apply equally, unless otherwise specified in the 'Comments' section.

The following Account Team Roles & Responsibilities Framework is an outcome of the work of this OKR and utilizes the DCI Matrix.

Account Team Roles & Responsibilities Framework - CSM/A

Journey Stage: Pre-Sales & Alignment

| Activities | Tasks | CSMA | CSA | AE | RM | SA | PS | Support | Executive Sponsor | Product and Engineering |
|-------|------|------|-----|----|----|----|------|----|-----|-----|
| Qualification | Leverage MEDDPPICC to ensure opportunity is qualified. |I||DRI||C,I|||||
| Technical Discovery | Continue the qualification of an opportunity coming from BDR/AE to understand whether there is a technical fit. |I||C,I||DRI||||
| Conduct a Day-in-the-Life of a Developer Session | Quick discovery to build recommendations for a developer's workflow. |I||C,I||DRI|||||
| Conduct a Reverse Demo Session | Quick customer-led demonstration of their development process or tools. |C,I - when there is already an engagement following an existing closed deal|C,I||C,I||DRI||||
| Run Custom Demo(s) | Series of tailored deep-dive demonstrations showcasing solutions to critical problems. |||||DRI|||||
| Deliver Hands-On Workshop/Labs | During the Tech Evaluation phase, we offer half-day pre-sales enablement workshops focused on GitLab use cases. These workshops are designed to accelerate user comfort and awareness for functionalities that have not yet been purchased, serving as an alternative to extensive Proof of Value (POV) exercises. |||||DRI|||||
| Conduct a Value Stream Workshop | Deep discovery to understand an organization's software delivery process and performance build current-state and recommended future-state for a specific application's value stream. |C,I - when there is already an engagement following an existing closed deal||C|I|DRI|I||||
| Run a Lunch and Learn Session | A short hands-on session showcasing areas of interest, performed in conjunction with the field marketing team. |C,I||||DRI|||||
| Oversee a Guided Trial | Customer-led trial that covers tactical features and implementation alongside guidance on how it aligns to assumed or known goals. |I - if the opportunity is fo