

opriately.

This flow is determined by:

1. the type of information,
1. its intended audience,
1. and timing sensitivity.

Furthermore, avoiding information overload is necessary to keep every stakeholder's focus.

To that end, we will have:

1. a dedicated Zoom call for all incidents. A link to the Zoom call can be found in the incident Slack channel posted in the corresponding incident announcement channel.
1. a Google Doc as needed for multiple user input based on the shared template
1. Incident Announcement channels for internal updates
1. regular updates to status.gitlab.com via status.io that disseminates to various media (e.g. Twitter)
1. a dedicated repo for issues related to Production separate from the queue that holds Infrastructure's workload: namely, issues for incidents and changes.

Incident Announcement channels

We have three dedicated incident slack channels where incidents are announced

- incidents : All incidents are announced here
- incidents-dotcom : All .com incidents are announced here
- incidents-dedicated : All Dedicated incidents are announced here

Status

We manage incident communication using status.io, which updates status.gitlab.com. Incidents in status.io have state and status and are updated by the CMOC.

To create an incident on status.io, you can use /woodhouse incident post-statuspage on Slack.

Status during Security Incidents

In some cases, we may choose not to post to status.io, the following are examples where we may skip a post/tweet. In some cases, this helps protect the security of self managed instances until we have released the security update.

- If a partial block of a URL is possible, for example to exclude problematic strings in a path.
- If there is no usage of the URL in the last week based on searches in our logs for GitLab.com.

States and Statuses

Definitions and rules for transitioning state and status are as follows.

State	Definition
Investigating	The incident has just been discovered and there is not yet a clear understanding of the impact or cause. If an incident remains in this state for longer than 30 minutes after the EOC has engaged, the incident should be escalated to the Incident Manager On Call.
Active	The incident is in progress and has not yet been mitigated. Note: Incidents should not be left in an Active state once the impact has been mitigated
Identified	The cause of the incident is believed to have been identified and a step to mitigate has been planned and agreed upon.
Monitoring	The step has been executed and metrics are being watched to ensure that we're operating at a baseline. If there is a clear understanding of the specific mitigation leading to resolution and high confidence in the fact that the impact will not recur it is preferable to skip this state.
Resolved	The impact of the incident has been mitigated and status is again Operational. Once resolved the incident can be marked for review and Corrective Actions can be defined.

Status can be set independent of state. The only time these must align is when an issues is

Status	Definition
Operational	The default status before an incident is opened and after an incident has been resolved. All systems are operating normally.
Degraded Performance	Users are impacted intermittently, but the impact is not observed in metrics, nor reported, to be widespread or systemic.
Partial Service Disruption	Users are impacted at a rate that violates our SLO. The Incident Manager On Call must be engaged and monitoring to resolution is required to last longer than 30 minutes.
Service Disruption	This is an outage. The Incident Manager On Call must be engaged.
Security Issue	A security vulnerability has been declared public and the security team has requested that it be published on the status page.

Severities

Incident Severity

Incident severity should be assigned at the beginning of an incident to ensure proper response across the organization. Incident severity should be determined based on the information that is available at the time. Severities can and should be adjusted as more information becomes available. The severity level reflects the maximum impact the incident had and should remain in that level even after the incident was mitigated or resolved.

Incident Managers and Engineers On-Call can use the following table as a guide for assigning incident severity.

Severity	Description	Example Incidents
~severity::1	 - GitLab.com is unavailable or severely degraded for the typical GitLab user  - Any data loss directly impacting customers  - The guaranteed self-managed release date is put in jeopardy  - It is a high impact security	

incident
 - It is an internally facing incident with full loss of metrics observability (Prometheus down)

Incident Managers should be paged for all ~severity::1 incidents| Past severity::1 Issues|
~severity::2	 - There is a recorded impact to the availability of one or more GitLab.com Primary Service with a weight > 0. This includes api, container registry, git access, API and web.  - GitLab.com is unavailable or degraded for a small subset of users  - GitLab.com is degraded but a reasonable workaround is available  - A widespread styling problem (frontend degradation) that makes the platform look unreliable.  - Any moderate impact security incident  - CustomersDot is offline Incident Managers should be paged for all ~severity::2 incidents	Past severity::2 Incidents
~severity::3	 - Broad impact on GitLab.com and minor inconvenience to typical user's workflow  - A workaround is not needed  - Any low impact security incident  - Most internally facing issues pertaining to blocked deployments (should a higher-severity incident be blocked by deployments, the severity for the blocker is still 3)  - CustomersDot is in maintenance mode	Past severity::3 Incidents
~severity::4	 - Minimal impact on GitLab.com typical user's workflow	Past severity::4 Incidents

Alert Severities

1. Alerts severities do not necessarily determine incident severities. A single incident can trigger a number of alerts at various severities, but the determination of the incident's severity is driven by the above definitions.

1. Over time, we aim to automate the determination of an incident's severity through service-level monitoring that can aggregate individual alerts against specific SLOs.

Incident Data Classification

There are four data classification levels defined in GitLab's Data Classification Standard.

- RED data should never be included in incidents, even if the issue is confidential.
- ORANGE and YELLOW data can be included and the Incident Manager managing the incident should ensure the incident issue is marked as confidential or is in an internal note.

The Incident Manager should exercise caution and their best judgement, in general we prefer to use internal notes instead of marking an entire issue confidential if possible.

A couple lines of non-descript log data may not represent a data security concern, but a larger set of log, query, or other data must have more restrictive access.

If assistance is required follow the Infrastructure Liaison Escalation process.

Incident Workflow

Summary

The entire incident lifecycle is managed through incident.io. All S1 and S2 incidents require a review, other incidents can also be reviewed as described here.

Incidents are reported and resolved when the degradation has ended and will not likely re-occur.

Incident Lead

The Incident Lead is responsible for ensuring that the incident progresses and is kept updated. This role is deliberately assigned after the start of an incident. The Lead should be chosen based on the type of incident, for example:

- Low Complexity Incidents: The team member most familiar with the affected system should lead (often the reporter)
- High Complexity Incidents: IMOC (for Sev1/2) or EOC (for Sev3/4) should typically lead due to coordination requirements, but product engineers and engineering managers are also capable of fulfilling this role.
- Delivery-Related Incidents: Release managers are often well-positioned to lead
- Security Incidents: Security team members should typically lead

Timeline

The incident Timeline is available on the incident in the incident.io web interface by changing "Highlights" to "All Activity" in the Activity section towards the bottom of the page. Items can be added to the timeline via the :pushpin: (·) or :star: (·) emoji reaction to a Slack post within the incident channel. If you react with the :pushpin:, a public comment will be left on the GitLab incident issue. If you react with a :star:, it will add an internal comment to the GitLab incident issue. Images attached to the Slack message will be added to the incident.io timeline, but will not be posted to the GitLab issue.

Labeling

We no longer use only GitLab labels to describe the status of an incident. The source of truth for any incident is incident.io. However, we do have incident.io set some labels based on the state of the incident.

Workflow Labeling

Label	Workflow State
~Incident::Active	Indicates that the incident labeled is active and ongoing. Initial severity is assigned when it is opened. This will be set when the incident is set to Active -> Investigating or Active -> Fixing
~Incident::Mitigated	Indicates that the incident has been mitigated. This label is applied if the incident status is set to Active -> Monitoring
~Incident::Resolved	Indicates that SRE engagement with the incident has ended and the condition that triggered the alert has been resolved. This will be applied when the incident is in the "Post-incident" or "Closed" stages of the incident lifecycle.

Other Incident Labels

These labels are added to incident issues as a mechanism to add metadata for the purposes of metrics and tracking.

Label	Purpose
-------	---------

| ---- | ----- |

| ~incident (automatically applied) | Label used for metrics tracking and immediate identification of incident issues. |

| ~blocks deployments | Indicates that if the incident is active, it will be a blocker for deployments. This label is set when the custom field "Blocks Deployments" in incident.io is set to yes. It is automatically applied to ~severity::1 and ~severity::2 incidents. |

| ~blocks feature-flags | Indicates that while the incident is active, it will be a blocker for changes to feature flags. This label is set when the custom field "Blocks Deployments" in incident.io is set to yes. It is automatically applied to ~severity::1 and ~severity::2 incidents. |

Duplicates

When an incident is created that is a duplicate of an existing incident it is up to the EOC to merge it with the appropriate primary incident.

Incidents can only be merged into open incidents, so if necessary you may need to briefly reopen the incident to merge.

Follow-up Issues

GitLab issues are created automatically when a "Follow-up" is created in incident.io. Any GitLab issue can be added as a Follow-up item by pasting the link into the incident Slack channel.

Follow-up items are created by default in the incident-follow-ups project and should be moved to the appropriate project after the incident is concluded.

Workflow Diagram

mermaid

graph TD

A(Incident is declared) --> |initial severity assigned| B(Active->Investigating)

A -. -> |If duplicate| Z(Merged)

B --> |"Fix identified"| C(Active->Fixing)

C --> |"Fix deployed"| D(Active->Monitoring)

D --> |"Incident resolved"| E(Resolved)

E --> |"S1 or S2"| F(Post-Incident Review)

E --> |"S3 or S4"| G(Incident Closed)

Near Misses

A near miss, "near hit", or "close call" is an unplanned event that has the potential to cause, but does not actually result in an incident.

Background

In the United States, the Aviation Safety Reporting System has been collecting reports of close calls since 1976. Due to near miss observations and other technological improvements, the rate of fatal accidents has dropped about 65 percent.

source)

As John Allspaw states:

- > Near misses are like a vaccine. They help the company better defend against
- > more serious errors in the future, without harming anyone or anything in the process.

Handling Near Misses

When a near miss occurs, we should treat it in a similar manner to a normal incident.

1. Open an incident issue, if one is not already opened. Label it with the severity label appropriate to the incident it would have caused, had the incident actually occurred. Label the incident issue with the ~Near Miss label.
1. corrective actions should be treated in the same way as those for an actual incident.
1. Ownership of the incident review should be assigned to the team-member who noticed the near-miss, or, when appropriate, the team-member with the most knowledge of how the near-miss came about.

SECTION: engineering/infrastructure/incident-management/incident-follow-ups.md

This page documents the process for managing follow-up items from incidents.

Terminology

- Follow Up Item: Action items, bugs, or improvements identified during an incident that need to be addressed after the incident is resolved.
- Incident Lead: The person responsible for coordinating the incident response and ensuring proper follow-up after resolution.
- Corrective Action: A follow-up item aimed at preventing similar incidents in the future (synonymous with InfraDev for labeling purposes).
- InfraDev: Infrastructure development issues that affect the production platform.

Default Issue Locations

We use different default projects for incident follow-up issues based on privacy requirements and team structure:

GitLab.com Team

- Default Location: gitlab-com/gl-infra/incident-follow-ups
- This project serves as the initial triage point for all follow-up items from GitLab.com incidents

Dedicated Team

- Default Location: gitlab-com/gl-infra/dedicated-follow-ups
- This private project preserves confidentiality for customer-specific incidents

Incident Lead Responsibilities

Incident Leads have the following responsibilities for managing follow-up items:

1. During the Incident:
 - Ensure all follow-up items are properly documented
 - Assign initial owners for each follow-up item where possible
2. After the Incident:
 - Review all automatically created follow-up issues within 1 business day
 - Verify that each issue has appropriate context from the incident
 - Ensure each issue is moved to the correct project for the team that will address it
 - Apply appropriate labels including team and group labels
 - Assign the issue to the correct team member or manager for triage

Issue Triage Process

For GitLab.com Incidents

1. Follow-up items are automatically created in `gitlab-com/gl-infra/incident-follow-ups`
2. To attach an existing GitLab issue to an incident as a follow up instead of creating a new one, paste the issue's URL into the slack channel
3. The Incident Lead reviews each issue and:
 - For infrastructure issues: Apply appropriate infrastructure team labels
 - For application issues: Move to `gitlab-org/gitlab` using the `/move` command
 - For service-specific issues: Move to the appropriate service project (e.g., `gitlab-org/gitaly` for Gitaly issues)
4. Apply appropriate group/team labels:
 - Engineering team issues should include the group label (e.g. `group::database`)
 - Infrastructure team issues should include the team label (e.g. `team::database`)
5. When moving an issue, ensure the `infradev` label is maintained on the issue

For Dedicated Incidents

1. Follow-up items are automatically created in `gitlab-com/gl-infra/dedicated-follow-ups`
2. To attach an existing GitLab issue to an incident as a follow up instead of creating a new one, paste the issue's URL into the slack channel
3. The Incident Lead reviews each issue and:
 - Applies appropriate team labels
 - Ensures the issue has a clear owner
 - Maintains confidentiality by keeping customer-specific information in the private repository
4. Only move issues if the Dedicated team requires a different location (for example, if it's an issue that needs to be addressed by a team outside of the Dedicated organization)

Tracking and Visibility

To ensure no follow-up items fall through the cracks:

1. Weekly Triage: The Production Engineering team performs a weekly triage of all open issues in the incident follow-up repositories to identify any untriaged items
2. SaaS Availability Reviews: The status of critical incident follow-up items is reviewed in SaaS Availability meetings
3. Engineering Teams: Engineering directors and managers are expected to regularly triage the infra dev issues assigned to their teams

Labels

The following labels are automatically applied to all incident follow-up issues in both repositories:

- infra dev: Applied to all follow-up issues
- corrective action: Applied to all follow-up issues

Additional labels that should be manually applied:

- Team labels (e.g., team::database, team::observability)
- Group labels (e.g., group::database, group::analytics)

Template and Issue Format

Follow-up issues contain:

1. A link to the original incident
2. A link to the Post Incident Review (PIR) if one has been created
3. Context from the incident
4. A description of the required action
5. Due date (if specified)

Incident Leads should ensure each issue contains:

- Clear acceptance criteria
- Sufficient context for someone unfamiliar with the incident to understand the issue
- Links to relevant logs, metrics, or other evidence from the incident

References

- Engineering Workflow - InfraDev
- Product Processes - InfraDev
- Incident Management - Corrective Actions

SECTION: engineering/infrastructure/incident-management/incident-manager-onboarding.md

Introduction

This page is meant to be the starting point for onboarding as an Incident Managers.

As a means to ensure a healthy Incident Manager rotation with sufficient staffing and no significant burden on any single individual we staff this role with Team Members from across Engineering.

An Incident Manager On Call (IMOC) has the following goals during a call:

- Identify/quantify impact to GitLab customers (metrics, customer support requests)
- Gather necessary folks to support area(s) of investigation/resolution
- Suggest politely that people not contributing leave the call, request that folks do so when the number of people in the call is a distraction (30 people as a somewhat arbitrary guideline).
- Validate recent releases and feature flags (can we roll back or change the flag?)
- Restate status for folks joining with a low level of shame (your possible misunderstanding is a chance for group clarity and learning)
- Drive to identify the fastest path to mitigation and interrupt as necessary
- Check in on the above on a regular cadence (S1 every 5 to 10 minutes)
- Communicate status to the right audience

Some of this may feel counter to GitLab Values; this is not designed or intended to diminish our values but to acknowledge and reinforce our need to mitigate customer impact as quickly as possible.

Incident Manager participants

The Incident Manager role will be staffed by all team members within this scope:

- Job Grades 8 and 9
- Development and Infrastructure departments
- All Staff Engineer, SRE, and Engineering Manager Job Families
- All specialties within the above Job Families (for example: Staff Backend Engineer, Staff Frontend Engineer, Frontend Engineering Manager, etc..)
- Not already engaged in another oncall assignment.
- Employed at the company for at least 3 months (same criteria for eligibility to join the Dev Escalation rotation)

As an Incident Manager, Team Members learn how we run GitLab.com and other GitLab SaaS environments. They help ensure the availability goals for GitLab.com by working with reliability engineers on call and development team members. The experience and awareness gained in this role leads to better understanding of building GitLab at scale and ultimately, a more reliable and scalable GitLab SaaS service and product.

Incident Manager Shifts

Shifts are 4 or 6 hours each at these times each day:

- 23:00 to 05:00 UTC
- 05:00 to 11:00 UTC
- 11:00 to 15:00 UTC
- 15:00 to 19:00 UTC
- 19:00 to 23:00 UTC

Each Incident Manager On Call assignment is for 4 consecutive days. The Incident Manager On Call shifts are based on individual preferences specified in the Incident Manager Onboarding issue. Shifts may include weekends based on this rotating assignment.

We have two scheduling layers implemented in PagerDuty, as a way to make weekend coverage more equitable and with the updated guidance for public holidays and weekend coverage.

1. Weekdays and Weekends - an Incident Manager's maximum 4-day shift can be scheduled any day of the week. This is how the rotating assignments have been scheduled to date.

1. Weekdays-only - an Incident Manager's maximum 4-day shift will only be scheduled on weekdays, which applies to team members who cannot cover weekend shifts. This layer has precedence over the "Weekdays and Weekends" layer, so that Incident Managers scheduled any day of the week will have their shifts scheduled less often.

1. Eligibility for weekday only shifts

1. Team members with United Kingdom, Germany and Ireland employment contracts dated before 2023-12-07, 2023-12-20 and 2023-06-06, respectively that do not wish to volunteer for weekend shifts and have not signed an amendment.

1. Team members with Austria, France or Italy employment contracts. Weekend shifts should not be taken due to local limitations.

1. Public Holidays should not be covered by team members with Switzerland or New Zealand employment contracts due to local limitations. If you are inadvertently scheduled for a public holiday please swap shifts

Special Coverage Days

It is more efficient to staff Incident Manager roles with fewer, but frequently longer, shifts during days where overall Team Member availability is lower due to Holidays or other global company events. These days will be designated as "Special Coverage Days" and volunteers will be sought to cover the relevant shifts. Examples of these days include New Years as well as the monthly Family & Friends days. Not all holidays will be covered in this manner as many holidays are local in nature and easier to cover simply by normal shift switches.

Special Coverage Days will be designated by creation of an Issue. The Issue will indicate the approach to scheduling for the specific event and will provide a table for volunteers to add their name. A DRI will be clearly noted and will be responsible for transferring the resulting coverage to PagerDuty.

During period with Special Coverage, indicated in the issue, Pager Duty schedule is void - incident managers are not expected to follow the PD schedule, however they are expected to collaborate DRI to provide full coverage for the given period.

Onboard as an Incident Manager

To onboard, create an Incident Manager onboarding issue.

Offboard as an Incident Manager

If your eligibility status changes or you have been exempted from Incident Manager On Call assignment, create an Incident Manager offboarding issue.

What to expect during your shift

Starting your on-call shift

Before your shift starts, verify your Slack alerts are working and your PagerDuty contact is up to date. Send a test page to make sure that you are receiving alerts correctly. You may get assigned to an on-call handover issue if your shift start time lines up with the start of the 8 hour SRE on-call shifts.

When your on-call shift starts, you will get notification(s) that your shift is starting (email or text, depending on your PagerDuty preferences). You will also get a Slack notification about being added to the @incident-managers user group.

When an incident happens

You are expected to respond to announcements in the incidents-dotcom Slack channel. Review the General guidelines for production incidents, verify the severity label on the issue, and update it if necessary.

When an S1 or S2 happens during a current incident, decide which incident has the highest customer impact and work on that issue. Reach out to infrastructure leadership escalation to get help with the lesser customer impact issue.

Ending your on-call shift

Before your shift ends, consider what tasks need to be handed over and make sure to proactively communicate them to the next on-call incident manager. This may mean paging the next on-call incident manager (through PagerDuty or the @incident-managers Slack handle) and bringing them into an active troubleshooting call, or it may mean providing instructions or context for situations you have handled that have some follow-up action. If there's a situation that occurred (or is actively happening) during your shift, it's your responsibility to prepare the next on-call incident manager.

Additional resources

- Reporting an incident
- IMOC Checklist
- IMOC Responsibilities
- Runbook IMOC
- Monitoring

Frequently Asked Questions

Who are the Incident Managers?

Incident Manager responsibilities are expected of Engineering Management and Staff+ Engineer job families in the Development and Infrastructure departments. Previously, Incident Manager responsibilities were fulfilled by a small group of managers in the Infrastructure department. All team members in these eligible roles are expected to participate in the Incident Manager

rotation.

In some cases, Senior Engineers may also participate as Incident Managers. This is particularly useful as Senior Engineers look for additional opportunities for growth in their role and as they prepare for future promotion opportunities. Senior Engineers may join the Incident Manager pool by opening an Incident Manager onboarding Issue and having their manager indicate approval in the Issue. As noted elsewhere, Team Members should only participate in one duty similar to this role, so should not also participate in Dev On-call at the same time (for example). When approving the request, managers shall be mindful of certain regions where a balance between Incident Manager and Dev On-call needs to be maintained because the eligible people overlap and both programs must be fully staffed.

Is this a volunteer effort?

No. All team members in eligible roles are expected to participate in the Incident Manager rotation. All team members who are eligible are to be scheduled into PagerDuty. Those who have upcoming shifts will need to prioritize completing their Incident Manager On-Call onboarding issues.

If anyone does not feel ready to become an Incident Manager, follow our process to swap shifts with someone else and work with your manager or the Incident Management Coordinator to resolve any challenges you are facing.

Note that we will consider exceptions on a case by case basis. Please talk to your manager if you do not think you can participate or would need an alternative schedule. Assignment and exemptions to participation will be coordinated by the VP of Infrastructure and the VP of Development

or their designees, and the requests require approval from the participants' reporting managers. This internal spreadsheet is being used to track exemptions from participating.

Development engineer: am I exempt from Dev On-call when I am on Incident Manager rotation?

Yes. If you are an engineer and you are in the Incident Manager rotation, you are exempt from the Dev On-call. You will not be expected to sign up for shifts in the monthly Dev On-call scheduling sheet. Please confirm that your email address is listed on the Development-Team-BE Google Spreadsheet, so that you are not auto-assigned a shift.

How do I prioritize Incident Manager duties with my other role expectations?

When you are scheduled to be Incident Manager On Call, this must be your 1 work priority. However, while it is the priority, there may not be any incidents. When there is not an active incident requiring Incident Manager leadership you maintain your normal work duties.

Currently, pages to Incident Manager On Calls happen roughly 0 to 4 times a week. During the time an Incident Manager is on call, there is a 15 minute SLA to respond to escalations. This can mean dropping out of meetings, 1:1s, etc. to join incident calls. This demand can make balancing hard, and you may have to reschedule meetings on short notice. Because of the interrupt-driven nature of incident management, it is advised that you shift sync meetings which would be difficult for you to leave to a later date.

In particular, please make sure to block out the time while you are on call in your calendar so

no interviews get scheduled during your shift. We need to be respectful of the candidate's time that they set apart to interview with us, an interview should not be interrupted by pages. If an interview does get scheduled during your shift, reach out to the recruitment team to reschedule.

What if I have family or other personal commitments during my shift?

The Incident Manager role requires a 15-minute response. However, there is no expectation that an Incident Manager is at their desk waiting for an incident during their entire shift.

You may have family or other short personal commitments which might complicate your response. In most cases, you should feel comfortable with any commitment that isn't longer than 45 minutes, as long as if during that commitment you can take a few minutes to engage in the incident slack channel to inform the rest of the team of your situation and timing.

If you have a standing commitment which is longer or does not allow for the informative response in slack it is suggested that you choose another shift. If you have an occasional commitment like this you shouldn't let that block you from a shift. One of the benefits of a larger team of Incident Managers is the ability to provide coverage for each other for unexpected or shorter duration commitments.

How often will I have an Incident Manager shift?

Four day shifts mean that, at the most, a Team Member participating in an Incident Manager shift with seven other IMs will have shifts approximately once every month. Each additional Team Members added to the pool in your region extends this timing by 4 days. The more Team Members involved, the less often any individual will have shifts.

How are Incident Managers scheduled?

Because new incident managers are being added to the rotation on an ongoing basis, schedules for upcoming months are not final and will shift as folks are added or removed from the rotation.

See the Incident Manager Coordinator responsibilities for details.

Can I shift my Incident Manager On Call days around?

Yes, this is one of the other benefits of having a well staffed pool of Team Members engaged as Incident Managers. Shift "trades" are easy to arrange in Pagerduty. It is your responsibility to ensure that your assigned shift is covered, but in extraordinary circumstances please reach out to the VP of Infrastructure for assistance. Swapping shifts is totally fine. Getting someone to cover for you for either planned vacation or a sudden/urgent family matter is something we should all do for each other.

What to do for covering a shift or asking for coverage:

1. Post to the im-general channel asking for help. Make sure to @mention people or @here if you are in an urgent situation. Let people know the days and times you will need help covering things.
2. Get the override scheduled in PagerDuty. Either person, the person asking or taking, can

put in the override.

Example 1, Scheduling yourself: Go to <https://gitlab.pagerduty.com/my-on-call/week> and click the shift for which you need an override. You should get a pop up which will let you pick the person covering you and the hours, which usually default to your whole shift.

Example 2, Covering for someone. Go to the schedule in PagerDuty and pick the shift and person you will override. You should get a pop up where you can pick yourself and the times for the override.

What role of Pagerduty shall I request?

A Professional Plus - Responder Role is sufficient to be an Incident Manager.
PagerDuty Roles reference

What if I am not available for my assigned shift?

Shifts are assigned based on the working hours that you selected during onboarding. Our current process is to swap shifts by asking for someone to take this shift in the im-general Slack channel.

What if I work a shift on a weekend or holiday?

When you are on-call, there is no expectation for you to be sitting at your laptop for the whole shift. You are required to be in a position to acknowledge a page and to join the Slack or Zoom call within 15 minutes. We still want to offer time-in-lieu for this disruption of making yourself available during the weekend and holidays.

Everyone is encouraged to explore options that work best for their personal situations in lieu of weekend shifts. When on-call you have the following possibilities:

1. Swap weekend days and weekdays.
1. Swap hours between weekend days and weekdays.
1. Take up to double the time off for any time worked during the weekend when the above two options don't work with your personal schedule.
 1. When an Incident Manager is in standby mode (e.g. not paged) during the weekend shift, they can take 1.25x time-off.
 1. When an Incident Manager is in call-back mode (e.g. being paged) during the weekend shift, they can take double the time-off.
 1. For those who reside in Australia, please refer to these guidelines of time in lieu in the handbook.
1. Other alternatives that promote work-life balance and have the least impact to your personal schedule.

Please create an OOO event in Workday and choose On-Call Time in Lieu.

It is important to abide by local labor laws and you are encouraged to understand if there are restrictions around your working time. The purpose of this information is to encourage you to take time off according to your schedule to account for the interruption to your weekend.

If you still have more questions, feel free to ask for guidance in the

im-general slack channel.

Can I expense my cell phone service used for on-call?

Yes, being on-call for incident management is considered vital to position for expenses purposes.

My direct report is serving as Incident Manager. What are my responsibilities as their manager?

Adding Incident Manager responsibilities to a direct report's role requires cooperation with their manager. The manager should support them when they are assigned Incident Manager duties, particularly when shifts overlap with weekends. Please make sure Incident Managers are empowered to delegate, postpone or drop some responsibilities given these additional responsibilities.

If the manager serves as an Incident Manager themselves as well, modelling similar behaviour is as vital as supporting their direct reports.

How do I see PagerDuty schedules from Google Calendar?

1. Locate the webcal feed for the Incident Manager - GitLab.com SaaS schedule in PagerDuty, which can be found in "Schedule Info". You can also link to your own Pagerduty shifts from My On-Call Shifts found in "Export Calendar".

1. Copy the webcal://... URL into Google Calendar by adding a calendar under the Other Calendars dropdown and select from url. Paste the webcal link.

1. Each Incident Manager On-Call shift should now show up as an event in this read-only Calendar view. Feel free to rename the Calendar to something legible (and not webcal://...), e.g. IMOC Shifts, My On-Call Shifts, etc.

Benefits of adding the PagerDuty IMOC schedule into Google Calendar:

- Search for yours or others' shifts, directly in Google Calendar.
- Block off your time by duplicating events from this calendar into your own, then set reminders for yourself accordingly.

How do I get notified when I am scheduled for on-call?

New schedules are announced in the im-general channel every month.

If you want to get notified when you are added to or removed from a schedule, or changes are made to your shifts, make sure to turn on the On-Call Boosters setting in the PagerDuty App on your phone.

Learning about Incident Management

We have a detailed handbook page that covers how we do Incident Management at GitLab.

SECTION: engineering/infrastructure/incident-management/tier2-oncall-onboarding.md

The Tier-2 SME on-call program enhances incident response by establishing a second tier of specialized support to complement the existing SRE EOC (Engineering On-Call) rotation. Subject Matter Experts (SMEs) provide domain-specific knowledge to help resolve complex incidents faster, improve MTTR (Mean Time To Recover), and increase ownership and accountability for service reliability.

Expectations

- Respond when paged by:
 - Tier-1 SRE EOC
 - Security
 - Self-Managed Support
- Assist with complex, domain-specific incidents that Tier-1 cannot resolve independently
- Maintain a 24/7/365 rotation with reliable coverage across time zones
- Acknowledge pages within 15 minutes and have a clear escalation path if the primary responder is unavailable
- Help define and maintain:
 - Escalation rules for your domain
 - Documentation and runbooks

Escalation Criteria

The Tier-1 Engineering On-Call (EOC) will perform initial triage and use available documentation before escalating to Tier-2 SMEs. However, timely escalation is encouraged when domain expertise is clearly needed, especially for critical incidents. The following guidelines help determine when escalation is appropriate:

By Severity Level

- S1/S2 Incidents: These typically require escalation when the Tier-1 team cannot resolve them independently. Due to their critical nature, Tier-2 SMEs should expect to be paged for these incidents when domain-specific expertise is needed.
- S3/S4 Incidents: These typically do not require escalation to Tier-2 SMEs during off-hours. However, Tier-1 may escalate S3/S4 incidents in specific circumstances:
 - When the customer impact is unclear and requires domain expertise to assess
 - When there's uncertainty about whether the issue might develop into a higher severity incident
 - When multiple lower-severity incidents combined create a potentially broader impact

Customer Impact Assessment

The primary consideration for escalation is actual or potential customer impact, regardless of the initial severity classification:

- If an incident has clear customer impact that Tier-1 cannot mitigate, escalation is appropriate even if initially classified as lower severity
- If Tier-1 needs help determining whether errors or unusual behavior in a service will affect customers, they may consult with Tier-2 SMEs

Onboarding Process for Teams

1. Define SME group
 - Identify team members who will participate in the rotation
 - Consider legal and employment restrictions for out-of-hours coverage
2. Create onboarding issue
 - Use the Team incident onboarding template
 - Document escalation criteria and severity levels
3. Set up on-call rotation
 - Configure schedules in incident.io
 - Ensure pager settings and notification rules are set for each member of the new rotation.
4. Document escalation procedures
 - Define escalation chain for unacknowledged pages (who is escalated to if no one responds within 15-minute period)
5. Complete training
 - Ensure all participants understand their role
 - Review incident response procedures

Next Steps

1. Double check your schedule to ensure coverage during holidays and team events

SECTION: engineering/infrastructure/platforms/_index.md

Mission

The Infrastructure Platforms section enables GitLab Engineering to build and deliver safe, scalable and efficient features for multi-tenant and single-tenant GitLab platforms (GitLab.com, GitLab Dedicated, and Self-managed).

Vision

To deliver on the mission, we are in the process of formalising the building blocks we need to work on.

```
<figure class="videocontainer">
  <iframe src="https://www.youtube.com/embed/Vui6iULzPw" frameborder="0"
allowfullscreen="true"> </iframe>
</figure>
```

Direction

Initiatives driven within the Platforms section, often spanning multiple quarters, are represented on the Infrastructure Platforms section epic.

How we work

Communication

Slack

Our main method of communication is Slack.

If you need assistance with a production issue or incident, please see the section on getting assistance.

Platforms

Channel	Purpose
infrastructure-platforms	We collaborate on department level items here. This channel is used to share important information with the wider team, but also serves to align all teams in Platforms with the common topic.
ginfrastructureplatformsleads	Communication for managers. Everyone interested is welcome to join this channel if they find the topics interesting.
confidential managers channel	Used to discuss staffing issues affecting all teams that require additional coordination. We default to using the public channel as much as possible.
infrastructureplatformssocial	Our social channel.

Dedicated

Channel	Purpose
gdedicated-team	Dedicated Group discussion channel. Please use this channel for discussions relevant to engineers across the Dedicated group
fgitlabdedicated	Dedicated function channel. Please use this channel to ask questions about features or ways of using the Dedicated product. Dedicated group will use this channel to make announcements relevant to wider groups
gdedicated-us-pubsec	Dedicated USPubSec team channel. Used to discuss topics that affect PubSec team only. For broader engineering discussions please use gdedicated-team
gdedicated-switchboard-team	Dedicated Switchboard team channel. Used to discuss topics that affect Switchboard team only. For broader engineering discussions please use gdedicated-team
gdedicated-environment-automation-team	Dedicated Environment Automation team channel. Used to discuss topics that affect Switchboard team only. For broader engineering discussions please use gdedicated-team
gdedicated-team-social	Dedicated social channel
dedicated-mr-review-stream	Visibility of new merge requests on Dedicated repos

Delivery

Channel	Purpose
gdelivery	Delivery Group channel
gdeliverystandups	
deliversocial	Social channel for the group.
releases	General communication about the current Release/Patch
fupcomingrelease	Detailed Release status / Release Manager channel
announcements	Release-Tools automation posts related to deployment activity

Production Engineering

Channel	Purpose
<code>#sproductionengineering</code>	General conversation for Production Engineering teams and requests coming in from other team members.
<code>#gproductionengineeringleads</code>	Channel for Production Engineering leads (staff+ and management)
<code>#ginfraops</code>	Team channel for Production Engineering Ops
<code>#ginfrasocial</code>	Social channel for Production Engineering
<code>#gfoundations</code>	Team channel for Production Engineering Foundations
<code>#gfoundationsocial</code>	Social channel for the Foundations team
<code>#gfoundationsalerts</code>	Non-urgent service alerts for Foundations owned services
<code>#gfoundationsnotifications</code>	Renovate notifications for Foundations owned projects
<code>#infra-terraform-alerts</code>	Terraform state drift alerts for SaaS infrastructure

Scalability

Channel	Purpose
<code>#gscalability</code>	General conversation for Scalability and requests coming in from other team members.
<code>#confidential-managers-channel</code>	Used for specific communication. We default to public channels.
<code>#gscalability-observability</code>	Team channel for general work in Observability.
<code>#gscalability-practices</code>	Team channel for general work in Practices.
<code>#scalability-social</code>	Our social channel.
<code>#scalability-id-project-name()</code>	We use project specific channels to keep make it easier to follow specific topics. Channel names follow this format.

The SaaS Platforms group is gradually directing requests for help to the `#saas-platforms-help` Slack channel.

This channel can be used if it is unclear which Infrastructure team the question should be directed to.

For more information, refer to the landing page for getting assistance.

The `#saas-platforms-help` channel is monitored by SaaS Platforms Engineering Managers and Staff+ engineers who triage any inbound requests. When triaging this channel, one should locate the team who can best answer this question and instruct the requestor to contact that team using the team's preferred contact method. When the requestor is connected to the right team, add a green check emoji to the message. Finally, if needed, update the getting assistance page with any changes.

Meetings

Once per week, we hold a Platforms leads call to align on action items related to career development, general direction or answer any ongoing questions that have not been addressed async. The call is cancelled when there are no topics added on the morning of the call.

In addition to the Platforms leads call, we have some recurring events and reminders that can

be viewed in the SaaS Platforms Leadership Calendar. Please add this to your Calendars to stay up-to-date with the various events.

Sr. Director of Infrastructure Marin Jankovski, likes to meet with new team members that join the organization. Marin sets up informal 1:1 coffee chats a few times a month with newer team members to get to know one another and see how they are doing. This process is organized by his EBA who will reach out to team members once he has the availability to meet. As this is a large team, it may take a while to get through everyone.

If someone needs to meet with Marin sooner than when the coffee chat is scheduled, you can reach out to his EBA Liki Simonot to set something up.

Grand Review

The Engineering Leads for each Stage, along with their Product Managers, hold weekly progress reviews to assess their groups' progress, share project updates, resolve blockers, and celebrate wins. Additionally, the Director of Product and the Senior Director of Infrastructure Platforms conduct a higher-level leadership review, where they go over summaries from these group-level meetings.

Weekly Schedule

- Wednesday: each Epic DRI updates the status section of their epics with the progress. It is important to surface:
 - risks and blockers impacting the project
 - projects that are completed, including a closing summary highlight. These epics will be closed during the Grand Reviews
- Thursday: Group Level Reviews conducted and added as threads in Infrastructure Platforms Top Level epic (see example)
 - Data Access: run by the Data Access Acting Sr. EM and the Group PM
 - Tenant Scale: run by the Tenant Scale Sr. EM and Group PM
 - Production Engineering: run by the Production Engineering Sr. EM and Group PM
 - Software Delivery: run by the Software Delivery Acting Sr. EM and Group PM
 - Developer Experience: run by the Developer Experience Director and a rotation of Product Managers
 - Dedicated: run by the Dedicated Sr. EM and a rotation of of Product Managers
- Friday: Leadership Review, run by the Sr. Director of Infrastructure Platforms and the Director of Product Infrastructure Platforms. Review the group-level summaries added as threads in the Infrastructure Platforms Top Level epic, then conduct a deep dive into one specific group to ensure comprehensive project coverage.
- Friday: Group level and the leadership level reviews are released together in infrastructureplatforms

The review is private streamed to the GitLab Unfiltered channel because the review covers confidential issues. All recordings are made available in the Platforms Grand Review YouTube Playlist

Infrastructure Platforms Leads Demo

The Infrastructure Platforms Leads Demo is an opportunity for sync discussions between Staff+ IC across the Infrastructure Platforms Group to highlight current ongoing efforts underway in

the teams they support.

All team members are welcome to join the call, but the emphasis is on Staff+ ICs to present and discuss the work they're focused on, the problems they're experiencing, and solutions they're considering.

The call is recorded to the Infrastructure Platforms Leads Demo Unfiltered Playlist. The agenda can be found in Google Docs.

While the intention is for the call to be made public on GitLab Unfiltered, the default is for it to be published as private.

At the end of the call, a quick vote is held between the attendees and if all agree that the content is SAFE, it can be published as public.

Requests for Help

On the landing page for getting assistance, we ask team-members who need assistance to raise Requests for Help using standard templates.

These issues are raised in the request for help issue tracker and are automatically assigned to the Engineering Manager of the relevant SaaS Platforms team.

The Engineering Manager is expected to:

1. Confirm that the question is not a duplicate and that the answer to the question is not already discoverable in the handbook or the tracker itself.
2. Confirm the urgency of the request.
3. Respond to the help request or assign to an engineer to help with the request.

Slack to GitLab Issue Tracker Integration

In an effort to enhance the tracking and resolution of requests directed to the Infrastructure team, we are evaluating a bot that converts Slack messages in infrastructurelounge channel into GitLab issues.

Workflow Overview

- Acknowledgement: An agent responds with the acknowledgedemoji (👍 in our case) to acknowledge a Slack message in the Infrastructure Lounge channel.
- Issue Creation: The Slack bot then creates an issue with the acknowledging agent assigned to it.
- Thread Attachment: The Slack thread corresponding to the message is also posted on the created GitLab issue.
- Label Assignment: Agents can further categorize issues by adding label emojis (ops, foundations, scalability-observability or scalability-practices) in the Slack message. This action automatically assigns the issue to the respective team: Ops, Foundations, Scalability-Observability or Scalability-Practices.
- Project Tracking: These converted issues are tracked under a dedicated project hosted at Infrastructure Lounge Slack Issue Tracker.
- Issue Closure: Agents/Requester can close the issue when resolved by adding any of the resolvedemojis (green-circle-check,whitecheckmarkor checkedin our case)

Configuration

Agents responsible for handling these issues are defined in a JSON file, which serves as a CI/CD variable. Currently, this file contains a static list of all members of the infrastructure department.

Project and Backlog Management

We use epics and issues to manage our work. Our project management process is shared between all teams in SaaS Platforms.

Tools

The Platforms section builds and maintains various tools to help deploy, operate and monitor our SaaS platforms. You can view a list of these tools in the Platforms Tools Index.

OKR

We use objective and key results to set goals in alignment with OKRs at GitLab. Our OKR process is shared between all teams in Platforms.

Hiring

Our hiring process is shared between all teams in Platforms.

Platforms Learning Path

All team members are encouraged to schedule time for personal development. The following links may help you get started with Platforms-relevant learning. Please add your own contributions to this list to help others with their personal development.

Learn about Platforms, and the Platforms Groups

Group Topic
----- -----
Platforms Product direction
Delivery Group Mission, Strategy, Team history
Scalability Group Mission, Strategy, Team history
Dedicated Group Mission

Learn about tools and technologies used within Platforms

1. Jsonnet tutorial
2. GitLab.com running on the Kubernetes platform

SECTION: engineering/infrastructure/platforms/hiring/index.md

Hiring in Platforms

Every eligible Platforms engineering team member should be qualified to conduct interviews:

All ICs should be qualified to conduct technical interviews up to and including their level

All ICs at senior level or above should be qualified to conduct peer interviews

Managers at EM level or above should be qualified to conduct hiring manager interviews

Senior Managers should be qualified to conduct Senior Leadership interviews

Interviewer Training Process

1. An interview training issue is created for the GitLab team member by their manager using this template

1. The team member completes the mandatory training identified in the training issue

1. The team member follows the steps in the Infrastructure Platforms Team Members section to add their details to the Infrastructure Platforms Shadowing Page

1. Hiring managers check the Infrastructure Platforms Shadowing Page when a new kick off issue is created to ensure shadowers are included in the interview panel

1. When shadowing is complete the team member adds their details to the relevant interviewer panel and assigns this MR to their manager to be merged.

1. The manager should update the relevant kickoff issues.

1. The team member closes their training issue

Hiring Manager Resources

Interviewer Panel is used to identify interviewers for new roles

Infrastructure Platforms Shadowing Page is referenced when a new kick off issue is created to ensure shadowers are included in the interview panel

Hiring Documentation & Resources is the SSOT for Platforms hiring resources including interview templates

SECTION: engineering/infrastructure/platforms/okrs/index.md

OKRs in Platforms

Creating OKRs

OKRs (or other items outside of projects) that require progress tracking should be updated every Wednesday.

When writing OKRs, the guidance is that:

Objective is defined as "What do you want to achieve?"

Key Results is defined as "How will you know when you've achieved the objective?"

As part of a KR, you can also have a sub point - which will likely tie to an epic. This would be an "Initiative", defined as "How are you going to achieve your key result?"

Objectives

The description for an Objective should have the following format:

markdown
Context

<context about the Objective and the problem it is trying to solve>

Linked Epics and Issues

<numbered list of relevant epics or issues>

If the description also needs other sections, these can be added after the block above.

Key Results

The description for an Key Result should have the following format:

markdown
Context

<context about how this Key Result contributes to meeting the goal of the Objective>

Scoring Criteria

<numbered list of epics or issues with the percentage contribution>

Keeping OKRs up to date

Objectives and Key Results should be updated every Wednesday.

For key results, the DRI must update the progress and the health of the key result.

For objectives, the DRI must update the health and provide a short status update.
The status update should be a one-line summary for each key result that belongs to the objective.

The progress of objectives should not be updated as this is achieved by automation.

The reason we do not have larger status updates is because we rely on the status information that is already present on the linked issues and epics.

We strive to keep the OKR process as lightweight as possible.

OKR Retrospective

At the end of the quarter, each team should perform a retrospective on the objectives.
There is no set format for the retrospective - it can be performed in a separate issue or on the objective itself.

The result of the retrospective is to update the description field of the objective using the format below:

markdown
Retro

Good

<2-3 bullet points summarising what went well during this OKR>

Bad

<2-3 bullet points summarising what didn't go well during this OKR>

Try

<2-3 bullet points summarising what we should try to do differently next time>

After the Retro section is finalized, the OKR can be closed.

SECTION: engineering/infrastructure/platforms/processes/_index.md

Processes

Process	Details
Calibration	Details and Schedule of Infrastructure Platforms annual calibration

SECTION: engineering/infrastructure/platforms/processes/calibration.md

Introduction

As part of our annual Talent Assessment process, each GitLab group performs calibration to ensure we are measuring performance consistently across teams. Calibration occurs in multiple phases, starting at the team level and then proceeding to group and finally stage level. The final calibration will include a People Business Partner who will enter our final calibrations.

Schedule

Session Number	Attendees	Calibration Focus	Session Date
-----	-----	-----	-----
Session 1	Team Manager	Team level calibration	self-organized, complete by mid February
Session 2	Group Managers	Group level calibration	self-organized, complete by end of February
Session 3	All Platform Managers	Informal stage level calibration	mid March
Session 4	All Platform Managers + People Business Partner	Formal stage level	

calibration | end of March |

SECTION: engineering/infrastructure/platforms/project-management/index.md

Project Management in Platforms

We use GitLab epics and issues to communicate the progress and status of our work. The Platforms epic is indexing the top level epic for each team, and links to active OKR's for a given quarter.

All teams in Infrastructure Platforms follow these guidelines so that it is easy for team-members to contribute to different projects if needed.

Projects are reviewed weekly in the Grand Review

Every Wednesday, the DRI for a project is expected to update the status block in the epic description to:

1. Indicate any project blockers.
2. Briefly highlight progress since the last update.
3. Indicate planned next steps, or mitigations required to progress.

This enables other engineers and other managers to have good information about projects in an asynchronous fashion.

Completed epics should remain "Open" with the ~"workflow-infra::In Progress" label. Update the status block in the epic description to summarize the project and share the completed status.

We use automation to collect the status information into top-level epics for each group, and to gather status updates from epic notes.

This automation runs several times each day, and can be triggered by running the pipeline listed on the project page.

Status updates on project epics

Comments are automatically generated on project epics to report status. To enable this functionality, please ensure the epic follows these steps.

When a project is finished

1. The epic should be left "Open" and retain the ~"workflow-infra::In Progress" label so that the project will still be listed for the next Platforms Grand Review.
1. The DRI is responsible for updating the epic description to ensure each section is up to date including the Participants section.
1. The DRI creates the closing status (final status update) which includes:
 1. The original problem
 1. A brief description of the changes made
 1. Information about the impact of the project now that it is completed

1. A link to the demo / walkthrough video
1. An instruction to direct the Grand Reviewers to close the epic
1. Emojis and / or shout-outs are also welcome but not required

Completed epics will be reviewed, celebrated, updated to done, and closed in the next Platforms Grand Review (Playlist accessible using GitLab Unfiltered account).

Note that closing statuses are helpful when producing the end-of-year team summaries.

Epics

Epics must always have the following sections:

- A Problem Statement
- Directly responsible individuals (DRI) responsible for the project completion
- Defined exit criteria
- Development Log, containing the previous Status Updates and any other relevant information to log
- Status Update, with the latest status update of the Epic
- Start date and estimated due date

Sections format

The DRI of the epic should be the assignee, and there should only be one assignee for an epic.

Use the following headings and structure in the Epic description:

Participants Section

This section is dynamic and includes people who are currently contributing to the Epic work. The DRI and the EM are responsible for updating the list of participants accordingly.

markdown

Participants

- @participant1
- @participant2
- @participantN

Status Section

The status section is generated automatically on sub-epics. See the section on status updates on project epics for more info.

Should this be an issue or an epic?

We use epics to show that a set of work can be combined under one theme to meet a specific goal or lead to a specific impact.

Epics can cover any set of work that is likely to take longer than two weeks, or has already

been running for longer than two weeks.

We can then report on our progress towards that goal regularly and clearly demonstrate the impact of our work when we are finished.

Specific Issue Types

Corrective Actions and Security issues

We also work on corrective actions and security issues labeled as ~"corrective action" or ~"security".

For these issues, severity:: labels are set to meet specific SLOs.

- For security issue see the Time to resolve table by severity
- Corrective Actions SLOs are currently based on definitions from Quality:

Labels

The Infrastructure Platform teams uses the following set of labels:

Description	Labels
The section label	devops::dataaccess group::GitLab Dedicated devops::gitlab-delivery devops::developer experience group::Production Engineering devops::tenant scale
The team label	group::durability group::gitaly group::git group::database frameworks group::database operations team::Observability team::runway team::Foundations team::Ops group::Self Managed group::build group::framework group::delivery-deploy group::delivery-release team::Environment Automation team::Switchboard team::US PubSec group::developer tooling group::development analytics group::feature readiness group::performance enablement group::test governance group::cells infrastructure group::organizations group::geo
Scoped workflow-infra:: labels	(see below)
Optional Scoped Service labels	Service::

Workflow labels

We leverage scoped workflow labels to track different stages of work.

They show the progression of work for each issue and allow us to remove blockers or change focus more easily.

While issues may not need to go through every state, the standard progression of workflow is from top to bottom in the table below:

State Label	Description
-----	-----
	Default label. Task is raised and effort is needed to determine the correct action, work required or team ownership
	Proposal is created following triage and put forward for a review. If there are no further questions or blockers, the issue can be moved into "Ready".

	Proposal is complete and the issue is waiting to be picked up for work.
	Issue is assigned and work has started. While in progress, the issue should be updated to include steps for verification that will be followed at a later stage.
	Issue has an MR in review.
	MR was merged and we are waiting to see the impact of the change to confirm that the initial problem is resolved.
	Issue is updated with the latest graphs and measurements, this label is applied and issue can be closed.

There are three other workflow labels of importance:

State Label	Description
-----	-----
	Work in the issue is being abandoned due to external factors or decision to not resolve the issue. After applying this label, issue will be closed.
	Work is not abandoned but other work has higher priority. After applying this label, team Engineering Manager is mentioned in the issue to either change the priority or find more help.
	Work is blocked due external dependencies or other external factors. Where possible, a blocking issue should also be set. After applying this label, issue will be regularly triaged by the team until the label can be removed.

Labels in gitlab-org group

Stage groups use type labels to label merge requests in projects in the gitlab-org group. If you need a stage group to perform work, it is best to apply the relevant stage group label when the issue is created.

Issue Trackers

This is a list of the group issue trackers:

Group	Tracker
-----	-----
	Dedicated issue tracker
	Delivery issue tracker
	Production Engineering issue tracker
	Scalability issue tracker

Triage Ops

For labeling and bot notifications under gitlab-com/gl-infra we use g