

d and adjusted throughout the quarter, in the following ways:

- Midway check-ins - a mini version of the global prioritisation process
- Monthly backlog trimming
- Re-adjustment after receiving an ad-hoc request that needs to be picked up

How to handle ad-hoc requests

It's normal to identify brand new research projects in an ad-hoc manner after prioritization happens, and we can still fold those new requests into the existing prioritized list. Here's how to do that:

1. The requester creates a research issue (template) and fills in the basic information.
 1. If the requester is a stakeholder in Product or Design, they assign the issue to themselves and the UX Researcher of their area. If the requester is a UX Researcher, they assign the issue to themselves and their Product and Design counterparts in relevant area.
 1. The stakeholders and the UX Researcher go through Step 2 - 4 as described above.
 1. If they decide to pick up the project, and the UX Researcher is at capacity, then they need to drop or off-load some current lower priority work and communicate that with relevant stakeholders.

Note: UXR leadership can be brought in to assist with the reprioritization as needed.

Maintaining the research priorities list

Throughout the quarter, the UX Researchers will maintain the global research priorities list, keep the status of the projects up to date, and document any changes and adjustments in the Quarterly UX Research Global Priorities sheet (Example in Google Sheets - internal access only).

Soliciting feedback on the process

Optional: Towards the end of quarter, the UX Researcher will open a retro comment thread to get feedback on how the process went for the team. This is an opportunity to reflect on the planning process, discuss how research went throughout that cycle, and identify any improvements to make in the next cycle.

Some questions you might consider in the retro comment thread:

- What went well?
- What did not go so well?
- Do you feel the most important research projects were prioritized and executed?

SECTION: product/ux/ux-research/resources-for-research-driv/index.md

Getting started

UX Researchers are one of the many GitLab Team Members who conduct user research. Other roles,

such as Product Managers and Product Designers, frequently conduct research with guidance from the UX Research team. All Research DRIs are supported by our UX Research Operations team to conduct efficient and high-quality research.

Who is the Research DRI?

- The Designated Research Individual (DRI) is the primary point of contact for overseeing and managing research projects.
- In GitLab, anyone can be the DRI of a UX Research project. If the Research DRI is not a UX Researcher, they are partnered with a UX Researcher acting in a supporting role.
- They ensure research efforts align with the project goals and stakeholder expectations.

Responsibilities of a Research DRI

- Act as the central coordinator for research-related activities.
- Define the scope and objectives of research initiatives.
- Ensure research outcomes address the key questions and provide actionable insights.
- Communicate findings effectively to stakeholders.
- Advocate for user needs and evidence-based decision-making in project development.

Research Process

How to Plan Research

Planning for research in GitLab is no different than in any other places. Research DRIs pair up with their corresponding UX Researchers on the following steps:

- Identify Goals: Determine what the research aims to achieve and define success criteria.
- Stakeholder Alignment: Collaborate with key stakeholders to clarify the scope and ensure alignment.
- Select Methods: Choose appropriate qualitative or quantitative research methods based on objectives.
- Create a Timeline: Establish key milestones and deliverables.
- Recruit Participants: Identify the target audience and develop a plan for recruiting research participants.
- Prepare Materials: Develop research guides, scripts, and data collection tools.
- Conduct Research: Execute the plan while ensuring ethical and unbiased practices.
- Analyze Results: Synthesize findings into actionable insights.
- Share Findings: Present results in a format tailored to the audience (for example - reports, presentations, blueprints, dashboards).

How to Set Up Research Projects the GitLab Way

GitLab is an async-first and remote-first company. We use the GitLab UX Research Project to track all research work and research operations efforts. Every project is a collaboration between the Research DRI, UX Researcher and Research Operations team.

Roles and Responsibilities by Phases

1. Scoping

- The Research DRI documents all context and goals in a Research Issue, and assign it to themselves, and their UX Research partner (if they are not UXR themselves).
- They then discuss with the UXR to assess and determine if research is necessary (more on how UXR prioritize our work).
 - If yes, they collaborate on the issue, set priorities, agree on support level, and due date.
 - If no, the UX Researcher will close the issue, and label it as "closed::will not do".

2. Drafting

- The Research DRI and the UX Researcher draft the research plan, the screener, and the script.
- All documents are peer reviewed, either by the UX Researcher partnering with the DRI, or another UX Researcher, if the DRI is a UXR themselves.

3. Recruitment kick-off

- Once the screener is peer reviewed, the Research DRI creates the recruiting request issue.
- The issue is then assessed by the Research Operations team on the recruitment complexity and assigned a recruitment grade.
 - If Grade 1, or Grade 2: The DRI can self-serve the recruitment by following a simple How-To guide provided by the Research Operations team.
 - Otherwise, the Research Operations team will set up and launch the recruitment.

4. Inviting Participants and Conducting the Research

- The Research DRI can monitor in real time whether the quota has been met directly in our recruitment tool.
- For Grade 1 or 2 recruitment requests, the DRI can invite participants for study themselves.
- For other recruitment requests, the Research Operations team does the invite instead.
- If cancellation or rescheduling is needed from our side:
 - Communicate early, explain the reasons and offer compensation if appropriate, for example when we didn't give enough notice.
 - Always ask for confirmation. If no reply, consider resending via multiple channels (e.g. email + SMS/WhatsApp), or attend the session at pre-agreed time to check.

5. Paying out

- Within 1-2 business days of a session or the study's completion, the Research DRI should inform the Research Operations team, and the UX Researcher they partnered with if they are not UXR themselves.
- If the DRI is a UX Researcher, they can pay out incentives for studies they lead.
- If the DRI is not a UX Researcher, the Research Operations team will pay out the incentives.

6. Analyzing and Reporting

- The Research DRI, and their UX Researcher partner if applicable, synthesise findings, validate key takeaways, and create the report.
- Though Google Slides is the most popular format, the report can really be in any form, for example, a doc, an issue, a board, or an diagram.
- It should be put through a peer review process in the uxresearchteamlounge channel. The peer

review may results in further analysis.

- The final report is shared in the uxresearchreports channel (using the template in the research issue itself), and any other relevant channels and sync meetings.
- It is expected that stakeholders and / or the Research DRI will follow up the report by identifying and issuing actionable insights for the Product and Design team to action on.

Useful to know

- A research project may be put in the backlog due to the team's capacity constraints. For more information on how we make that decision, see how UXR prioritize our work.
- A research issue may be closed / terminated before being turned into an actual project, if we agree that the research is not needed.
- The Research Operations team only support research recruitment requests of projects that have been vetted by a UX Researcher.

Templates and Tools

- Screening Question Repository (internal access only)
- Research Issue Template
- Recruiting Request Template

SECTION: product/ux/ux-research/rite/index.md

Rapid Iterative Testing and Evaluation (RITE) is a usability testing method where you evaluate a solution to a usability problem multiple times in a rapid and iterative manner. The goal is to identify not just usability problems, but also to react quickly on identified issues and test new solutions that cater to them. The end result of a RITE study is an experience (along with solutions to usability issues) that has been fully usability tested. This translates to an experience that we can be highly confident about shipping (in the context of usability), which helps to remove uncertainty around whether a proposed solution will be usable.

How a RITE study works

Conducting a RITE study is quite simple and is similar to conducting a traditional usability test. The differences are: 1) the sample size, and 2) immediately addressing issues as they are found. The diagram below illustrates the workflow to follow when conducting a RITE study.

To summarize:

- Step 1: Start with 3 participants
 - If issues are found, address them, and repeat Step 1 with 3 new participants
 - If issues are NOT found, proceed to Step 2
- Step 2: Add 2 more participants
 - If issues are found, address them, and start at Step 1 again with 3 new participants
 - If issues are NOT found after testing 5 total participants, the study is done!

Elements of RITE

In comparison to other usability testing protocols, RITE consists of the following elements:

- It's rapid and iterative: RITE tests are run regularly with a smaller number of participants. Each usability problem identified you aim to fix immediately and test again with a new set of participants. This process continues until no new usability issues are identified.
- Classification of issues: To move quickly with the test sessions, usability problems identified in a session are categorized into the following categories:

Category	Description
What to do	
----- -----	-----
----- -----	-----

A	Issues with an obvious cause and an obvious solution that can be implemented in the prototype quickly, such as text changes, re-labeling buttons. Implement solutions immediately and this will be the version for the next test session.
B	Issues with an obvious cause and an obvious solution that can't be implemented quickly/by the time of the next test session. Start working on the solution to address these issues and test in an upcoming session.
C	Issues with no obvious cause or due to other factors, such as task instructions. Collect more data in the upcoming session until they can be promoted to Category A or B.

- Domain knowledge and decision making: The decision makers of the team should be present or catch up on the session insights in order to quickly react on the learnings from each session. Domain knowledge is essential to estimate if the observed problems are likely to be a problem for others.

Roles and expectations to make RITE successful

RITE studies promise a rapid and iterative process for evaluating and improving design prototypes. It can only succeed if all contributors to the process are aware of their responsibilities and follow the schedule (timing and scheduling are critical to a successful RITE study). Make sure to inform yourself of your individual responsibilities and decide with your team what schedule best fits you.

- Product Designer: Make time to prepare a design prototype, set up a test in [usertesting.com](https://www.usertesting.com) based on the prototype, and then launch it. Watch the sessions as soon as they are available and analyze insights. Share your observations in the team sync and decide what changes need to be made with the Product Manager. Document all identified usability problems by placing them into the three categories (A-C). Based on the learnings, you need time to work on future iterations. In the example schedule below, we aim to test two iterations per week, which means during this week, the majority of your time will be spent on running RITE.
- Product Managers: While optional, make time to watch each research session as soon as a study gets picked up. Share your observations in the team sync, and decide what changes need to be made with the Product Designer for the next iteration. In the example schedule below, we aim to test two iterations per week.
- Engineering: Building great products is a team effort, and having exposure to real users regularly is crucial. While optional, it's recommended to make time to watch some or all of the

sessions and contribute your insights in the team sync.

Sample size

The sample size for RITE studies is both unique and critical to follow. It's fairly simple:

- N = 3 for each iteration and N = 5 to complete the process.

You only need 3 participants per iteration because as more iterations are made, the N increases, which provides that opportunity to capture additional usability issues. Over time, a larger number of users provides feedback on the product, and you fix the problems you identify as you go.

The RITE process is completed if you don't identify any new usability issues in an iteration. However, a sample size of 5 participants is recommended to increase the likelihood of identifying the majority of usability issues. Therefore, if your iteration doesn't bring in any new usability issues, add two more participants to ensure that all usability problems were identified.

Metrics

Since a RITE study is a usability testing method, the same measures that are used for usability testing are used for RITE testing: Effectiveness, Efficiency, Satisfaction, and Usefulness. The detailed measures to use can be found on the usability testing handbook page.

Number of tasks

There's no magic number when it comes to how many tasks to include. A guideline is to have 3 - 4 tasks to ensure that participants don't get tired and exhausted. Another aspect to consider is that unmoderated test sessions should be 15 - 20 minutes max and moderated sessions 60 minutes max. If you find your study may be too long for unmoderated testing, you might have to split your tasks across multiple usertesting.com sessions.

The key thing to remember when writing your tasks is that they reflect realistic user goals. Taking the JTBD into account when creating your tasks helps keep the focus on user goals.

A sample RITE study approach

A RITE study can take place in as little as a single day, or up to as long as it takes. The below example spans 5 days and is merely used to illustrate the process. If your team prefers a different pace, feel free to adjust the timeline accordingly. The example schedule is built to fit unmoderated usability testing using usertesting.com. If you prefer to run a moderated session, allow additional time for recruiting.

Day 1 - Preparation

Gather your team, and agree on a timeline on the following:

- When will the test session go live in usertesting.com?
- When will everyone watch the recordings?

- When will you input insights in Dovetail, and who will do it?
- How will you decide on the decisions and adjustments to make in the next iteration? For example, you might discuss async in Dovetail or use a sync meeting.

Prepare prototypes and the test script (this should be the same or similar to what you used for the baseline usability test).

Day 2 - Test Day

Run a pilot session, watch the recording, and adjust your test plan as needed.

Next, launch your test in [usertesting.com](https://www.usertesting.com) to three participants. Depending on how quickly the test comes back, analyze insights and organize identified usability problems using the RITE classification. Feel free to use a copy of this Figjam board.

Day 3 - Sync with team

Everyone in the team should have a chance to watch the three sessions and become aware of the identified usability problems. Share your learnings and decide on the scope for the next session. Adjust and prepare your prototype accordingly. Remember that Category A issues should be fixed before the new test round.

Day 4 - Test Day

Same as Day 2: Launch your test in [usertesting.com](https://www.usertesting.com) with three participants. Analyze insights and classify identified issues.

Day 5 - Sync with team

Same as Day 3: Everyone in the team should have a chance to watch the session and contribute to deciding how to act upon the learnings.

You can repeat this process in the same milestone or in an upcoming one, depending on how long it takes to find solutions for issues in Category B, as well as other activities that need to be taken care of. Document usability issues in Category B (Issues with an obvious solution that can't be implemented immediately) in GitLab for future planning of work.

The process is considered complete when usability issues are no longer found in a sample of 5 participants, and you're confident that your solution is user friendly, which may not happen within one milestone.

Frequently Asked Questions

Can I conduct a RITE study in a single day?

Yes, it's certainly possible. With careful planning, you can use [usertesting.com](https://www.usertesting.com) to complete a RITE study in a single day, whether it's moderated or unmoderated. The key is careful planning and scheduling.

Is it possible to make RITE testing my default usability testing method?

As long as the requirements are being met, of course! You will have more confidence around the usability performance of an experience tested through RITE than through a traditional usability study. Why? Because you tested your solutions along the way - and you know those perform well!

I don't have time to run a clean test with 5 participants. What should I do?

At that point, the test is no longer a RITE study and has the make up of a traditional usability study, where solutions that are addressed post-study are pushed out untested.

SECTION: product/ux/ux-research/strategic-research-at-gitlab/_index.md

What is Strategic research?

Strategic research is focused on answering future facing questions about users that help inform the long term product direction and roadmap. This kind of research can consist of one or more connected projects to better understand a larger topic area. Usually, the topic is broad enough to cover multiple stage groups and/or stages. Thus, strategic research can be high impact for the product teams and company as a whole.

{{< youtube "pYJEiC9Cs-A" >}}

We can use Strategic research to:

- Understand how to apply a solution across different territories or groups of people. For example, consider how a different industry, like lawyers curating legal documentation, could leverage source control management for versioning.
- Explore additional constraints to the current problem to make it more future proof -- for example, adding the number of concurrent users to support on a feature.
- Consider upstream and downstream experiences -- for example, enhancing the Verify experience for GitLab users by using Auto DevOps to help them get started more easily with Continuous Integration.
- Identify emerging market trends. For example, if customers are signaling an interest in their maturity for adopting DevOps, uncover ways can GitLab coach people in their maturity using Machine Learning or AI.

How does GitLab think about Strategic research?

Strategic research is research that is:

- Conducted to directly inform the strategy of a team and/or organization
- Aimed at answering broader research questions, not specific to a single stage area or feature
- More closely tied to problem validation than solution validation
- Independent from the monthly release cycle
- Focus is 6 months to 3 years ahead to best support the long term product roadmap

What it is NOT:

- Solution validation
- Specific to a single stage group (in most cases)
- Problem validation research tied to a specific issue/feature

What does Strategic research look like?

- Informs the UX and/or product direction
- Spans the user journey, rather than product group
- Seeks to understand and challenge current assumptions of feature functionality
- Engages cross-stage or multi-stage stakeholders
- Defines unmet needs
- Results in the addition, refinement, or removal of personas

How is Strategic research measured?

- Success will be measured by the impact of the research, at the project level. As strategic research matures at GitLab, we will look for common measures of success that could be applied to all strategic research projects.

Examples of Strategic research topics:

- Understanding stage usage trends
- Expanding understanding of our customers' motivations, goals, challenges, and needs
- Identifying our customers' biggest IT initiatives/goals
- Understanding how GitLab can overcome DevOps adoption pains
- Understanding more about cross-stage use cases and workflows
- Identifying ways Ops can drive adoption across multiple stages and stage groups
- Creating Actionable Insights connected to cross-stage group issues to be scheduled in the next 1 year

Examples of Strategic research from recent projects:

1. Release Manager Persona Addition was a great example of a cross-stage and multi-group strategic research effort that resulted in an update to the Roles and Personas page. This positively impacted the entire Ops Section and helped us reach new audiences in the Ultimate tier.

1. KubeCon Ops Product Direction Survey was a large scale project that was able to inform the future product roadmap within Ops. This was done through obtaining valuable feedback about what CI/CD workflows are most challenging for GitLab users and learning what preference users had in mind for their ideal secrets solution.

Strategic research outputs

The outputs from Strategic Research can be found in issues, changes made to documentation, and in published changes to a category direction page (which includes 1 year and 3 year visions), such as:

- GitLab Stage Group 1-year Plans
- GitLab 3-year strategy
- GitLab's Direction Personas

Resources for Strategic research

- Developing Research Hypothesis
- New Persona Template

SECTION: product/ux/ux-research/strategic-research-at-gitlab/gitlab-adoption-research-program.md

What is GitLab Adoption Research?

As a solution, GitLab is considered a DevOps Platform. The DevOps platform is a single application that eliminates complex integrations, data chokepoints, and toolchain maintenance, resulting in greater productivity. This single application is made up of our GitLab stages. In order for us to better understand DevOps Platform adoption, we must understand how and why users are adopting more stages.

In our research on adoption, we have learned new namespaces adopting 3 or more stages have significant differences from those namespaces not adopting multiple stages. First, the namespace has a significantly higher free to paid conversion rate, and the more stages a customer uses, the deeper the engagement, the stickier GitLab becomes, which also makes it the harder for customers to leave. Additionally, more stage usage creates more monetization opportunities in terms of upsell as well expansion for our business.

To better serve our DevOps Platform Solution and users of our end-to-end capabilities, we have created this research program.

How does the GitLab Adoption Research Program work?

This program DRI is Jackie Porter, Director (@jreporter) and is a Strategic Research Initiative.

What are the GitLab Adoption Research Program outputs?

1. Readouts of adoption research
1. Organization-based themes for Persona updates
1. 1-3 year time horizon direction page updates

Current GitLab Adoption Research

1. Application Types Impact on DevOps Adoption

Past GitLab Adoption Research

1. Facilitate Cross-stage Adoption - Issue
1. Reasons for cross-stage adoption - Issue
1. SMB Pursuit of the Golden Journey

SECTION: product/ux/ux-research/surveys/_index.md

Surveys are a powerful tool which outputs qualitative and quantitative data at scale. They are fairly easy to execute, but at the same time one of the most difficult methods to execute correctly.

Surveys are not unique to UX; they are used across many disciplines. For UX Research, we use surveys to collect data from users (or future users) on their demographics, attitudes, opinions, needs, and behaviors.

The two types of survey questions

We can classify the types of survey questions into two groups:

1. Quantitative survey questions

- The goal of quantitative survey questions is to understand data/a question around how many people. Examples include: a satisfaction rating, if users prefer one option over another, what features users prefer, etc.

- You analyze quantitative surveys by counting the number of people who selected each option presented and then use descriptive statistics.

1. Qualitative survey questions

- The goal of qualitative survey questions is to understand more about how people feel behind a given research topic. Examples include: how users feel about a particular experience or any type of open-ended question.

- You analyze qualitative survey questions by tagging the responses with themes and then counting the themes, similar to a thematic analysis.

It's sometimes useful to include both quantitative and qualitative survey questions within a single research effort. For example, you can start with qualitative surveys to establish an understanding of the scope of the topic and inform survey question choices to use within a quantitative survey. Or, you can use a quantitative survey question and then follow up with an open-ended qualitative survey question to understand why they selected their answer.

When to use a survey

When determining which research method to use, you should start with your research questions and envision how those can be answered, given the available research methods.

Perhaps the most important item to consider when determining if you want to use a survey: if your research questions require open-ended questions to answer them, then using a survey probably isn't the best method for you. That's because too many open-ended questions lead to a variation in responses and often a longer survey. Sometimes respondents don't give enough depth or breadth and get fatigued with open-ended questions. Also, this structure of survey can be time intensive to analyze. If your research questions don't require open-ended questions, then a survey could be a good method for you.

Quantitative research questions	Qualitative open-ended research
questions	
:-----	:-----

How satisfied are GitLab users with a given feature?	Why are users either satisfied or unsatisfied with a given feature?
Which DevOps workflow is the most common?	Why are different DevOps workflows challenging?
What feature is the most popular?	How will different features impact user workflows? Why?

Additionally, a you might use a survey if:

- You need a large number of respondents to answer your research questions with a high degree of confidence. This may include validating insights from a qualitative research study to make decisions more confidently.
- You need your research questions answered in a relatively short amount of time
- You want an efficient way to achieve the above
- You don't need deep qualitative data
- Your research questions can be answered through quantitative data

If you're unsure if a survey is right for your needs, it's suggested that you get crisp on your research questions to understand what kind of data is needed to answer those questions. That exercise will often make it easier to reveal the best research method to use.

Pros and cons of a survey

A survey comes with pros and cons, which you should consider before committing to the method.

Pros:

- Can be completed by respondents at their own convenience
- Low cost relative to the number of respondents
- Able to reach a large number of respondents
- When designed well, can be efficient to analyze
- Time-efficient in terms of sampling a large number of respondents
- Provide a standard response from respondents (in other words, data comes in for each question in the same format), making analysis more streamlined when looking for patterns

Cons:

- Difficult to get right because users can interpret questions in so many ways and we can introduce biases when writing survey questions
- Easy to get wrong
- Can introduce biases when writing survey questions
- Often chosen as the wrong research method
- Usually have a low completion rate, requiring more effort around sampling/recruiting
- Some questions are sometimes skipped by respondents
- It's easy to send out a survey to a large group without knowing if it is going to the "right" people

Self-reported data

Surveys rely on self-reported data, which is important to note. As a guideline, if we already

have that information, we shouldn't ask it; as it's a tax on the respondent to answer those survey questions and can be less accurate. Examples of information we may already have are: license type, user tenure, or company size. Additionally, respondents can also make mistakes or possibly lie, making the responses less reliable. In cases when we have that information, we should utilize it. Or, if we don't have that information, and want to avoid self-reported data, we should be using a different research method to obtain it (example: a user interview where you're observing behaviors).

How to conduct a survey

If you've decided that a survey is the right method to use, you can follow the steps below to conduct a survey:

1. Define goals, objectives, and hypotheses
1. Determine your sample and sampling method
1. Design your survey
1. Peer review
 - Have at least one other person review your survey. It's ideal if they're experienced in conducting surveys and know to look for common pitfalls.
1. Pilot your survey
1. Create a recruitment request issue to distribute your survey
1. Analyze the data

SECTION: product/ux/ux-research/surveys/designing-your-survey.md

Designing a survey isn't easy. It requires thoughtful planning to make sure the flow of the survey is easy to follow, the questions are clear to the respondent, and that the output will yield accurate data.

Here are some of the more important tips to consider when designing a survey:

1. Only include survey questions that directly relate to your research questions. It can be exciting to learn about your users and survey writers have a tendency to include questions that they are curious about, but that won't impact their product decisions or planning directly. Best practice is to start with a planning table where you list your research questions in one column, the survey questions that you'll use to answer your research questions in another one, and what you plan to do with your data in a third column. This will help you to think through whether or not your questions are nice-to-have or necessary in order for you to answer your research questions.
1. Ensure the questions are clear and cannot be easily confused. Since surveys are unmoderated, you won't be there to provide clarification. That means that your survey and accompanying questions need to be simple and clear to your respondents. A great way to determine this is by piloting a survey through UserTesting.com, where you can watch respondents think aloud while taking your survey. It's also ok to provide definitions to terms when it makes sense.
1. Use neutral language. Avoid positively or negatively worded questions to reduce any bias in your results. You'll be surprised how easy it can be to inadvertently write your questions as leading.
1. Make sure each question covers a single topic. Avoid double-barreled questions by breaking

out questions that are really 2+ questions in 1. Not only are double-barreled questions difficult to answer for respondents, they're hard to analyze and act upon.

1. Use standard survey questions. Luckily, surveys have been around for quite some time, resulting in standardized questions that are commonly used in the industry such as the System Usability Scale. Before spending time designing a survey, check to see if there are standardized survey questions available that may help in answering your research questions. If so, use those rather than reinventing the wheel.

1. Avoid too many open-ended questions. While you may be tempted to ask many open-ended questions within a survey, there are some factors to consider:

- Limited detailed data - The data you obtain through open-ended questions may result in limited detail, and it certainly won't be as rich as data obtained from a user interview. In short, respondents sometimes aren't willing to go into detail when typing out responses.
- Left with more questions than answers - Related to the above, you may not end up with the answers you're looking for if you rely on open-ended survey questions to answer complex research questions.
- Open-ended questions are often skipped - Filling out an open-ended survey question takes time and can be viewed as an annoyance by respondents. Because of that, they're often skipped.
- Fatigue - Too many open-ended questions can quickly lead to fatigue, resulting in survey abandonment.

Now, sometimes it's necessary to ask open-ended survey questions. When doing so, just be mindful of the above factors, and explore other survey question types first. If you end up with a lot of open-ended questions, reconsider your methodology. It may be more appropriate to start with an interview, and then run a survey to get a larger volume of answers.

1. Avoid asking questions that force the respondent to estimate because it's hard to estimate accurately. What helps is to provide a designated time frame when asking common questions, like "How many times have you XYZ?" Without a scoped time frame, each participant will come up with their own and provide an estimate. Instead, be crisp about a scoped timeframe for the user to think back about. This will result in more accurate data across a consistent time frame.

1. Make sure the survey flows well. Ask yourself if the ordering of the questions makes sense. Questions about the same topic should be grouped together to avoid context switching. A best practice is to start with your most important questions at the beginning of the survey, in case respondents drop off.

1. Think about any built-in logic you're using. Using logic can help make a survey cleaner, appear shorter, and add some useful capabilities. Logic can sometimes negatively affect the flow by negatively altering it, so make sure the flow makes sense for every logic-based scenario.

1. Don't make all the questions required. You may think that all questions should be required, since they're all important to you. Doing so could result in survey abandonment or inaccurate data on some questions that respondents would otherwise like to skip. Only flag required questions that are critical to your research questions and/or survey logic.

1. Give respondents options that feel right. For multiple-choice questions, provide a way for respondents to answer with 'other' or 'not applicable'. This is so you don't force respondents to answer a question inaccurately, which can result in inaccurate data. A bonus: you may learn something about your respondents if you see a significant amount of them answering with 'other' or 'not applicable'.

1. Keep it short. Respondents tend to complete shorter surveys faster than longer surveys. If your survey is more than 20 questions, it's probably too long. It's strongly advised to keep the scope of the survey to only answer the research questions, and eliminate any 'nice-to-have'

questions. You'll be surprised at how often 'nice-to-have' questions will try to make their way into your survey. If you don't have an immediate need for the data of a given survey question, eliminate it. You'll save yourself a lot of time in the analysis, too.

1. Let respondents know how much is left. It's important to let respondents know how much of the survey is left to complete. You can do this by showing a progress meter or a percent complete indicator that's already built into the survey tool you use.

1. Think carefully before using an embedded survey question in an email. Consider the potential advantages and drawbacks of embedding survey questions directly into email distributions. While it may offer convenience and increased visibility, it can also present challenges such as response bias and limited formatting options. For instance, in our previous experience with embedded surveys in our paid NPS surveys, despite achieving a good response rate, we encountered a notable increase in one-click detractors (respondents who selected the answer to the question in the email, but did not complete any other questions within the remaining survey questions). This could indicate a tendency for respondents to provide quick, reflexive responses without fully considering their feedback, possibly due to the ease of responding within the email.

Here is an example of an embedded survey question:

SECTION: product/ux/ux-research/surveys/how-to-identify-low-quality-data-when-designing-and-reviewing-your-survey.md

A major challenge in UX Research is ensuring data is representative of the users you are trying to sample. Unlike with interviews, where you can assess the quality of a response by talking to a person for a set period of time, survey data is harder to verify because it is usually anonymous and there is a greater quantity of individual responses to parse through. When conducting a survey, an important goal is to limit the amount of low quality data (for example: data provided by bots or AI instead of humans, false/inaccurate data provided by humans), so that you can be more confident in your survey results. This handbook page covers various ways you can improve your survey design and analysis to make you more confident with the quality of data in your final dataset.

Ways to design your survey to identify invalid survey responses

Attention checks

Attention check questions are designed to ensure participants are paying attention and responding thoughtfully. Consider including an attention check question to more easily filter during cleaning and analysis.

Example: "Please select 'agree' for this question to confirm you're paying attention." If the respondent selects 'agree', this indicates they are reviewing and responding to questions carefully, so this respondent shouldn't be removed.

Example: "The following question is to verify that you are a real person. Please enter the word 'Red' into the box below." If the respondent responds 'Blue', this suggests they are likely not paying attention and should be removed from the dataset.

Attention checks should be used strategically. They are best suited for complex surveys, long

surveys, and surveys with a large number of multiple choice questions. They are not suited for short surveys, surveys with a large number of open responses, and screener surveys.

Bot detection and fraud detection

We suggest enabling settings for bot detection AND fraud detection in Qualtrics. Comparing scores from both survey settings within the data will help you determine whether or not to filter out participants.

Example: If a respondent has a bot detection score of 0.2 and a fraud detection score of 100, this suggests responses are problematic and could be removed.

Example: If a respondent has a bot detection score of 0.9 and a fraud detection score of 10, this suggests the feedback is legitimate.

Survey Setting	Column Name in Qualtrics Data	Minimum Value	Maximum Value	Interpretation of Score
----------------	-------------------------------	---------------	---------------	-------------------------

Bot Detection	QRecaptchaScore	0	1	Score of less than 0.5 means the response is likely a bot.
---------------	-----------------	---	---	--

Fraud Detection	QRelevantIDFraudScore	0	130	Score greater than or equal to 30 means the response is likely fraudulent and a bot.
-----------------	-----------------------	---	-----	--

What to look for when cleaning survey data

Incomplete survey responses

Incomplete responses can skew data analysis and interpretation. It's essential to define what constitutes an incomplete response. For instance, if a survey has ten questions and a participant only answers three, you may want to remove that response.

Example: A participant starts the survey but drops out after the demographics section, leaving the rest of the questions unanswered. This response should probably be dropped.

Example: A survey respondent completes the NPS question, but fails to provide an open-ended response. This response may not need to be dropped.

Participants that don't match our desired criteria

Filtering out participants who don't meet the criteria ensures that the data collected aligns with the survey's objectives. Define your criteria clearly to facilitate this process.

Example: In a survey aimed at paid users, responses from users on a free plan need to be excluded.

Example: Screener survey data from respondents who do not initially meet study criteria should be kept in case the criteria needs to be adjusted later on, so respondents who might not be a perfect fit could participate.

Responses that are completed too quickly

Rapid completion may indicate that participants are not engaging thoughtfully with the

questions, potentially leading to unreliable data.

Example: A survey with several complex questions that is completed in less than 30 seconds should be considered suspicious and a candidate for removal.

Straightlining or other pattern based responses

Straightlining occurs when respondents consistently choose the same response option without considering the question. It can indicate inattentiveness or a lack of genuine engagement.

Example: A respondent selects "strongly agree" for every question in the survey. There is no perfect solution to straightlining or other pattern based responses. If there is only one or a small number of these responses it is best practice to remove those participants because their data is untrustworthy. If you find a large number of these kinds of responses in your data you can adjust your survey to avoid matrix style questions with too many options or to present one question per page.

Example: A respondent who picks a mix of "strongly agree" and "agree" throughout the survey shouldn't be excluded.

Inconsistent responses

Inconsistencies may indicate misunderstanding, lack of attention, or even intentional deception. Investigate and reconcile conflicting information.

Example: A respondent states they only use GitLab, but mentions performing pull requests (GitHub terminology) in their role instead of merge requests (GitLab terminology). This feedback suggests they are not a GitLab user and could be removed from the dataset.

Example: A respondent states they use GitLab and have experience doing merge requests in their groups and projects (GitLab specific terminology). This response is consistent with someone who has used GitLab, so they should remain in the data.

Nonsensical or suspicious (i.e. AI generated) open ended feedback

Open-ended feedback should be coherent and relevant. Nonsensical or suspicious responses may indicate automated or fraudulent submissions.

Example: A response consists of random characters or does not address the question asked might be considered feedback left by AI. This would suggest they are a candidate for removal.

Example: A response that directly answers the question and sounds like a human wrote it would likely be considered non-AI generated feedback, so the respondent should not be excluded.

Common AI generated phrasing that could suggest fraudulent or insincere responses:

"I apologize for the inconvenience. Unfortunately, I am unable to provide real-time data at this moment."

This response demonstrates typical AI behavior by offering a polite apology but failing to provide a direct answer and specifically mentions real-time data. AI-generated responses often default to formal language and may avoid committing to real-time interactions, leading to vague or evasive responses.

"As of my last knowledge update..."

This response displays AI's tendency to rely on qualifiers and pre-programmed responses. AI-generated content often includes disclaimers or phrases indicating reliance on outdated information, reflecting the system's inability to provide real-time updates or nuanced understanding.

"Certainly, here is a summary of the requested information."

This response exhibits characteristics typical of AI-generated content by using formal language and offering to provide information in a structured format. AI-generated responses often prioritize delivering data or summaries in a straightforward manner, lacking the spontaneity and contextual understanding inherent in human communication. The use of the word "certainly" followed by a predetermined action (providing a summary) suggests a scripted response rather than a genuine human interaction.

SECTION: product/ux/ux-research/surveys/piloting-your-survey.md

When you feel your survey is complete, you're ready to pilot it. This means sending it out to a small number of people to understand how respondents feel about it and whether they're able to understand the questions as you hoped they would.

Piloting a survey is an important step for many reasons:

- You can learn if your respondents are interpreting your survey questions as you designed them.
- You can identify any bugs or logic issues.
- It saves you time, if you end up having to re-run the survey due one of the above issues.
- If you have a small population to sample from, you can't afford to re-run a survey due to preventable errors.

When piloting a survey, you're essentially running a qualitative research study on your own survey, which is interesting in itself! Since you're looking for qualitative feedback, that means you don't need a large sample size to run your pilot. It's recommended to aim for 5-8 respondents in total, depending on the importance of findings, and the length and complexity of the survey.

There are two recommended ways to pilot a survey:

1. Unmoderated - With an unmoderated pilot, you can use UserTesting.com to have qualified respondents take the survey, with instructions to think aloud while giving their own interpretation of what each survey question is asking of them. Each participant session is recorded, so you can watch the results at your own pace. This is the most time-efficient option.
1. Moderated - With a moderated pilot, you essentially conduct 1:1 interviews with qualified respondents taking the survey, with instructions to think aloud while giving their own interpretation of what each survey question is asking of them. The nice thing about a moderated approach is that you're present to ask follow up questions (example: 'What do you think would be a better way to ask that question?'). This option takes more time to plan, schedule, and execute on, but it does allow you to gain further insight and clarification.

Upon completion of a pilot, you'll probably need to make some iterations to the flow, wording,

and/or options within your survey. Be sure to run a follow-up pilot to ensure your changes are fully understood; don't assume they will be! Keep piloting until you find your pilot respondents all understand the questions, wording, and flow of the survey.

SECTION: product/ux/ux-research/surveys/sample-sizes-for-surveys.md

In a perfect world, we'd sample every single user within the population. However, we can't do that. Why? It's inefficient, expensive, time-consuming, and probably impossible to pull off. Instead, we try to attain a representative sample of the population from the population we want to focus on. To do this, we take a slice of that population, referred to as our 'sample'.

Furthermore, we compute measures from the sample and use them as estimates of the measures for the entire population. For example, we can't measure every users' SUS score, even though that true SUS score does exist. Instead, we compute the SUS score from a sample of users. While it's not going to be the exact SUS score for the population, it's close within some degree of error (or margin of error).

Factors to consider when determining your survey sample size

One of the most common questions around sampling is 'how many people should we sample?' The correct answer is: 'It depends'. There are several factors to consider when determining how large your sample should be. The main factors we look at are: confidence/accuracy, who you are sampling/how many people you sample, population, and the type of analysis. Further, not all studies need a high level of accuracy, but the information is still very valuable/useful, allowing you to make informed decisions. Here's what we mean:

- Scenario 1: If your research questions require a high degree of accuracy, such as a granular measure, you'll want to sample more people. Example: SUS (n=200)
 - Why? SUS is measured over time, because we want to pick up on changes. When we see changes in SUS scores, we want to be confident that it is due to a true change and not because we were lucky or unlucky with our sampling. Smaller samples can be more sensitive to errors.
 - SUS is measured across the entire population of users; therefore, we have a much larger pool of people to sample from. To ensure we have accuracy in the sample, we need a larger sample size.
 - Company-wide goals are set around SUS. Thus we want a high degree of accuracy. As the sample size increases, your error (or margin of error) in other words, your likelihood of lucky or unlucky sampling goes down, and the observations you make with your data are more accurate.
- Scenario 2: If your survey requires a larger sample, you may run the risk of oversampling companies with more GitLab users/licenses (i.e., large enterprises). Example: CSAT (n=400)
 - Why? Based on the results of our survey, we want to be able to say that our survey sample is representative of our larger population. If more than 1% of responses in a large scale sample (n > 500) come from a single company, then we have lower confidence in our results because that company is being over-represented in our sample compared to how they would be if we sampled our entire user base (i.e., population).
 - Note: These guidelines do not apply if you are targeting a very specific or hard to reach user group.
 - When trying to maintain the balance between accuracy and sampling, check your survey email distributions to limit how many individuals you target from a single company. You can do this

by placing limits on how many people you query from a single company email domain or by filtering out specific email domains between email distributions through tools like Rally and Qualtrics. That way, your results are less likely to be influenced by those groups of users.

- Scenario 3: If your research is trying to understand how many people are going through a certain workflow, or how people would rank certain features, you probably don't need as many people to sample as the above Scenario 1. Example: survey to understand how users are ranking features (n=30)

- Why? Not all users use every workflow, so your population is not the entire GitLab user base. Rather, it's all users of the workflow you are researching. A smaller population means you don't need as large of a sample.

- You are not collecting a measure that is being compared over time. In other words, the measure is not as granular compared to Scenario 1.

- You don't need a high level of accuracy to make a decision. Knowing from 30 users how they rank features is enough for you to make a confident decision.

- Scenario 4: If your research is trying to understand what kinds of tools users' incorporate in their workflow, you will likely have more open-ended questions in the survey. Example: survey to understand tools in a workflow (n=30)

- Why? You don't need as large of a sample for accuracy, since you are not computing a measure. Rather, you're trying to get an idea of what responses users give.

- You might be looking to target certain tools that users list and want to conduct follow up interviews to understand how they are using the tool.

- A sample size of 30 is commonly used because that number is large enough to accrue a normal distribution in scores even from populations where there is a skewed distribution. In short, this number helps us prevent bias in our results.

Guidelines for survey sample size

To help make it easier to determine sample sizes, here are some guidelines:

- Ranking survey

- Sample size target: 30

- Quantitative survey

- Sample size target: 100

- Qualitative survey

- Sample size target: 30

These are general guidelines, but the size will vary depending on your study and the level of accuracy needed. You can use a sample size calculator, but you will need to know your population size.

Not every result needs to be statistically significant. You can sometimes gather enough responses to have meaningful inferences while not achieving statistical significance. What's more important is to have a representative sample of the population of interest.

Reporting your survey sample size

Justification for your sample size is often needed to help other team members understand why you chose that sample size and what it means in the context of your results. When creating a report for your survey, we recommend adding a slide (search Survey Sample Size Overview Template in Google Drive to access) within the Appendix section of your deck to communicate the

following information to stakeholders:

1. Why did you survey that number of respondents?
2. Were your findings statistically significant?
3. What was your population size?
4. What does the confidence level and margin of error mean?

Here is an example of a filled in report slide for a survey:

SECTION: product/ux/ux-research/surveys/common-types-of-survey-questions/index.md

There are many types of survey questions to choose from when designing your survey. Determining which kind of survey question to choose to use depends on the data you need to collect. Below are some of the more common survey question types that we use at GitLab.

- Open ended - Typically a comment box or a single line where respondents can type in responses.
 - Best for: qualitative data such as providing feedback
 - Things to watch for: don't use too many of these types of questions, time-consuming to analyze, participants tend to skip them and/or leave limited detail
- Multiple choice - Offers respondents one or more options to choose from. Intuitive for respondents to complete. Can also be used with images as choices.
 - Best for: you have informed data to support the options, the options are distinct
 - Things to watch for: offering too many choices
- Rating scale - Allows the respondent to select a number as their response.
 - Best for: research questions that require numerical measures as data (example: scores, ratings, etc.)
 - Things to watch for: the scale needs to be defined (example: if the scale is '1-5' define what 1 and 5 mean)
- Likert scale - A type of a rating scale question that lets respondents choose words/phrases that align with their feelings.
 - Best for: measuring feelings, attitudes, or opinions
 - Things to watch for: the scale needs to be balanced (has an equal number of positive and negative items)
- Matrix - A set of questions that all have the same scale and are presented as a grid. Often used in conjunction with Rating scale or Likert scale questions.

- Best for: consolidating questions that have the same scale into an efficient experience
- Things to watch for: presenting too many questions in the matrix, keeping the scale the same, making sure the questions are clear and related

- Ranking - Allows respondents to rank a list of items, typically by preference.
 - Best for: understanding how users rank certain items in relation to other items (example: ranking possible features to build)
 - Things to watch for: offering too many items to rank (very common) which can impact how much time it takes respondents to rank items, making sure respondents understand each item clearly (can be accomplished by piloting the survey first).

Survey question examples

Question type	Goal of this sample question	Poor wording
version		Problem with the wording
Better wording version		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
Open-ended	Acquire feedback on the GitLab experience	"How great is the GitLab experience?"
		Positively biased
		"Tell us about your experience using GitLab."
Multiple choice	Understanding how many hours respondents use GitLab per day	"How many hours do you use GitLab per day?" (1-2 \ 2-3 \ 3-4 \ 4-5 \ 5-6)
		Overlapping options, lack of scoped timeframe, no option for 'other'
		"Within the last week, how many hours have you used GitLab per day?" (1-2 \ 3-4 \ 5-6 \ Other:)
Rating scale	Net Promoter Score	"How likely are you to recommend GitLab?" (0-10)
		Scale is not defined
		"How likely are you to recommend GitLab?" (0 (not likely) -10 (very likely))
Likert scale	Understand satisfaction of the GitLab experience	"How was your GitLab experience?" (Awesome \ Amazing \ Great \ Good \ OK)
		Unbalanced scale
		"How satisfied were you with the GitLab experience?" (Very dissatisfied \ Dissatisfied \ Neutral \ Satisfied \ Very satisfied)

 ## SECTION: product/ux/ux-research/surveys/qualtrics/index.md

Creating a survey

1. Select Blank survey project
1. Give your project a name.

1. Enter survey questions. On the right hand side of the screen you will find:
 - Question types and formatting (number of answer choices, positioning and so on).
 - Validation (force respondents to answer a question or request that they consider answering the question before leaving the page).
 - Actions (such as page breaks, skip logic and display logic).
1. Update the Look & Feel of your survey.
 - Please refer to Look & Feel settings for guidance.
1. Adjust your Survey Options as appropriate.
1. Under the Tools menu, you will find a host of useful features such as:
 - Collaborate - Allows you to share your survey with another Qualtrics user so you can edit it and/or analyze the results together.
 - Spell check - Supports English (US) only.
 - Generate test responses - Generate automated dummy responses for your survey to see what the dataset and report will look like before sending the survey to actual participants.
 - Check survey accessibility - Determine whether your survey is accessible for respondents who use screen readers, and receive suggestions for improving your survey's accessibility.
 - Analyze Survey - Open ExpertReview's suggestions for your survey.
 - Import/export survey - From/to a QSF, TXT or DOC file.
1. When you have finished building your survey, Preview it to ensure that everything is working the way you expect it to.
1. Once you are ready to launch your survey, click Publish.

Styling (Look & Feel settings)

The following settings should be applied to your survey:

1. Theme: Blank
1. Layout: Modern
1. General
 - Progress bar: Without Text
 - Progress bar position: Top
1. Style
 - Primary colour: 554488
 - Secondary colour: 554488
 - Font: Arial
 - Foreground contrast: Medium
 - Question spacing: Compact
 - Question text: 14px bold
 - Answer text: 14px
1. Motion
 - Page transition: None
 - Additional check boxes: Unticked
1. Logo
 - GitLab First Look Logo (stored in the UX Research & Product graphics library within Qualtrics)
 - Style: Banner
 - Placement: Left
 - Max height: 70px
 - Mobile scaling: 66%
1. Background

- Background type: Color
- Background color: ffffff
- Foreground contrast: Medium
- Questions container: Off

Avoiding spam responses from social media

Distributing a survey on social media is a great way to quickly expand your audience, reducing the time and cost of research. However, be aware that surveys shared on social media are often targeted by bots or fake participants to submit dozens or hundreds of responses in a short amount of time, affecting the integrity of the data. This especially affects surveys with rewards.

Some bots are more sophisticated than others. For example, basic bots will "speed" through the survey or provide illogical responses to open-ended questions; more sophisticated bots will intentionally take more time to complete the survey or even use language from the survey itself to compose logical responses to open-ended questions.

There are multiple things you can do to prevent or filter out most of these responses. We say "most of these responses" because the only true way to avoid spam in your surveys is to not share them on social media but rather to send email survey invitations with individual links that can only be used once, leveraging resources such as Marketo or the Data Warehouse to contact targeted sets of GitLab users we have on file.

Prevent spam responses before sharing the survey

Before sharing your survey, use the following approaches to prevent bots and fake participants. Note that while Qualtrics captcha question and fraud detection features help, they are insufficient in preventing more sophisticated bots.

- Add a captcha question at the beginning of the survey.
- Enable all fraud detection features: Prevent multiple submissions, Bot detection, Security scan monitor, and RelevantID.
- Add a "honeypot" question, a question hidden from humans but "visible" to bots. Some bots will answer this question, allowing you to screen them out or filter out their responses.
 1. Add an optional single-choice question (checkbox) somewhere in your survey. The question and location is not relevant as it will be hidden from humans. For example, "Are you human?" [] Yes.
 1. Add the following JavaScript to that question, by replacing the existing Qualtrics.SurveyEngine.addOnload function:

```
js
Qualtrics.SurveyEngine.addOnload(function()
{
  jQuery(''+this.questionId).hide();
});
```

1. Preview the survey to check that the question is hidden when the page loads.
- Add a required choice question that checks respondent's attention and logic. For example,

directly state which answer they should choose or ask which year it is.

- Add illogical options to multiple-choice questions. For example, in the question "Which of these tools do you use for code review?" add an illogical option like [] Tanuki Code.

- Ask the same required question at two separate points of the survey. Try to use a simple multiple-choice question that people can easily answer. For example, their job title or their team size.

- For all three question types above, if you want to screen out respondents based on their answers, see the Branching Respondents Out of the Survey section under the Screening Out Respondents in the Survey Flow section. After adding the survey flow logic described in that section, you can then decide to just flag those responses for later analysis or immediately discard them.

In addition, if your survey is shared on multiple social media channels, you can track f