

ts input, the ticket can now be marked as Solved. Use a resolution of Other, and add the text CSE Referral in the description.

1. If the discussion in the CSE channel results in the CSE advising they can not help with the ticket, continue to work it as a support ticket.

Helping with a ticket that has an Assigned Support Engineer

Some customers have an Assigned Support Engineer (ASE) who will own their tickets. This will be apparent from the org note in the ticket. For such customers, see the ASE workflow.

FAQ

How many new tickets should I pick up each week?

Each week, every Support Engineer should aim at least to meet, and preferably to exceed, the appropriate baseline from the following list for the number of first responses:

Ticket type	FRT goal	
-----	-----	
Self-managed	5 tickets	
SaaS	5 tickets	
SaaS Accounts	20 tickets	
L&R	15 tickets	

The goal is to ensure fair distribution of tickets across the team. Intermediate and senior engineers excluding those in core L&R and Associate SE roles are expected to develop as well-rounded contributors by consistently working on a balanced mix of SM, SaaS, and Dedicated tickets. These expectations are informed by overall ticket volume, team size, and an average PTO rate of 15%. We continuously monitor ticket volume and staffing levels to ensure these guidelines remain accurate and appropriate. For further details please review the Support Engineer Performance Indicators FY26 issue.

What if I can't meet baseline?

There will be times when you may be leading an escalation, working through challenging tickets, or focusing on non-ticket work, and you cannot take on new ticket assignment. This is okay! The key is to make sure you have coordinated with your regional team and Manager so they are aware of the risks to our team.

What can I do when I'm stuck and I really need help?

Follow the how to get help workflow for guidance.

I'm going to be absent, what do I do with my assigned tickets?

Please see the Support Team Member Time Off page for guidance on taking time off.

What if the customer is absent and wants to keep the ticket open until they return?

In this scenario:

1. Acknowledge the customer's request and inform them of your plan.
1. Set the ticket status to on-hold.
1. Change the ticket type to Task.
1. Set a reminder for yourself to check in with the customer 1-2 days after their planned return date.

Because of the behavior of On-Hold tickets the ticket will remain in the on-hold state until the customer returns. This is useful if the customer will be unavailable for 1-2 weeks. If a customer will be unavailable for a longer period of time, consider asking them to handover the ticket to a colleague, or open a new ticket upon their return.

May I reassign a ticket to someone else?

There are some situations in which it is appropriate to reassign a ticket. A few common situations are listed below, but others may arise. If you encounter a scenario where you need to reassign a ticket that isn't detailed in the handbook, please discuss with your manager, and then add it below!

If you'll be out of office

If you're out of office, either planned or unplanned, follow the guidance in Support Team Member Time Off.

If you need an expert

If you've determined that specific expertise outside your own is required to resolve the ticket, pair with an expert so that you can get the ticket resolved and learn in the process. If that person determines that they need to take the lead due to the advanced or complex nature of the problem, then:

1. Send a message to the customer informing them:
 - You've asked another support engineer with relevant expertise to take the ticket
 - You've reviewed the ticket with that engineer
 - You'll stay involved in order to help in any way you can
1. Assign the ticket to the expert

If you have too many tickets

If you've become overloaded with tickets, you may look to find one or more other SEs to take some of your tickets. Be sure you discuss each ticket before reassigning it to gain agreement and so that the other support engineers don't have to start from scratch.

Once you've found a new assignee:

1. Add an internal note summarizing the ticket and noting the reassignment
1. Send a public reply to the customer informing them of the reassignment
1. Reassign to the new engineer

What if a customer has confirmed a ticket can be closed but the ticket has no assignee?

While we should ensure that all tickets always have an assignee from FRT to close, there are some scenarios where a ticket is ready to be closed but has no assignee. For example, a customer may open a ticket and inform us that they have resolved the issue before we send an FRT. In these scenarios:

1. Inform the customer that you are changing the ticket status to solved
1. Assign the ticket
 - If a solution was provided, assign to the engineer who provided the solution
 - If the issue was not solved as the result of a direct suggestion, assign it to an engineer who has significantly contributed to the ticket throughout its life cycle
 - If the customer solved the ticket before we could respond, take assignment of the ticket yourself. After all, you're doing the work of acknowledging the customer's update, and solving the ticket!
1. Mark the ticket as solved

How many issues or incidents should I handle in a single support ticket?

As detailed in Support General Policies
it is GitLab's policy to handle each unique issue or incident within a single support ticket.

How can I open a new ticket on behalf of a customer?

As detailed in the Support General Policies
it is GitLab policy to handle each individual incident, problem or issue within a single support ticket. If a situation arises where you need to open a new ticket on behalf of a customer, you must use the Create new ticket plugin within the Zendesk Super App to ensure proper routing and SLA assignment:

1. Within a ticket, navigate to the Zendesk Super App and select the Create new ticket plugin.
1. Select the appropriate form type. This will most likely match the request type of the originating ticket.
1. Enter the customer's details like subject, problem description and other relevant information. Use the customer's existing Zendesk Ticket to obtain the necessary customer information to complete the form.
1. Submit the ticket using the Create Ticket button.

Once the new ticket has been submitted, the app will provide you with the corresponding ticket ID. You can click the ticket ID to directly navigate to the newly created ticket.

If you plan to immediately start working the ticket, please assign the ticket to yourself.

Can I use output from an LLM in ticket replies?

While you can use a tool like GitLab Duo to help when researching and when writing a reply to a customer, there are a few important things to keep in mind:

- Verify the information provided by the LLM to ensure that it is accurate using the docs and the source code or by checking with a human subject matter expert
 - You should pay special attention to making sure whether things like environment variables, configuration options, UI settings and documentation URLs actually exist
 - You should double check that any bugs or feature availability identified aligns with the GitLab version/edition/tier/install method that the customer is using
 - You should make sure that you can understand, can explain and have tested any commands or code snippets generated by an LLM and shared with a customer
 - You should not send text from an LLM to a customer verbatim without at least verifying it for technical accuracy first
- Be mindful of GitLab's Data Classification Standard and the kind of input you have in mind when selecting a tool
 - When working with customer data, always consider whether you need to anonymize information before sharing with any LLM
 - Take special care with customer screenshots, logs, and configuration files which may contain sensitive or identifying information
- LLMs can be useful for refining and improving a response you have drafted to a customer (such as checking for clarity, completeness, and tone), however responsibility for the final response always lies with the Support Engineer, not the tool.

Read the General Purpose AI Tool Usage Guidelines for additional guidance.

SECTION: support/workflows/working-on-zendesk-content.md

General process

To make it easier for Support Engineers to contribute to things that are managed by Support Operations, data and code have been separated for some items. This approach removes the need for Customer Support Operations to review every change proposed for Zendesk views, macros, triggers and automations.

For any changes to the content of Zendesk views, macros, triggers and automations consider their impact first. If your changes go beyond something like fixing a typo or updating a link and have broader repercussions for the team, consider the Change Management in GitLab Support handbook page. Otherwise you can suggest your changes directly via a merge request in the relevant managed content project.

For creating new content or renaming existing content (e.g. a new macro, or renaming a view), an STM issue is required as Customer Support Operations has to handle these actions. The managed content projects only allow managing content of existing items.

Managed content projects

Refer to the corresponding projects and create a merge request with your changes:

- [gitlab-com/support/zendesk-global/automations](#)
- [gitlab-com/support/zendesk-global/macros](#)
- [gitlab-com/support/zendesk-global/triggers](#)
- [gitlab-com/support/zendesk-global/views](#)
- [gitlab-com/support/support-pages](#) (also known as articles)

Whether you're proposing minor changes directly or something bigger following an RFC discussion, tag a Support Manager for review and ask them to merge the changes. Once changes have been merged in one of these projects, you'll notice that the Sync Stage badge says Awaiting Deployment on the project overview page. This indicates that there's changes in the project that will be rolled out during the next regular deployment.

Macros

- If you want to edit the wording of a macro, make an MR to the repo (Zendesk Global or Zendesk US Government).
- If you want to make non-wording changes (renaming it, add/removing options, etc.), use the Support Super Form 'Edit a Zendesk macro' option.
 - NOTE If a managed content file is needed (as you selected a Public or Private comment would be used), the Support Super Form will create a placeholder file for you (if said file does not exist). For the quickest resolution, it is best to not create the managed content file before your submission of the Support Super Form.
- If you want to deactivate a macro, use the Support Super Form 'Edit a Zendesk macro' option.
 - NOTE If the macro has a managed content file, only deactivate the file after Support Ops confirms the macro has been deactivated in Zendesk. This step is important to prevent sync errors in the ZD Macro repo.
- If you want to create a new macro, use the Support Super Form 'Create a Zendesk macro' option.
 - NOTE If a managed content file is needed (as you selected a Public or Private comment would be used), the Support Super Form will create a placeholder file for you (if said file does not exist). For the quickest resolution, it is best to not create the managed content file before your submission of the Support Super Form.
- To discuss a macro (but not actually action on anything), use a support-team-meta issue.

Organizations

Modifications to Zendesk organizations can be made through the use of the Support Super Form

- The Support Super Form will walk you through making your modification. If an option does not appear in the list of available choices, it cannot currently be modified.

To make changes to a Global organization's note you would use the Zendesk Global Organizations project to modify an organizations notes.

For US Government Organizations, all organization notes are managed manually by the Customer Support Operations team. Due to the sensitive nature of the organizations, please reach out to

Support operations via slack to make changes.

Articles

- If you want to edit the wording of an article, make an MR to the repo.
- If you want to create a new article, use the Support Super Form 'Create a Zendesk article' option.
- To discuss an article (but not actually action on anything), use a support-team-meta issue.

Previewing variable replacement

When using liquid variables it can be useful to verify that they will be working as you expect them to. A simple way to do this is to create a test ticket, pasting the content of your automation, trigger, or macro in there and observing the outcome.

SECTION: support/workflows/working-with-issues.md

Creating, updating and escalating GitLab issues correctly is an important part of providing quick and accurate customer support. The support team uses the processes and escalation points described on this page when dealing with GitLab issues.

Issue Prioritization

In general, the Product team will prioritize all issues (not just customer requests) based on types of issues and the direction of the product.

The Support Team plays a role in communicating the impact to customers of issues and feature requests. By using appropriate templates, adding labels, and adding relevant information in descriptions and comments, the team can communicate which issues affect customers along with their priority and severity. By participating in the scheduling effort for each release, the Support Team represents an additional voice of the customer in product development.

Additional Context and Notes

Having the Product team comment in the issues directly follows our core value of being transparent and will help customers understand the context around why / when their issues are being resolved, and it provides direct feedback from customers to the Development and Product teams.

Working this way, it is possible that a customer reported issue is not picked up for a while (scheduling first, then time to work on a fix, then schedule for release, etc.). However, the idea is that this is OK because most truly urgent issues will in fact be regressions that don't have this scheduling problem. If a bug isn't a regression, that means it has existed for more than a month when the customer notes it, and thus we've gone at least a full month without someone reporting the issue as urgent.

Issues are not scheduled for a particular release unless Product adds them to a release milestone and they are assigned to a developer. We aim to be realistic about scheduled

deliverables and will avoid scheduling issues that cannot be delivered in a given release.

Searching for Issues

Before creating an issue or commenting on an existing issue you'll need to search for any existing or related issues.

Here are some tips:

- Search using an external search engine, like Google.
- Issues are created in a number of projects in addition to the main GitLab Rails project, of which a handful of the most common are:
 - The Omnibus project is where issues relating to packaged GitLab and Docker are created.
 - If an error arises during gitlab-ctl reconfigure this is probably the best place for an issue to go, unless it relates to a database migration which usually should be in the Rails project issue tracker.
 - Issues relating to deploying GitLab with Helm go in the GitLab Chart project.
 - Look in the Gitaly project for issues relating to Gitaly and Gitaly cluster.
- It's often worth searching across the whole gitlab.org group, particularly if you're not sure which project issues might have been raised in.
- Advanced search tips:
 - Combined with searching across the whole group, advanced search has the advantage of also searching Epics.
 - Once you have results, you can filter out closed issues. This is often useful if you're just checking for outstanding bugs or feature requests.
 - If you're searching for an error or for other text generated by the product, this will also search all our code repositories as well.
 - If you're looking for a particular error or string, put it in "double quotes".
 - See other available syntax options from the Advanced Search Docs
- Searching in the issue view is usually more efficient if you can filter the results:
 - If you are looking for a change that was implemented in a particular GitLab release, filter on Milestone.
 - We create a milestone for each monthly GitLab release.
 - Consider also searching merge requests by milestone as well. If you think something changed in a particular release, this will identify actual code changes in that release. The issue might be closed against a later milestone when the whole feature was finally delivered.
 - If you know which product area is responsible for a feature, filter on the ~group:: label.
 - The features by group page is good for finding out the correct group.
 - Alternatively, look at the labels used on related issues and MRs you may have found already.
- Sort order: If you get a lot of results in the issue view, look at the options for sorting. For example: Created date might be more relevant if you suspect a recent regression in the product.
- You are likely to spend some time searching for issues you've read before.
 - Consider adding a reaction on all issues that you come across, and any you raise. For example use bookmark, reminder ribbon, or pick something really unusual! You can then filter on My-Reaction and search within those results.
 - Try searching your browser's history.
 - For quickly finding things based on your own comments, consider setting up the User Contribution Search tool. It indexes all your contributions and gives you a fulltext search

interface.

- Read more about searching in GitLab in the product documentation.

Adding Comments on Existing Issues

Regardless of the type of issue, please include any relevant information along with a link. Also check that the correct labels have been applied.

Please see the product handbook to see what information product wants us to provide for feature requests

Optional:

Search for Zendesk tickets based on the GitLab Issues custom field. Copy and add the link to the search query as a comment to the issue, specifying that it's an internal link to a ZenDesk search of related tickets. This will give the PM more insight into how many customers have encountered an open bug, or have requested a certain feature.

Note: To obtain the field ID for the GitLab Issues field, you can reference this repository file.

Adding Labels

Using the appropriate labels is critical to ensuring visibility of issues and to get them on relevant PM's radar.

Required:

- Group: Unsure of which? Check out documentation metadata, label descriptions, DevOps Stages, features list by stage or similar existing issue.
- Type Labels should be added by the template, but add them if any are missing.
- ~customer (if a ZenDesk link is added, the bot will add this automatically if you forget)
- ~regression .x if applicable; for high-impact ones, add ~"Next Patch Release" and ping the relevant lead and subject area experts

For ~customer+ any label that allows severity (most commonly ~"type::bug") labeled issues, a Severity estimate is required. If it is missing, please add them to attract PM attention to the issue:

- Severity: Follow the definitions to the best of your ability when assigning severity. If it's an ~severity::1, mention the PM and consider posting in the appropriate Slack channel as well. As Support often has a better idea of the impact on the customer(s), please explain the impact in a comment when you assign the Severity label. Feel free to have the customer add a comment as well, adding any other context they feel might be important.

Optional, but highly recommended:

- Stage
- ~Reproduced on GitLab.com if applicable

For L&R related issues, please see the specific guidance on Fulfillment issues.

Escalate New and Existing Issues

The Support team can directly ping the PM on the issue or in the product Slack channel (see DevOps Stages) in case this may help with communication. Comments might include asking for an update on behalf of the customer, or discussing the severity or priority especially if it needs to be increased.

Creating Issues

Whenever possible, reproduce the issue and file a public issue with a reference to the ZD link as an additional example.

When reporting a problem, use the Bug template, then fill out as much of the information as possible. Ensure you add labels.

When writing issues, consider adding questions as a comment after creating the issue. For example: "@PM please provide feedback on this issue. Are we interested in fixing/implementing this? How critical do you think it is?"

Avoid Re-opening Closed Issues

If you find an existing closed issue that is the same or similar to the customer problem at hand, you probably should not re-open the closed issue. The closed issue may actually be unrelated. Or maybe the closed issue describes the same problem you see now, but has a different root cause. Even if it's the same identical problem, you probably should not re-open the issue because that issue might have a previous milestone and so now the re-opened issue would have an invalid milestone in the past.

Instead, consider creating a new issue, and mention/link the existing closed issue to it. Let the Product and Engineering teams triage the new issue further and don't worry about creating duplicates.

Maintain confidentiality

If an image, log output, etc. is required for the issue, try to produce your own test image. If you are unable to reproduce the issue and you wish to use the image/information provided by the customer make sure you obtain permission from the customer since the image/information may (inadvertently) include sensitive information like names, group names, user names, or code.

Public issues are always preferred, but if customer logs or other information needs to be included and the customer is willing to share it internally, but not publicly then make the issue confidential.

Information gathering

For the Application and environment information section of issue templates, use:

- Omnibus: `sudo gitlab-rake gitlab:check`

- Source: `sudo -u git -H bundle exec rake gitlab:check RAILS_ENV=production SANITIZE=true`

and

- Omnibus: `sudo gitlab-rake gitlab:env:info`

- Source: `sudo -u git -H bundle exec rake gitlab:env:info RAILS_ENV=production`

Creating Feature Proposals

As per our Statement of Support, the Support Team will generally ask the customer to create feature requests. Feature requests with direct feedback from customers are valuable as customers are often best equipped to explain their use case, requirements, and needs. Ask customers to create the feature request issue and share the link with us. Once an issue link is available, add labels and relevant details in the comments, and link the source.

If you create a feature proposal on behalf of a customer, please follow the same process as Creating Issues by using the Feature Proposal template. After the issue is created, share the link in a reply encouraging the customer to follow and contribute to the issue.

Note: GitLab has limited development resources. Additionally, we must think about how widely applicable a feature may be to other users. Requests that are very specific to one company's workflow are likely to be rejected. Even if a feature seems widely applicable, we may leave the feature proposal dormant for some time and see if other users and customers chime in that they are also interested. Features that garner interest from multiple organizations will be considered more rapidly. Of course, there are always exceptions to these 'rules'. This note is meant to set the expectation that feature proposals may not be implemented quickly.

Functional escalation points

Service/Product	Escalation Types	Escalation Point
Assignment		
----- ----- -----		
----- -----		
GitLab	Bug reports or Feature proposals	Bug or Feature proposal
Omnibus GitLab	Bug reports, Feature proposals	< https://gitlab.com/gitlab-org/omnibus-gitlab/issues/new > Omnibus GitLab specialist
GitLab Runner	Bug reports, Feature proposals	< https://gitlab.com/gitlab-org/gitlab-runner/issues/new > GitLab CI specialist
GitLab Workhorse	Bug reports, Feature proposals	< https://gitlab.com/gitlab-org/gitlab/-/issues/new?issue%5bdescription%5d=%2flabel%20%7eworkhorse > Maintainer of gitlab-workhorse

See the GitLab team page for assignments

Operational escalation points

Service/Product	Escalation Type
Escalation Point	Assignment
----- -----	
----- -----	

----- |

| GitLab Infrastructure | Anything related to the running of GitLab.com, performance, something breaks | <<https://gitlab.com/gitlab-com/infrastructure/issues/new>> | Production

Lead/Senior Production Engineer |

| Support Engineers with GitLab.com console access| Anything related to the use of GitLab.com, operations that can't be performed with admin access | See Console escalation requests | Use ~"GitLab.com Console Escalation" label |

| GitLab Support | Any and all questions in relation to providing customer service for GitLab users and customers. | <<https://gitlab.com/gitlab-com/support/support-team-meta/issues/new>> | Support Team Lead/Senior Support Engineer |

See the GitLab team page for assignments

Omnibus GitLab

- Related to Omnibus GitLab packaging only.
- GitLab omnibus release packages

GitLab Runner

- Information on GitLab Runner
- Runner documentation

GitLab Workhorse

- Information on GitLab Workhorse
- Description "GitLab-workhorse is a smart reverse proxy for GitLab. It handles "large" HTTP requests such as file downloads, file uploads, Git push/pull and Git archive downloads."

GitLab Infrastructure

- Reach the infra team on Slack
- Old blog post on infrastructure

General Product Feedback

In the case where general product feedback is received and it is not clear whether it is related to or belongs in an issue, feel free to convey the feedback to the product team as outlined in our Product Feedback section.

SECTION: support/workflows/working-with-pods.md

What is Support Pods?

Support Pods is a flexible framework which any group of team members can use to set up self-organizing interest-based groups to work on any set of problems they wish to take on.

The full Support Pods handbook and a list of active Support Pods can be found in the Support Pods handbook page.

How do I get involved?

You can follow the instructions here to start or join a Support Pod.

Integrating Support Pods in day to day work

Support Pods are currently not integrated into our ticket handling workflows. This is intended to allow each Support Pod to determine their own approach to working with tickets and customer issues.

In general, each Support Pod's approach to work should keep to the Key Principles and Priorities & Impact guidelines for working on tickets, and contribute to the team's achievement of positive customer outcomes.

Each Support Pod may have specific recommendations on integrating its work into your day-to-day. These recommendations should be documented in the Support Pod's README in the Support Pods project.

SECTION: support/workflows/working_with_security.md

Overview

Previously, GitLab utilized the email address <security@gitlab.com> to report and inquire about security concerns. However, as GitLab and our Security Department has grown and expanded, we were unable to provide the high level of service our customers deserve utilizing that singular queue. Users who reach out to security@gitlab.com will now receive an auto-reply providing them with specific instructions for reporting the various types of security concerns. If the auto-responder does not answer their questions or a security-related ticket is submitted to Support you can utilize the macro Security::All Security Questions to provide the user with detailed instructions.

Workflows

Please do not transfer to security and refer to the relevant workflow for the following:

- Reinstate Blocked Accounts
- Abuse Report Review Request
- Log Requests
- Reporting a New Security Incident

You can also utilize the Security::All Security Questions macro for more details on the language. If the workflows above and the Macro do not resolve the customer's concern, please

post a link to the ticket in the sec-fieldsecurity Slack Channel.

Triage Workflow

Vulnerability disclosures are no longer triaged via ZenDesk. When in doubt, please involve the Application Security team.

- If the report would be Informative or would be an ~"type::feature", you may ask the requester to file an issue themselves. You can start with the "General -> Ask to create an issue" macro.
- If the report appears to be a valid security issue, or if in doubt:
 - If the severity may be an ~severity::1 or otherwise needs immediate investigation:
 - Engage the Security Engineer on-call
 - Manually create an issue for report following the creating a security issue instructions
 - Otherwise, reply to the reporter to see if they wish to either:
 - Report via HackerOne
 - Create a confidential issue themselves following the creating a security issue instructions
 - Have a team member create the issue and reply with their preference for acknowledgement.
 - To ensure timely handling of the report, let the reporter know that if no response within 3 business days, a GitLab team member will create the issue.

Request for Information on Security Update or Patch

Sometimes, a customer will request details on a security update that was released. e.g. "Should I worry about this? What's this patch about?".

A summary of GitLab CVEs for specific versions is available in Customer Success' "What's New Since" tool.

If the customer is asking about a security vulnerability published as part of a release, the only information we can provide is what is in the security blog post. For more information on security communication, please see the security incident communication page.

Security will make the issue public if possible after a set number of days.

If you believe more information should be made available in the blog post, or to a specific customer, please open a confidential issue in the security communication tracker.

Note: Confirmed mitigation strategies are typically added to the security blog post. If none are listed, the only recommendation is to upgrade to a version where the issue is fixed.

While some suggestions (such as disabling a feature) may seem like they would mitigate an issue, without validation from the security team, we cannot be fully certain.

Following the Responsible Disclosure Policy, see below on reporting a security issue.

General Guidelines for Security Staff

1. All members of the Security Team should have followed the Light Agent request process as part of their onboarding and should use these accounts to make internal comments. When it's necessary to reply or change the status of the ticket use the shared account in 1Password.
1. Always cross-link any relevant ZenDesk, HackerOne, and GitLab issues, in internal comments where appropriate and available. ZenDesk links in GitLab issues should be noted as "GitLab internal only": Reported via ZenDesk (GitLab internal only): <https://gitlab.zendesk.com/.../xxxxx>
1. When responding directly via email, you must include the requester's email address as a recipient, since emails replies to ZenDesk from Light Agent accounts will only be processed as internal notes and not sent to the requester.
1. If you email directly, you are responsible for setting the ticket to "Solved" or "Pending", or finding someone with permission to do so.

If the customer has already created an issue

In the case that the customer has already filed an issue for the vulnerability:

1. Mark the issue is confidential
1. Add security, customer, and bug or feature proposal labels
1. Assign Severity and Priority Labels

If the customer has not yet created an issue

See Creating New Security Issues

SECTION: support/workflows/working_with_support_ops.md

Overview

Customer Support Operations sits under Security, alongside teams like Corporate Identity and IT End User Support. As a corporate provider, Customer Support Operations is specifically tasked with providing support for Customer Support.

This page exists primarily as a quick-reference, linking to sources of truth and making sure Customer Support has the information needed to get things done!

Quick Links

Resource	Link
-----	-----
Issue Tracker	Customer Support Operations Issue Tracker
Bug Report Template	Create Bug Report
Feature Request Template	Create Feature Request

Incident Handling	See Incidents
Slack Channel	support-operations
System Criticality Sheet	Customer Support Operations System Criticality

> Important: All issues for Customer Support Operations should be created in the Customer Support Operations Issue Tracker.

Self-Service Options

Before creating an issue, check if your request can be self-serviced. The following items can be handled directly by Support team members:

- Zendesk Managed Content - Updates to macros, triggers, and other Zendesk content can be managed directly by Support team members with manager approval. See: Working on Zendesk Content
- Support Team YAML Files - Changes to support team configuration files can be made through direct merge requests. Customer Support Operations must approve to merge. See: Support Team YAML Files

> Note: If an issue is opened for something that could have been self-serviced, the issue will be closed with instructions on how to self-serve the change. If you need help, reach out in support-operations and we can guide you through!

Workflows

Ticket where customer is asking for support portal changes

Occasionally, a customer might mis-file a ticket and send a request for support portal changes using a Support form. In these cases, please change the form to Support Ops in Zendesk and Customer Support Operations will handle it from there. Please do not send a public reply (as that removes our SLA clock from the ticket).

Asking general questions

For general questions and assistance, reach out via the supportoperations Slack channel. From there, a member of the Customer Support Operations team will respond and work with you to help address the question.

Basic Issue Flows

All requests for Customer Support Operations should be created in the Customer Support Operations Issue Tracker using the appropriate template. With the exception of a few self-service options, nearly everything starts with an issue!

If what you want to get done is:

- On the roadmap? Create a Feature Request issue!
- Not on the roadmap, but seems easy? Create a Feature Request issue!
- Maybe easy, but maybe hard? I definitely know what I want though. Create a Feature Request issue!
- Reporting a bug? Create a Bug Report issue!

- ... well, actually, I don't totally know yet. Create an RFC issue in support-team-meta (or work otherwise within the Change Management in GitLab Support framework)

If you're creating an issue, the Customer Support Operations team expects that:

- You're empowered to make decisions. That is, you're a DRI at the appropriate level to steward your change or you've been delegated that responsibility by an appropriate DRI.
- You're responsive. You're prioritizing working on the thing you want so that it can be shipped on time, and to spec.
- You'll be ready to validate and sign off. When Support Operations deploys something to staging, you need to be ready to test (or delegate / organize testing) and sign off on acceptability before it can go live in production.

More about Working with Customer Support Operations.

More about stages that Customer Support Operations uses to represent progress.

Bugs

Bugs should be reported using the Bug Report Template in the Customer Support Operations Issue Tracker (project path: `gitlab-com/gl-security/corp/cust-support-ops/issue-tracker`).

When creating a bug report:

1. Select the Bug template from the dropdown
2. Fill in all required fields, especially steps to reproduce
3. Select appropriate severity

Bugs are generally of a high priority to fix, and will get scheduled in the current iteration. Depending on workload, severity, fix complexity and system criticality may get scheduled in the future.

More about the Bug Report issue flow.

Feature Requests

The focus for feature requests is on aligning them with our existing Support Roadmap. For larger requests, you may need to work with the leadership team to escalate issues to re-plan any conflicts and protect strategic items. For smaller requests, they can often be added in as we go, as long as they don't get in the way of delivering current roadmap items.

Feature requests should be submitted using the Feature Requests Template in the Customer Support Operations Issue Tracker (project path: `gitlab-com/gl-security/corp/cust-support-ops/issue-tracker`).

When creating a feature request:

1. Select the Feature template from the dropdown
2. Clearly describe the business need and expected outcome
3. Include any relevant stakeholders using the `/cc @username` syntax

More about the Feature Requests issue flow.

Incidents

Things do break! Customer Support Operations is here to help.

First, check the Customer Support Operations System Criticality Sheet to determine the severity of the issue.

Incident Reporting Process

1. Start by posting in the support-operations Slack channel (during Global Support Hours)
2. Wait for a response. Continue down this list if you get no response for:
 - 24 hours for issues with Administrative systems
 - 8 hours for issues with Business Operational systems
 - 30 minutes for issues with Business Critical Systems
 - 10 minutes for issues with Mission Critical Systems
3. Create a new PagerDuty incident with "Customer Support Operations" as the Impacted Service (See: Support Operations On-Call Workflow)
4. Be available to work with the on-call to develop a plan to resolve the problem

When creating an incident issue, include:

- Clear description of what's happening
- Impact on users/customers
- Any error messages or screenshots
- Steps you've already taken to troubleshoot

More about how Customer Support Operations handles incidents.

General tips

1. Start early - if you are the DRI for a key strategic item on the Support Roadmap, start working with Ops early. Create a tracking issue so a plan can be formed and blockers discovered early.
1. Don't "save up" work - if items are important to the business, they should be tracked. If they're not tracked, they won't get prioritized.
1. Have strong opinions on outcomes - avoid strong opinions on technical implementation.
1. Don't be quiet - as a DRI, if a technical implementation isn't looking like it's going to meet the criteria, speak up whether it's in the Validation stage or not. It's better to ship the right thing a bit late than the wrong thing quickly.

Deployments and Delivery

Unless otherwise communicated, changes ship on the 1st of the month. This is to align any changes to metrics with month boundaries.

Once you've signed off on work, make sure you're following the steps in Change Management in Support - Rolling out a Change to make sure everyone on the team is aware.

SECTION: support/workflows/writing_workflows.md

Overview

This page is meant to provide some guidance on writing workflow pages including creating new or updating existing workflows.

It is meant to be easy to read and quick to scan in the spirit of the guidance provided on this page.

Why this is important

While having information written down has many benefits, for people to find what they need, it needs to be well organized too.

Organization is particularly important in the workflows because often steps for a process only apply in some cases.

Making it easier for a team member to follow a process while also allowing them to only read the necessary parts to complete a task is needed for efficiency.

Creating, updating, and improving workflows fits with [GitLab values]/handbook/values/) and emphasizes many of our operating principles and approaches including:

- Handbook first
- See something, say something
- Everyone can contribute

This very page came out of a discussion on increasing support efficiency.

Basic guidelines

Good workflows have:

1. the information related to the title.
 - If you're including new information, which workflow does it fit best in? If it doesn't fit anywhere, consider a new workflow.
1. information broken up in small sections.
 - Is what you're adding a discrete step for a process? Make sure it's part of a list. Specific use case? Use a heading to separate it.
1. well organized and descriptive headings to make looking for information easier.
 - What level heading is appropriate? Are there many others (example h3) already? If so, consider adding a new higher level heading (example h2).
1. links to other relevant workflows or documentation.
 - A lot of our workflows are related, but we should avoid duplicating information. Link to it instead.
1. a reasonable length.
 - If you think a workflow page is too long, then it might be. Consider if it can be broken up, or if it needs to be reorganized.

1. visual aids, such as flowcharts, whenever helpful.
 - Consider adding a mermaid chart or image to help visual learners.

For general writing practices including inclusive language and use of markdown, consider following the GitLab documentation guidelines.

Example from okay to better

Let's take the Support Time Off page as an example (even though it's not a workflow page, it's process related).

The original version organized the information by using number of days off as the "divider".

1. Less than half a day
1. Half day or more
 - 1. Use Workday
 - 1. PagerDuty shifts
 - 1. Meetings
1. 1-2 days
 - 1. Check with customers
 - 1. Groom queue
1. 3+ days
 - 1. Notify team
 - 1. Change ticket workflow
 - 1. Find new assignees for your tickets

That's understandable, but difficult to scan which parts are relevant to an individual, because not all Support team members have PagerDuty shifts, or assigned tickets (think Support Ops).

Additionally, some pieces duplicated company handbook information.

In a revision, the page was reorganized to focus on responsibilities so that team members could skip sections that are not relevant to them.

1. Marking time off
 - with link to relevant company handbook section
1. Coverage for PagerDuty
1. Coverage for tickets
 - 1. Refine queue
 - 1. Check with customers
 - 1. Additional actions for 3+ days
 - 1. Notify team and find new assignees
 - 1. Change ticket workflow

This is a small, but hopefully good example of how a difference in organizing the content can make it easier and more efficient for the reader.

Creating a new workflow

When creating a new workflow, please keep the above guidelines in mind, and consider this

process:

1. Ordering and organizing with an outline.
1. Get someone to review the outline.
1. Writing the full content.
1. Get one or more reviewers focusing on one of the following for their review:
 1. Organization and structure.
 1. Clarity and conciseness.

Template

There is no specific template for a workflow (except the metadata noted below). Some possible places to start:

1. Use another workflow as a template.
1. Use the GitLab documentation structure and template.

Metadata

The following metadata and table of contents pieces should be at the top of every workflow page.

Notes:

1. The only optional (can be removed) metadata line is subcategory. All others must be filled in.
1. See the existing workflows for existing categories.

text

First heading

SECTION: support/workflows/zendesk-instances.md

Emergency Requests for Zendesk US Government

Our US Government Customers will be provided with an email address they can email to in case of emergency. The emergency email triggers an emergency page that our engineers will respond to within the time outlined in our Support Page.

Working Tickets in Zendesk US Government

Please see Working with US Government Support tickets for information about working tickets on this instance.

SECTION: support/workflows/zendesk-ticket-basics.md

Ticket status

Each ticket in Zendesk has a status that indicates its current state.

Ticket statuses and their descriptions

Status	State	Notes
-----	-----	-----
New	The ticket has just been submitted and has had no replies.	
Open	The ticket has had one or more replies, and the user is waiting on GitLab Support to provide the next reply.	
Pending	Support has replied to the ticket and is waiting on the user to respond. If the user does not respond for 20 days, Zendesk will change the status of the ticket to Solved.	
On-Hold	Support is working on the ticket or is awaiting information from another GitLab team See Behavior of On-Hold tickets	
Solved	The ticket has been solved When a user replies to a Solved ticket, Zendesk reopens it. We automatically change a Solved ticket's status to Closed after 7 days.	
Closed	The ticket is archived When a user replies to a Closed ticket, Zendesk opens a new ticket with a note that relates it to the closed one.	

Make sure to reply to the user

If a user's reply is the last one in the ticket, make sure to send a public reply when changing the ticket status.

Changing a ticket's status (except for Solved) without replying will not stop the ticket's breach clock. See SLA clock for more details.

Avoiding status change automations

By default, our Zendesk automations will do the following:

- After 7 days in Pending, it will send a notification to the user we are still awaiting a response from them.
- After 14 days in Pending, it will send another notification and mark the ticket as Solved.
- After 7 days in Solved, it will close the ticket.

While this is normally the right workflow, there might be situations in which you need to prevent this from occurring. To do so, use the appropriate Zendesk labels:

Label	What it does
-----	-----
skipautosolve	This tells Zendesk to refrain from moving the ticket to Solved automatically
skipautoclose	This tells Zendesk to refrain from moving the ticket to Closed automatically

NOTE: If the ticket has been reopened after already auto-solving and we want to prevent autosolve from happening again, the autosolve and autosolvemessage tags will be present. These do NOT need to be removed when adding the skipautosolve tag.

Behavior of On-Hold tickets

We have a number of automations around tickets in an On-Hold status:

- When a ticket is set to On-Hold, it will be automatically assigned to you by the trigger Automatically assign on-hold ticket to the engineer who put it to the on-hold status.
- Tickets without an assignee will be automatically reopened by the trigger Automatically reopen on-hold tickets without assignee.
- Tickets with an assignee will be automatically reopened in 4 days by the automation Reopen on-hold tickets after 4 days, unless the ticket is of type Task.

SLA clock

If a customer's reply is the last one in the ticket, do not set it to any status silently (except for Solved), because the SLA clock will continue to run and the ticket may breach SLA silently. Instead, send a confirmation, greeting, or other message, while also changing the status.

Ticket subject

Ensure that the subject of a support ticket is both descriptive and accurate. You can edit the Subject to fix typos or make the problem clearer. Some examples include:

- gitlab error 500 on login -> gitlab error 500 on login due to no partition of relation "auditevents"
- My Account was Blockes -> My Account was Blocked
- git reconfigure with below errors -> git reconfigure with letsencryptcertificate errors

Other ticket fields

Ticket fields help us capture important data that will help us improve the user experience.

Depending on the view in which you are working and the form that has been selected for the ticket, you might need to fill out some ticket fields manually. As a high percentage of our tickets are solved or closed automatically through our workflows, it is important to begin your work on a ticket by setting appropriate values in all of the required () fields and relevant non-required fields.

Adding CCs to a ticket

Do not try to add external contacts to the CCs of a ticket, even if the customer asks you to. As a matter of security policy, customers have to add CCs themselves. Direct the customer to the corresponding instructions on how to do so.

You can still add internal contacts (like other SEs, the customer's CSM, ·) yourself.

Handling Large Files

Zendesk has a fixed maximum attachment size of 50MB per file. If you need a user to share a larger file than this, then see [Provide Large Files to GitLab Support](#) for information on how to do so.

Merging tickets

WARNING: Any attached files in the to be merged tickets will be shared across the tickets. Everyone in CC on both of these tickets will receive the files.

When Merging Tickets, leave Requester can see this comment unchecked in the ticket that's being merged into (the second ticket from the top) in order to maintain the SLA. If the merge comment is made public, Zendesk considers it a response and removes the SLA. The ticket that was merged into another ticket is closed while the status of the target ticket is unaffected.

NOTE: Any ticket merge is final -- there is no option to undo it.

ZenGuard - Action Warning System

ZenGuard is a Zendesk application deployed on Global Zendesk only. The app provides warning dialogs and confirmation prompts for high-risk actions that could lead to unwanted results. It helps prevent common mistakes that require tickets to be reopened or recreated.

What ZenGuard protects against

The app provides warnings and blocks certain actions:

- Due date issues: Setting due dates in the past, or too far in the future
- External collaborator risks: Adding end-users as collaborators/CC
- Status change without response: Setting a ticket to pending with only an internal note
- On-hold timer reset: Sending a public reply on an on-hold ticket that will not reset the SLA timer
- Missing public reply: Setting ticket to on-hold without a public reply
- Form change closure: Changing ticket form which will cause automatic closure

How ZenGuard works

When you attempt a potentially problematic action, ZenGuard will:

1. Display a warning dialog explaining the potential issue
2. Block the action if it cannot be bypassed (marked as "this cannot be bypassed")
3. Allow bypass for certain warnings by refreshing the app or proceeding after confirmation, the action blocked can be bypassed

Common scenarios and solutions

Bypassing warnings

Some warnings can be bypassed by either of the following::

- Refreshing the ZenGuard app in your browser
- Clicking through the confirmation dialog (where permitted)

Note that critical safety checks (like setting pending without any response) cannot be bypassed.

How to refresh ZenGuard

1. Click in to the ZenGuard app using either the icon on your apps list (right side of Zendesk), or if you can't see it, click the + to pin a new app and select ZenGuard.
2. Click the reload button in the title of the app.

Troubleshooting

If ZenGuard prevents a legitimate action:

1. Check if there's a bypass option in the warning dialog
2. Ensure all required fields are properly filled
3. For macro-related issues, wait a moment and resubmit
4. If the problem persists, open a new issue in the Support Operations project

Related resources

- RFC: What Zendesk "pitfalls" are we most worried about when working tickets
- Support Operations issue: Feature Request: Zendesk Action Warning System
- Support Deployment issue: ZenGuard App Deployment - Zendesk Action Warning System

SECTION: support/workflows/zendesk_agent_personalization.md

Zendesk agent personalization

As a support engineer, there are several things you can personalize in Zendesk. For example, during your onboarding you are guided to update your profile picture. There are aspects things you can modify by editing your support-team yaml entry. This page describes what items you can personalize, and how to change them.

For more information about the Zendesk fields that are editable via the support-team.yaml, see the support team wiki.

Display name

By default, your display name in Zendesk will be Firstname Lastname. You can choose to personalize your display name that appears in your ticket replies and in your signature, by setting up an alias.

To change your display name to use an alias:

1. Edit your support-team project entry page.
2. Find the section for the instance of Zendesk you want to update (main for Global or us-federal for US Government).
3. Update the alias line to the name you would like to use as an alias.

Be aware that there are other ways that your name may be displayed to customers that you also might want to consider, for example:

Your Zoom display name,
Calendly Personal Link URLs,
GitLab username.

GitLab handle displayed in your Zendesk signature

By default, your GitLab handle will not be shown in your Zendesk Signature. You can choose to have it displayed if you prefer, as follows:

1. Edit your support-team project entry page.
2. Find the section for the instance of Zendesk you want to update (main for Global or us-federal for US Government).
3. Update the showinsignature:, gitlabhandle: section from false to true.

Salutations

You can set a personalized salutation that will always be inserted before your signature with your ticket updates. For example Regards,. We generally recommend omitting this (by not creating a default salutation) and instead choose a salutation that matches the situation in the ticket you are working on. However, if you do want to set this up, you can do as follows:

1. Edit your support-team.yaml page.
2. Find the section for the instance of Zendesk you want to update (main for Global or us-federal for US Government).
3. Update the salutations: line to your preferred phrase.

SECTION: support/workflows/zendesk_organizations_and_users_overview.md

SECTION: support/workflows/regions/_index.md

SECTION: support/workflows/regions/cross_region.md

Overview

Although each region within GitLab Support has its own specific responsibilities each day, and has the freedom to manage those responsibilities in whatever ways work best for them, the regions must also provide a consistent experience to customers. To do that, the regions must coordinate and collaborate with each other.

The purpose of this page is to present clearly the workflows used by the regions in working together and all the information about regional structures that support those workflows.

Regional queue-related responsibilities

Regional queue hours

During regional queue hours, Support Engineers are expected to be actively monitoring and taking new tickets from the queue. These are hours when the region is responsible for the queue. When two regions' hours overlap, they share that responsibility.

Clean queue

Each region is responsible for leaving a clean queue for the next region by the end of their regional queue hours. A clean queue for Region X has no unassigned tickets in it with a preferred region of Region X. Note that if a region has a cleanup hour at the end of its day, that hour will come after their regional queue hours end.

Cleanup hour

A cleanup hour is a time during which the region is not responsible for taking tickets from the queue, and is instead responsible for making sure that their end-of-day activities, such as initiating rehomes, are done.

The purpose of the cleanup hour is to make it reasonable for an SE to take a new ticket from the queue before cleanup hour by having the knowledge that their normal workday includes an hour after their FRT responsibilities end. They would be able to take a new ticket even at the very last minute of

regional queue hours, and still have time to construct and send a good first response during the "free" hour that is "cleanup time."

Cleanup hour will occur for different people in a region at different times during the day, based on each individual's working hours. We only represent cleanup hour in the table at the end of a region's day, with the understanding that for all other hours, the · represents both "regional queue hours" (FRT time) and "individual cleanup hour".

Extra capacity hour

An extra capacity hour is a time during which some Support Engineers in a region might be working on tickets even though the region is not officially responsible for taking new tickets from the queue. We display these hours to help people to know when they might be able to find people working outside of regional queue hours in other regions.

Queue-related responsibilities schedule

This table defines the queue-related responsibilities of Support Engineers in each region for each hour of the day. The hours are given in UTC. Each hour for each region is marked with exactly one of the symbols defined here.

Legend

Symbol	Meaning
·	Regional queue hour
·	Extra capacity hour - some Support Engineers might still be working on tickets
·	Cleanup hour
--	No coverage during that hour

Region	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
APAC	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·
EMEA	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--
AMER	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

SECTION: support/workflows/team/performance_review.md

Overview

Support Engineering currently uses the following performance factor worksheets as an aid to evaluating performance:

- Support Engineer
- Senior Support Engineer
- Staff Support Engineer
- Support Engineering Manager
- Senior Support Engineering Manager

Using the Worksheets: Tips

1. You might want to take notes as you determine each line item in each of your direct reports' worksheets. One good place for those notes is the space in the worksheet between the responsibility (or value or remote work) title and the weighting. This makes it easy to deliver the review to your direct report entirely from the worksheet.

Updating this page

Updates to this page must be reviewed and approved by the VP, a Director or a Senior Manager of Support Engineering before merging.

No changes should be merged from 16 October to 31 December. This is to avoid changes during the Formal Talent Assessment cycle which begins on 1 November and concludes by the end of each year.

Performance factor description

When making changes to this section, make sure that you also change the corresponding entries in the "performance factor worksheet" tab in the corresponding Google Sheet.

Support Engineering Manager

Growth, development and hiring of team members (Weighting: 20%)

- Developing:
 - You are still getting to know your direct reports, and have not yet established a foundation to discuss career or learning goals.
 - You are still getting familiar with your direct reports' career or learning goals, and provide minimal input helping your direct reports set and progress their career or learning goals.

- You're becoming aware of your regional team's hiring needs, and/or are not yet ready to conduct interviews and grade take-home assessments.
- Performing:
 - You encourage and coach your direct reports to set their own career or learning goals in alignment with Support Engineering's business needs.
 - You are aware of your regional team's hiring needs and contribute to hiring activities when the opportunity presents itself. You are ready to conduct interviews and grade take-home assessments when called upon to do so.
- Exceeding:
 - Extends Performing descriptor.
 - You play an active role in your direct reports' growth by connecting them with learning resources, creating opportunities for personal development and partnering with others to remove or smooth over obstacles to growth.
 - You shape hiring strategy and process to mitigate or resolve people resource-related gaps and weaknesses in Support Engineering and partner with Talent Acquisition to ensure effective execution of the hiring plan.

Process development and improvement (Weighting: 20%)

- Developing:
 - You are just starting to build an understanding of Support Engineering processes and workflows, and where friction exists in how we structure and interact with work.
 - You are still inwardly focused and haven't started engaging with your direct reports and your peers to understand the day-to-day challenges of Support Engineering work.
 - You're still finding your voice and aren't fully comfortable expressing your opinions or encouraging constructive participation in public discussions on processes and process improvement.
 - You're still working to get context and understand where you can contribute. As a result you do rarely propose, or play an active part in helping progress, process improvement initiatives.
- Performing:
 - You stay up-to-date on the challenges Support Engineering team members face in their day-to-day work and are able to understand and describe the conflicting priorities that cause them.
 - You take the initiative to surface issues constructively and carefully use your voice as a manager to create psychological safety that enables others on the team to contribute to discussions.
 - You propose and drive improvement initiatives, or play an active role in helping others progress implementation of their initiatives.
- Exceeding:
 - Extends Performing descriptor.
 - You make use of effective iteration and decision making to gain support of and buy in from stakeholders both within Support

Engineering and in the wider GitLab organization to drive improvement initiatives. You carefully monitor the impact of improvement initiatives and are quick to adapt or reverse implementation if the desired outcome is not achieved.

- You coach and enable your direct reports and other team members to own and independently drive improvement initiatives.

People care and performance management (Weighting: 30%)

- Developing:
 - You are just starting to build an understanding of most of your direct reports' productivity, engagement with work and personal well-being. You are just starting to discover their learning styles and coaching needs and as a result may not be able to quickly identify changes or needs in the same.
 - You are beginning to develop an awareness of the quality of your direct reports' work with customers and the impact of their contributions to the Support Engineering team.
 - Identifying and taking action on underperformance is challenging.
- Performing:
 - You keep up to date with your direct reports' productivity, engagement with work and personal well-being, and provide a listening ear, encouragement and coaching when necessary.
 - You coach your direct reports to increase and/or maintain the quality of their work with customers and their contributions to the Support Engineering team.
 - You identify underperformance of your direct reports and act quickly and appropriately to manage it.
- Exceeding:
 - Extends Performing descriptor.
 - You identify issues which might impact your direct reports' productivity, engagement with work and personal well-being and work with your peers and your manager to address the issues before they become problems.
 - You enable your direct reports and amplify their voice where appropriate so they can model being a manager of one and values alignment in the areas of productivity, engagement with work and personal well-being to the Support Engineering team.

Managing customer service out